

bricks Programming Reference v3.4.0

- [About this publication](#)
- [Contents](#)
- [Part 1. The bricks programming model](#)
 - [Chapter 1. Overview](#)
 - [Chapter 2. The EXEC CICS command environment](#)
 - [Chapter 3. The map DSL](#)
- [Part 4. EXEC CICS command reference](#)
 - [Chapter 4. Terminal I/O commands](#)
 - [Chapter 5. Program control commands](#)
 - [Chapter 6. System services](#)
- [Part 5. EXEC CICS file and queue commands](#)
 - [Chapter 7. KSDS file commands](#)
 - [Chapter 8. KSDS browse commands](#)
 - [Chapter 8a. File Control on tmp_dir Sequential Files](#)
 - [Chapter 9. Temporary storage and transient data commands](#)
 - [Chapter 10. Recovery and condition handling](#)
 - [Chapter 11. The Execute Interface Block \(EIB\)](#)
 - [Chapter 12. Response codes](#)
 - [Chapter 13. Commands not implemented](#)
- [Part 7. EXEC CICS WEB command reference](#)
 - [EXEC CICS WEB — server side \(Phase 1\)](#)
 - [EXEC CICS WEB — client side \(Phase 2a\)](#)
 - [EXEC CICS DOCUMENT — chunked body builder](#)
 - [EXEC CICS WEB — Phase 3b additions](#)
- [Part 3. The REXX language](#)
 - [Chapter 14. REXX program structure](#)
 - [Chapter 15. Variables and stems](#)
 - [Chapter 16. Control flow](#)
 - [Chapter 17. PARSE templates](#)
 - [Chapter 18. Conditions and SIGNAL ON](#)
 - [Chapter 19. Built-in functions](#)
- [Part 2. The COBOL language](#)
 - [Chapter 20. COBOL source format](#)
 - [Chapter 21. DATA DIVISION](#)
 - [Chapter 22. PROCEDURE DIVISION](#)
 - [Chapter 23. The EIB block in COBOL](#)
 - [Chapter 24. EXEC CICS in COBOL](#)
 - [Chapter 25. Copybooks](#)
- [Part 6. EXEC SQL command reference](#)
 - [Chapter 26. Embedded SQL \(COBOL and REXX\)](#)
 - [Chapter 27. Restrictions and deferred features](#)
- [Part 8. Sample programs](#)
 - [Chapter 28. Pre-installed sample transactions](#)
 - [Chapter 29. Worked examples](#)
- [Appendix A. Adapting to terminal size \(mod 2 vs mod 4\)](#)

- [Appendix B. Pitfalls and idioms](#)
- [Appendix C. Quick command reference card](#)
 - [EXEC CICS — terminal I/O \(Chapter 4\)](#)
 - [EXEC CICS — program control \(Chapter 5\)](#)
 - [EXEC CICS — system services \(Chapter 6\)](#)
 - [EXEC CICS — KSDS files \(Chapter 7-8\)](#)
 - [EXEC CICS — temporary storage / transient data \(Chapter 9\)](#)
 - [EXEC CICS — recovery + condition \(Chapter 10\)](#)
 - [EXEC SQL \(Chapter 26\)](#)
 - [EXEC CICS WEB — server side \(Phase 1\)](#)
 - [EXEC CICS WEB — client side \(Phase 2\)](#)
 - [EXEC CICS DOCUMENT](#)
 - [Standard copybooks](#)
 - [EIB fields \(Chapter 11\)](#)
 - [SQLCA fields \(Chapter 26\)](#)

BRICKS

Transaction Server

Application Programming Reference

Order number: BX26-1500-01

Edition: First Edition

Date: 2026

Applies to: Bricks Transaction Server, version 1.5 and later

Document	Bricks Application Programming Reference
Audience	Application programmers writing REXX or COBOL transactions for Bricks.
Companion	README.md — installation, configuration, and operations.
Operator manual	ISPF_editor.md — the in-3270 source editor reference.

About this publication

This reference describes the application-programming interface of the **Bricks Transaction Server**: the COBOL and REXX languages it accepts, the EXEC CICS command set those programs may issue, the EXEC SQL embedded database surface, the EXEC CICS WEB command family for HTTP service, the BMS-flavoured map DSL used to build 3270 panels, the VSAM and temporary-storage queue surfaces, and the catalogue of sample programs shipped under runtime/.

Operational and installation topics — running bricks, editing bricks.cnf, signing on through CSSN, the CEMT master terminal, the bricksload stress tester, and the /metrics endpoint — are documented in the companion **README.md**.

Who should read this publication

This manual is for application programmers who write transactions that run under Bricks. Familiarity with the IBM CICS programming model (pseudo-conversational dispatch, the **EIB**, **COMMAREA**, BMS maps, KSDS files, temporary-storage queues, the SQLCA, and embedded SQL) is assumed. Where Bricks deviates from the IBM behaviour the difference is documented explicitly.

Conventions used in this publication

The following notational conventions appear throughout the manual.

Symbol	Meaning
UPPERCASE	A keyword. Coded literally.
lowercase	A value supplied by the programmer (a name or expression).
[]	An optional clause.
{ a b }	
...	The preceding clause may repeat.

Command syntax is shown in the EXEC CICS ... END-EXEC form. The bare-string form ("VERB OPTIONS" under ADDRESS CICS, REXX only) is described in *Chapter 2. The EXEC CICS command environment*; both forms dispatch identically.

Reference layout

Each command described in this manual follows the same five-section layout, in this order:

1. **Format** — the command syntax as a code block.
2. **Description** — what the command does and any important runtime behaviour.
3. **Options** — every keyword that follows the verb, with its meaning and any value constraints.
4. **Conditions** — the EIBRESP values (or SQLCODE values, for EXEC SQL) the command may return, together with the cause of each.
5. **Example** — a short, runnable fragment illustrating typical use.

Programming Note. Where the same verb behaves differently in COBOL and in REXX, the difference is called out in a *Programming Note* like this one. Where both languages share identical behaviour, the example uses whichever form reads most clearly.

Contents

This publication is organised into nine parts, an appendix series, and a quick-reference card. Read **Part 1** first; the remaining parts are reference material consulted as needed.

Part 1. Bricks programming principles

- [Chapter 1. Overview](#)
- [Chapter 2. The EXEC CICS command environment](#)
- [Chapter 3. The map DSL](#)

Part 2. The COBOL language

- [Chapter 20. COBOL source format](#)
- [Chapter 21. DATA DIVISION](#)
- [Chapter 22. PROCEDURE DIVISION](#)
- [Chapter 23. The EIB block in COBOL](#)
- [Chapter 24. EXEC CICS in COBOL](#)
- [Chapter 25. Copybooks](#)
- [Chapter 27. Restrictions and deferred features](#)

Part 3. The REXX language

- [Chapter 14. REXX program structure](#)
- [Chapter 15. Variables and stems](#)
- [Chapter 16. Control flow](#)
- [Chapter 17. PARSE templates](#)
- [Chapter 18. Conditions and SIGNAL ON](#)
- [Chapter 19. Built-in functions](#)

Part 4. EXEC CICS command reference

- [Chapter 4. Terminal I/O commands](#) — SEND MAP, RECEIVE MAP, CONVERSE, SEND TEXT, RECEIVE
- [Chapter 5. Program control commands](#) — RETURN, XCTL, LINK, ABEND, START, RETRIEVE
- [Chapter 6. System services](#) — ASSIGN, ASKTIME, FORMATIME, QUERY SECURITY, VERIFY PASSWORD, INQUIRE SYSTEM, ADDRESS, GETMAIN/FREEMAIN, INQUIRE/SET (FILE/TASK/TRANSACTION/TERMINAL/PROGRAM)
- [Chapter 10. Recovery and condition handling](#) — SYNCPOINT, SYNCPOINT ROLLBACK, HANDLE CONDITION, IGNORE CONDITION, HANDLE AID, HANDLE ABEND
- [Chapter 11. The Execute Interface Block \(EIB\)](#)
- [Chapter 12. Response codes](#)
- [Chapter 13. Commands not implemented](#)

Part 5. EXEC CICS file and queue commands

VSAM-style KSDS files and temporary-storage / transient-data queues.

- [Chapter 7. KSDS file commands](#) — READ, WRITE, REWRITE, DELETE
- [Chapter 8. KSDS browse commands](#) — STARTBR, READNEXT, READPREV, RESETBR, ENDBR
- [Chapter 8a. File Control on tmp_dir Sequential Files](#) — browse + positional READ against ESDS-style sequential files in the tmp_dir sandbox
- [Chapter 9. Temporary storage and transient data commands](#) — READQ TS/TD, WRITEQ TS/TD, DELETEQ TS/TD, the tmp_dir sandbox

Part 6. EXEC SQL command reference

- [Chapter 26. Embedded SQL \(COBOL and REXX\)](#) — SELECT INTO, INSERT, UPDATE, DELETE, COMMIT, ROLLBACK, CONNECT TO, cursors (DECLARE / OPEN / FETCH / CLOSE), WHENEVER, null indicators, the SQLCA, the SQLCODE catalogue, SYNCPOINT integration.

Part 7. EXEC CICS WEB command reference

- [EXEC CICS WEB](#) — server side
- [EXEC CICS WEB](#) — client side
- [EXEC CICS DOCUMENT](#) — chunked body builder
- [EXEC CICS WEB](#) — Phase 3b additions

Part 8. Sample programs

- [Chapter 28. Pre-installed sample transactions](#)
- [Chapter 29. Worked examples](#)

Appendixes

- [Appendix A. Adapting to terminal size \(mod 2 vs mod 4\)](#)
- [Appendix B. Pitfalls and idioms](#)
- [Appendix C. Quick command reference card](#)

Note on physical order. The chapter numbers in this contents list are the canonical numbers used throughout the publication. Because some parts have been re-ordered into the logical grouping above, a chapter you find on Part 4 of the contents may sit physically later in the file

than a chapter listed in Part 5. Cross-references everywhere in the manual use the chapter number, not the file position.

Part 1. The bricks programming model

Chapter 1. Overview

A bricks **transaction** is a 4-character TRANSID listed in `runtime/transactions.conf`. Each transaction names a **program** — either a REXX source file in `runtime/rexx/` or a COBOL source file in `runtime/cobol/` — and an optional access-control list:

```
HEL0:rexx:hello.rexx
HELP:rexx:help.rexx:public
QAGE:rexx:qage.rexx:public,users,admin
HELC:cobol:hello.cob:public
GUST:cobol:gust.cob:public
```

A program that runs EXEC SQL against something other than the **first** database in `runtime/databases.conf` adds an explicit binding as the 5th colon-separated field:

```
BANK:cobol:bank.cob:public,users,admin:bank
```

See *Database binding* in [Chapter 26](#) for the full contract; the short version is that an absent 5th field defaults to the first row of `databases.conf` and a misbinding manifests as `SQLCODE = -204` (relation does not exist) on every dispatch.

Organising programs into sub-directories

Once a deployment grows past a handful of programs, the flat `runtime/rexx/ / runtime/cobol/` layout becomes hard to scan. You can group programs into sub-directories at any depth and reference them with a relative path from the language root:

```
runtime/cobol/billing/invoice.cob
runtime/cobol/billing/statement.cob
runtime/cobol/cards/visa/authz.cob
runtime/rexx/admin/userprov.rexx
```

```
INVC:cobol:billing/invoice.cob:public
STMT:cobol:billing/statement.cob:public
AUTH:cobol:cards/visa/authz.cob:admin
USRP:rexx:admin/userprov.rexx:admin
```

Forward slashes work on both POSIX and Windows. CEDA PROGRAM walks the language roots recursively and shows each file with its relative path (so `billing/invoice.cob` appears in the **FILE** column, matching what you write in `transactions.conf`). CEMT INQUIRE TRANSACTION shows the same form in the **PROGRAM** column, eliding a long middle (`billing/.../authz.cob`) so the leading application and the file basename always stay visible. Hidden entries (dotfile-prefixed, e.g. `.git/`, `.swp` files) are skipped during the catalogue walk; recursion is capped at 8 levels.

CEDA TRANSACTION accepts a relative subdirectory path in the **PROGRAM** form field. `..` traversal is rejected; so are empty path segments (`a//b.rexx`). Absolute POSIX paths are accepted unchanged for the rare case of running one-off programs from outside the runtime tree. A path containing a `:`, a newline or a whitespace-preceded `#` is **refused** — including a Windows-style `C:\app\prog.cob`, whose drive colon used to write a `transactions.conf` line the parser read back as a different record, shifting the groups and

database fields one slot to the right and leaving the transaction unloadable, unreachable, and bound to a nonexistent database.

`brickscompile` accepts either a single file or a directory. In directory mode it walks recursively (same depth cap + hidden-skip rules as `CEDA PROGRAM`) and reports per-file pass/fail; exit code is 0 only if every file parses cleanly.

Program caching

The first time a TRANSID is dispatched, its program is parsed and the AST is cached (see *Parsed-program cache* in the README); subsequent dispatches skip the parse. Each running task has its own heap and stack — there is no shared mutable state between concurrently running tasks of the same TRANSID.

Programs interact with three layers:

1. **The 3270 terminal**, through `EXEC CICS SEND MAP / RECEIVE MAP` (and the unstructured `SEND TEXT / RECEIVE` pairing). Maps are defined in the BMS-style DSL described in [Chapter 3](#).
2. **Persistent data**, through `KSDS` file commands ([Chapter 7](#)) and `TS-queue` commands ([Chapter 9](#)). Both are backed by an embedded `bbolt` B+tree at `data/files.boltdb`.
3. **Other programs**, through `LINK` (synchronous, with `COMMAREA`), `XCTL` (transfer of control, no return), and `RETURN TRANSID` (pseudo-conversational chaining). The Execute Interface Block (EIB) exposes session and request state ([Chapter 11](#)).

REXX and COBOL share **the same EXEC CICS dispatcher** (`cics/handler.go`). Every command documented in [Part 2](#) is therefore available, with identical syntax and identical semantics, from either language.

A typical pseudo-conversational task has the shape:

```
ADDRESS CICS

EXEC CICS RECEIVE MAP('CUST1')           END-EXEC  /* read prior screen */
... validate / compute ...
EXEC CICS SEND    MAP('CUST1') FROM(SCR.) ERASE END-EXEC
EXEC CICS RETURN  TRANSID('CUST')        END-EXEC  /* chain back to self */
```

The dispatcher invokes the program once per ENTER. State that must survive between invocations rides in the `COMMAREA` on `RETURN`; the chained task receives it through `EIBCALEN` and the `DFHCOMMAREA` data area.

Chapter 2. The EXEC CICS command environment

Surface forms

Bricks accepts two equivalent surface forms inside an `ADDRESS CICS` scope.

1. IBM-canonical `EXEC CICS ... END-EXEC` — the form CICS programmers recognize. Available from REXX and COBOL.

```
ADDRESS CICS

EXEC CICS ASSIGN  USERID(USR) TERMID(TRM)      END-EXEC
EXEC CICS SEND    MAP('HEL01') FROM(SCR.) ERASE END-EXEC
EXEC CICS RETURN                                     END-EXEC
```

In REXX, a small preprocessor (`rexx/preprocess.go`, called from `rexx.Parse`) rewrites each `EXEC CICS ...` `END-EXEC` block into the equivalent quoted-string command before lexing. Multi-line bodies are collapsed into a single command string; trailing newlines preserve source line numbers in error messages. Comments and string literals are left alone.

In COBOL, the parser collects every token between `EXEC CICS` and `END-EXEC` and reconstructs the body verbatim for the same `cics.ParseCommand` REXX uses.

2. Bare string under `ADDRESS CICS` — terser; available **only in REXX**.

```
ADDRESS CICS
"ASSIGN USERID(USR) TERMID(TRM)"
"SEND MAP('HEL01') FROM(SCR.) ERASE"
"RETURN"
```

Both forms route through `cics.ParseCommand` and dispatch to the same handler. After every command the handler writes `EIBRESP`, `EIBRESP2`, and `RC` into the program's frame.

Argument-passing rules

Inside the parentheses of an option:

- A bare identifier is treated as a **variable name**. The handler reads the value at runtime, and (where the option is an output) writes the result back into that variable.
- A quoted string literal (`'...'` or `"..."`) is treated as a **literal value**. The handler does not write to literals; passing a literal to an output-only option is rejected with `INVREQ`.

This distinction matters most for length and key fields: `LENGTH(LEN)` both reads the requested length on input and writes the actual length back; `LENGTH(80)` only sets the input length.

Programming model summary

Concept	Bricks implementation
Task	One invocation of one <code>TRANSID</code> (<code>session.TxCB</code>).
Program	A parsed REXX or COBOL source file, cached at L1/L2.
Pseudo-conversational chain	<code>EXEC CICS RETURN TRANSID(t) COMMAREA(d)</code> .
Conversational	Loop in the program; each <code>RECEIVE MAP</code> blocks on input.
Inter-program call	<code>LINK PROGRAM(name) COMMAREA(var)</code> (synchronous).
Storage shared across tasks	KSDS files, TS queues.
Storage shared within one task	REXX variables / COBOL <code>WORKING-STORAGE</code> only.
Implicit task end	Program runs off the end, or <code>STOP RUN</code> (COBOL).
Forced task end	<code>EXEC CICS ABEND</code> .

Response-code handling (`RESP`, `RESP2`, `DFHRESP`)

Every `EXEC CICS` verb accepts the IBM-canonical `RESP(var)` and `RESP2(var)` clauses. After dispatch the runtime writes the numeric response into the named host variables, in addition to the always-set `EIBRESP` / `EIBRESP2` frame fields. Programs may therefore write either style:

```
EXEC CICS READ FILE('CUST') INTO(REC) RIDFLD(KEY)
                                RESP(WS-RC) END-EXEC.
IF WS-RC = DFHRESP(NORMAL)    PERFORM USE-REC.
IF WS-RC = DFHRESP(NOTFND)    PERFORM CREATE-REC.

EXEC CICS WRITE FILE('LOG') FROM(BUF) RIDFLD(KEY) END-EXEC.
IF EIBRESP = DFHRESP(DUPREC) PERFORM HANDLE-DUP.
```

DFHRESP(NAME) is a compile-time function: the COBOL parser resolves it to the numeric constant from `runtime/cobolcopy/DFHRESP.cpy` (and from the matching table in `cobol/parser.go`). An unknown name is a parse-time error with a "see DFHRESP.cpy for the list" hint.

Length as an input bound

LENGTH(n) on READQ TS, RECEIVE INTO, LINK COMMAREA, and RETURN COMMAREA is treated as IBM's in/out parameter — the caller's value caps the bytes returned (or passed forward, for LINK/RETURN); the post-dispatch write-back stores the actual byte count. A LENGTH(80) on READQ TS against a 200-byte item now hands the program 80 bytes and reports 80, matching real CICS instead of overflowing the buffer.

On READ FILE, READNEXT, and READPREV (file-record family), LENGTH is OUTPUT-only — the store decides record size; there is no operator-supplied truncation bound. This is a bricks-defensible simplification for the KSDS-fixed-length model; real VSAM variable-length permits LENGTH-as-max-buffer. The write-back-then-read-as-input pattern (call 1 writes len(rec1) into LEN, call 2 re-reads LEN as a cap) would silently corrupt subsequent reads, so bricks ignores the input value entirely on these verbs.

Chapter 3. The map DSL

Bricks ships its own line-oriented, BMS-flavoured map DSL (parsed by `mapdsl/`). Maps live as `*.map` files in the directory configured by `maps_dir` (default `runtime/map/`). Each file contains one or more MAP ... ENDMAP blocks. Comments start with `*`.

Format

```
[MAPSET <name>]                -- optional, namespace for the file
MAP <name> SIZE rowsxcols
  [FIELD AT r,c LEN n <attrs> "literal"]
  [INPUT <fieldname> AT r,c LEN n <attrs> [DEFAULT "value"]]
  [STOP AT r,c]
  [CURSOR AT {fieldname | r,c}]
ENDMAP
[MAP <name> SIZE rowsxcols      -- multiple MAP blocks per file
  ...
ENDMAP]
```

Statements

MAPSET <name> (*optional, must be the first non-comment line*) Names the mapset — the namespace every MAP block in this file belongs to. The catalog uses MAPSET.NAME as its key so two different files can declare a map named MAPA without colliding. When the line is omitted, the mapset defaults to the file's basename (uppercased, extension stripped): a file named `help1.map` has implicit mapset `HELP1`.

MAP <name> SIZE rowsxcols The map header. Names must be unique **within their mapset** (case-insensitive); typical sizes are 24x80 (mod 2) and 43x80 (mod 4). A single .map file may declare any number of MAP blocks back-to-back — useful when porting a BMS mapset that contained several DFHMDI macros under one DFHMSD.

FIELD AT r,c LEN n <attrs> "literal" A display-only field that paints the literal at row r, column c, length n. Useful for labels, headings, and panel chrome.

INPUT <fieldname> AT r,c LEN n <attrs> [DEFAULT "value"] A named input field. The name is how SEND MAP FROM(STEM.) and RECEIVE MAP INTO(STEM.) route data to / from this position (STEM.<fieldname>). Use PROT to make the field write-protected but still named — useful for output values painted by SEND MAP FROM(STEM.).

STOP AT r,c An autoskip stop attribute that ends the preceding input field.

CURSOR AT <fieldname> or CURSOR AT r,c The home position for the cursor when the map is sent. The named form resolves to (field.Row, field.Col + 1) — i.e. one byte to the right of the leading 3270 attribute byte, which is the writable cell. Forward references are allowed; names are resolved after the whole map is parsed. When omitted, the renderer falls back to the first input field.

ENDMAP Terminates the map.

Converting legacy BMS sources

The bricksconvert CLI utility (under cmd/bricksconvert) converts IBM CICS **BMS map source** (DFHMSD / DFHMDI / DFHMDF macros) into this DSL — the recommended path for porting an existing CICS application's screens onto bricks without rewriting every panel by hand. See the [BMS conversion section in the README](#) for usage details.

Attributes

PROT, UNPROT, BRIGHT, DIM, UNDERLINE, HIDDEN, NUMERIC, MDT, BLINK, REVERSE,
COLOR=BLUE|RED|PINK|GREEN|TURQUOISE|YELLOW|WHITE.

UNDERLINE is the canonical IBM BMS spelling (HIGHLIGHT=UNDERLINE). UNDERSCORE is accepted as a legacy alias for hand-written bricks DSL maps that already use it; new maps should prefer UNDERLINE.

Catalogue lifecycle

mapdsl.NewCatalog(dir) parses every *.map file at startup and self-refreshes on subsequent edits. Each Lookup(name) stats the directory plus the source file backing the requested name; on mtime change the directory is reparsed and swapped atomically. A failed re-parse keeps the prior catalogue in place — the operator can find the broken file with CEMT PERFORM RESCAN MAP (see the README).

Sending a map by mapset

EXEC CICS SEND MAP('NAME') MAPSET('SET') FROM(STEM) resolves through the qualified MAPSET.NAME index, so:

```
EXEC CICS SEND MAP('MAPA') MAPSET('AM01') FROM(AM01) ERASE END-EXEC.  
EXEC CICS SEND MAP('MAPA') MAPSET('AM16') FROM(AP16) ERASE END-EXEC.  
EXEC CICS SEND MAP('MAPB') MAPSET('AM99') FROM(AM99) ERASE END-EXEC.
```

each picks the right map even when several files share the same bare map name. When `MAPSET(...)` is omitted, the catalog tries a bare-name lookup; if that name appears in more than one mapset the runtime returns the short diagnostic `SEND MAP: map "MAPA" is ambiguous; specify MAPSET` rather than silently picking one. Programs that only ever reference unique bare names (the typical case for the runtime/map shipped with bricks) keep working without any code change.

`RECEIVE MAP` accepts the same `MAPSET(...)` clause with identical resolution semantics.

Example

```
* CICS sign-on screen
MAP CSSN SIZE 24x80
  FIELD AT 1,28 LEN 24 PROT BRIGHT COLOR=TURQUOISE  "BRICKS SIGN-ON"
  FIELD AT 6,15 LEN 9  PROT                                "Userid:"
  INPUT USERID  AT 6,25 LEN 8 UNDERSCORE COLOR=GREEN
  STOP          AT 6,34
  FIELD AT 8,15 LEN 9  PROT                                "Password:"
  INPUT PASSWORD AT 8,25 LEN 8 HIDDEN UNDERSCORE COLOR=RED
  STOP          AT 8,34
  CURSOR        AT USERID
ENDMAP
```

Part 4. EXEC CICS command reference

Part 4 begins here. This part documents the core non-file, non-SQL, non-WEB EXEC CICS verbs. File and queue commands are in **Part 5**; EXEC SQL is in **Part 6**; EXEC CICS WEB is in **Part 7**.

This part documents every EXEC CICS command bricks implements. Each command page follows the same layout: **Format**, **Description**, **Options**, **Conditions**, **Example**.

Commands are grouped by function:

- [Chapter 4. Terminal I/O](#)
- [Chapter 5. Program control](#)
- [Chapter 6. System services](#)
- [Chapter 7. KSDS files](#)
- [Chapter 8. KSDS browse](#)
- [Chapter 9. Temporary storage](#)

Cross-cutting reference:

- [Chapter 10. Recovery and condition handling](#)
- [Chapter 11. The EIB](#)
- [Chapter 12. Response codes](#)
- [Chapter 13. Commands not implemented](#)

Chapter 4. Terminal I/O commands

SEND MAP

Send a BMS-style map to the 3270 terminal and wait for the operator to press an AID key.

Format

```
EXEC CICS SEND MAP(name)
                [MAPSET(set)] [FROM(stem.)]
                [ERASE | ERASEAUP | DATAONLY] [MAPONLY]
                [CURSOR(position)]
                [FREEKB] [ALARM] [FRSET]
                [RESP(var)] [RESP2(var)]
END-EXEC
```

Description

Loads the named map from the catalogue, populates each named field with the value of <stem>. <fieldname> (when FROM is supplied), paints the screen, and waits for the operator to press an AID key (ENTER, PFn, PAn, CLEAR). On return, the captured response is stored on the TCB so a subsequent RECEIVE MAP can pull modified field values back into the program. EIBAID and EIBCP05N are updated with the AID character and 1-based cursor position.

IBM-canonical clauses honoured: MAPSET(set) qualifies the lookup to disambiguate same-named maps across mapsets (see Chapter 3); ERASEAUP wipes unprotected fields only (approximated as ERASE in bricks); MAPONLY paints just the map's literals + INPUT defaults, ignoring any FROM(stem). FREEKB, ALARM, and FRSET are accepted for source compatibility — FREEKB and FRSET are subsumed by the fixed WCC the go3270 library writes; ALARM is parsed but not yet tunnelled to the wire.

If FROM is omitted, the map renders using its INPUT DEFAULT values. If FROM is a literal string, every named field on the map is filled with that literal (rare; supported for parity with BMS).

Three paint modes

Flag	Field set	Blocking?	When to use
ERASE	Every field in the map	Yes — waits for AID	Initial paint, full refresh.
(neither)	Only fields the program populated in FROM(stem)	Yes — waits for AID	In-place update inside a conversational loop (rare; usually ERASE is fine).
DATAONLY	Only fields the program populated in FROM(stem)	No — returns immediately	Background refresh while a different task is mid-input.

ERASE clears the screen first and emits **every** field in the map (with either the program's FROM value or the map default). Without ERASE, bricks does a **partial repaint**: only fields the program explicitly populated in FROM(stem) are sent; fields the program didn't touch stay on the terminal exactly as the operator left them. Use ERASE for the initial paint of a screen or a full refresh; omit it for in-place partial updates.

DATAONLY — background partial paint

DATAONLY is the bricks idiom for **background, fire-and-forget partial paint**, used to refresh part of the screen — typically a clock or other status fields — while a different task on the same terminal is blocking inside SEND MAP waiting for the operator's next keystroke. It layers two things on top of the plain no-ERASE partial-paint behaviour:

1. **No-block**: DATAONLY makes SEND MAP return immediately after writing the screen, without entering the AID-wait. The call is fire-and-forget — the operator's next AID is captured by whatever other SEND MAP or RECEIVE MAP happens to be blocking for it.
2. **Scheduler-safe**: when an EXEC CICS START INTERVAL(...) timer fires on a terminal whose current task is in input-wait, the bricks scheduler dispatches the queued task **inline** in the timer goroutine instead of poisoning the conn read deadline (which would abend the in-flight task with an i/o timeout). The inline-dispatched task is restricted to write-only verbs: SEND MAP DATAONLY, file I/O, TS queue, ASSIGN, START. Any blocking verb (RECEIVE MAP, SEND MAP without DATAONLY, CONVERSE) returns INVREQ in that context — the scheduler goroutine must not block on conn.Read.

This together lets a transaction self-schedule a periodic refresh via:

```
EXEC CICS START TRANSID('CHAT') INTERVAL(000002) FROM('TICK') END-EXEC
EXEC CICS RETURN END-EXEC
```

When the START fires, the same transaction is dispatched a second time. The dispatched copy detects "I am a tick" by checking the RETRIEVE payload, paints only the fields it wants to refresh (SEND MAP ... DATAONLY — the operator's input field is not in the populated set, so it's never repainted), schedules the next tick, and returns. The main task on the same terminal stays blocked in its outer SEND MAP throughout. See runtime/rexx/chat.rexx for the worked example.

DATAONLY and ERASE are mutually exclusive — the runtime rejects the combination at parse time. Note: bricks does not implement the real-CICS DATAONLY attribute-byte semantics (application-supplied attribute bytes via X'00' sentinels); programs that need a synchronous partial repaint use the plain no-ERASE form, and programs that need a background refresh use DATAONLY with the START-INTERVAL pattern above.

Options

MAP(name) — *required* Name of a map in the catalogue (case-insensitive). MAPFAIL if absent.

FROM(stem.) Stem variable whose tails name the map's input fields. The stem's trailing dot is optional: FROM(SCR) and FROM(SCR.) are equivalent.

ERASE Clears the screen before painting AND emits every field in the map. Without this flag, only fields the program populated in FROM(stem) are sent — fields the program didn't touch stay on the terminal unchanged (partial repaint).

DATAONLY Fire-and-forget partial paint: SEND returns immediately without waiting for an AID. Implies no-ERASE partial-repaint semantics. Mutually exclusive with ERASE.

CURSOR(position) 1-based cursor position. Overrides the map's CURSOR AT clause.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Map sent, AID received.
MAPFAIL	36	Named map not found in the catalogue.
IOERR	17	The terminal write failed.

Example

```
SCR.GREETING = 'Hello, ' || USR
EXEC CICS SEND MAP('HEL01') FROM(SCR.) ERASE END-EXEC
```

Runtime field attribute overrides

A map's COLOR= clause is the *default* colour for each field. Programs can override the colour of any named field at runtime by setting a sibling element named <FIELD>-C in COBOL or <stem>.<FIELD>_C in REXX, before the SEND MAP. The override wins; an empty / unset value means "use map default."

COBOL

```
COPY DFHCOLOR.
01 SCR.
   05 STATUS    PIC X(10).
   05 STATUS-C  PIC X(9) VALUE SPACES.
...
IF BALANCE < 0
  MOVE DFHRED   TO STATUS-C
ELSE
  MOVE DFHGREEN TO STATUS-C
END-IF.
EXEC CICS SEND MAP('BALC') FROM(SCR) ERASE END-EXEC.
```

REXX

```
if balance < 0 then scr.status_c = 'RED'
    else scr.status_c = 'GREEN'
'EXEC CICS SEND MAP(BALC) FROM(SCR.) ERASE'
```

Accepted colour mnemonics: BLUE RED PINK GREEN TURQUOISE YELLOW NEUTRAL DEFAULT (also the DFH-prefixed canonical names from DFHCOLOR.cpy: DFHBLUE DFHRED DFHPINK DFHGREEN DFHTURQ DFHYELLO DFHNEUTR DFHDFT). Unknown strings degrade to the device default.

REXX symbol caveat. REXX evaluates stem tails by value, not literal: if your program has assigned a bare variable STATUS_C anywhere, then scr.status_c = 'RED' binds to the value of STATUS_C, not the literal tail STATUS_C. Defensive form is scr.'STATUS_C' = 'RED' (quoted tail).

Highlight (-H) and attribute byte (-A) overrides follow the same convention and are planned for a follow-up release.

See runtime/cobol/balc.cob and runtime/rexx/balc.rexx for complete worked examples.

RECEIVE MAP

Retrieve the operator's input from the most recent SEND MAP.

Format

```
EXEC CICS RECEIVE MAP(name)
                        [INTO(stem.)]
END-EXEC
```

Description

Pulls the response stored by the most recent SEND MAP into a stem, one tail per named field on the map (<stem>.<fieldname>). When INTO is omitted, the default stem is MAP..

A RECEIVE MAP with no matching prior SEND MAP returns MAPFAIL.

Options

MAP(name) — *required* The same map name passed to the prior SEND MAP.

INTO(stem.) Destination stem. Default MAP..

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Fields retrieved.
MAPFAIL	36	No prior matching SEND MAP, or named map not found.

Example

```
EXEC CICS RECEIVE MAP('CUST1') INTO(IN.) END-EXEC
ACTION = IN.ACTION
CKEY   = IN.CUSTNO
```

CONVERSE

A single verb that fuses SEND MAP and RECEIVE MAP against the same map. Semantically identical to issuing the two halves back-to-back; provided because real-world CICS programs written between the 1970s and the 1990s used CONVERSE heavily, and porting them shouldn't require a SEND/RECEIVE rewrite.

Format

CONVERSE has two IBM-canonical forms — the 3270 mapped form and the raw-text TS form. Bricks supports both.

Mapped form (3270):

```
EXEC CICS CONVERSE MAP(name) [MAPSET(set)]
                        [FROM(area) | FROMMAP(area)]
                        [INTO(area)]
                        [ERASE] [CURSOR(pos)]
                        [RESP(var)] [RESP2(var)]
END-EXEC
```

Text form (TS):

```
EXEC CICS CONVERSE FROM(area) [FROMLENGTH(n)]
                    INTO(target) [TOLength(var) | MAXLENGTH(n)]
                    [STRFIELD] [DEFRESP]
                    [RESP(var)] [RESP2(var)]
END-EXEC
```

In the mapped form: FROM(area) and FROMMAP(area) are accepted spellings for the outbound stem; either works, both feed the SEND half. INTO(area) is the inbound stem; it feeds the RECEIVE half. MAPSET, ERASE, CURSOR pass through to both halves.

In the text form (no MAP): bricks dispatches SEND TEXT then RECEIVE INTO. FROMLENGTH(n) caps the outbound bytes; TOLength(var) / MAXLENGTH(n) cap the inbound buffer. STRFIELD (structured-field stream) and DEFRESP (definite- response) are accepted silently — bricks doesn't tunnel either wire bit today but program source ports cleanly.

Behaviour

CONVERSE issues SEND MAP first, then RECEIVE MAP against the same MAP(name). If the SEND half returns anything other than RESP-NORMAL, the RECEIVE half is **skipped** and the SEND's response code is returned — matches real CICS semantics (no screen ever made it to the terminal, so there's no operator response to wait for).

Example

```
EXEC CICS CONVERSE MAP('TIM1') FROM(SCR) INTO(SCR)
                        ERASE
END-EXEC.
IF EIBAID = PF03 THEN
    EXEC CICS RETURN END-EXEC
END-IF.
```

This replaces the longer pair:

```
EXEC CICS SEND MAP('TIM1') FROM(SCR) ERASE END-EXEC.
EXEC CICS RECEIVE MAP('TIM1') INTO(SCR)      END-EXEC.
```

Live example: `runtime/cobol/timc.cob` uses CONVERSE for both the input screen and the reminder screen.

SEND TEXT

Free-form text output without a BMS map.

Format

```
EXEC CICS SEND TEXT FROM(area)
                        [LENGTH(n)]
                        [ERASE]
END-EXEC
```

Description

Real-CICS SEND TEXT. The body is treated as a flat row-major buffer: every `cols` bytes (default 80) lands on the next row, starting at row 0. **3270 has no LF/CR**, so programs must pad each logical line to the column width and concatenate (REXX `LEFT(s,80)`, COBOL `PIC X(80)` group children).

`LENGTH(n)` truncates the body to `n` bytes; bytes past `rows*cols` are dropped. `ERASE` clears the screen before painting; without it, existing fields stay behind. Like `SEND MAP`, the call paints synchronously and waits for an AID, so the operator has time to read the screen before `RETURN`.

Works identically from REXX and COBOL.

Options

FROM(area) — *required* The buffer to paint. Can be a REXX variable or a COBOL group item.

LENGTH(n) Truncate the body to this length.

ERASE Clear the screen first.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Body painted, AID received.
INVREQ	16	FROM missing.

Condition	EIBRESP	Cause
IOERR	17	The terminal write failed.

Example

See [Chapter 29. Worked examples](#), the GETC program, for the canonical multi-row SEND TEXT pattern.

RECEIVE

Retrieve the unedited terminal line typed by the operator at the blank prompt — typically used to pick up command-line arguments.

Format

```
EXEC CICS RECEIVE INTO(buffer)
                    [LENGTH(len)]
END-EXEC
```

Description

Returns the unedited terminal line the operator typed at the blank prompt (TRANSID prefix included), e.g. EXAM 1 2 3. Single-shot per task: a second RECEIVE in the same task, or any RECEIVE in a chained RETURN TRANSID task, returns E0C (RESP=6).

When LENGTH(var) is a bare variable, the actual byte count of the line is written back. The receiving COBOL field's PIC X(n) width handles padding/truncation via standard MOVE semantics; programs detect truncation by comparing len against n.

Options

INTO(buffer) — *required* Destination variable for the line.

LENGTH(len) When a bare variable, receives the actual length on return.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Line returned.
E0C	6	No input available (already consumed in this task chain).

Example

```
EXEC CICS RECEIVE INTO(BUF) LENGTH(LEN) END-EXEC    /* "GETC 100" */
PARSE VAR BUF TID CKEY .

EXEC CICS RECEIVE INTO(WS-INPUT) LENGTH(WS-LEN) END-EXEC
UNSTRING WS-INPUT DELIMITED BY ' '
      INTO WS-TID WS-A WS-B WS-C
END-UNSTRING.
```

Chapter 5. Program control commands

RETURN

End the current task; optionally chain to another TRANSID with a COMMAREA.

Format

```
EXEC CICS RETURN [TRANSID(id)]  
                  [COMMAREA(data)]  
END-EXEC
```

Description

Sets `tcb.NextTransid` and `tcb.Commarea` and ends the task. With no TRANSID, control falls back to the blank prompt. With a TRANSID, the dispatcher invokes that TRANSID next, with the supplied COMMAREA available to the new task as DFHCOMMAREA and its length as EIBCALEN.

COBOL program termination & the end-of-task syncpoint. EXEC CICS RETURN is the canonical way a CICS program ends, and the only one that carries the pseudo-conversational TRANSID / COMMAREA handoff. STOP RUN, GOBACK, EXIT PROGRAM, and running off the end of the PROCEDURE DIVISION also end the task cleanly (bricks tolerates them LE-style), and bricks runs the **implicit end-of-task syncpoint** — which commits the task's EXEC SQL unit of work — for those paths too, so a program that mutates SQL and then STOP RUNs does **not** lose the update. STOP RUN is nonetheless the CICS anti-pattern: it carries no pseudo-conversational handoff and, in a CALLED sub-program, would terminate the whole run unit. Use EXEC CICS RETURN in CICS code (with GOBACK, not STOP RUN, as the unreachable post-RETURN terminator the compiler wants).

Options

TRANSID(id) The next TRANSID to dispatch.

COMMAREA(data) Bytes to flow into the next task.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Task ended.

Example

```
EXEC CICS RETURN TRANSID('CUST') COMMAREA(SAVEAREA) END-EXEC
```

XCTL

Transfer control to another program. The current program does not resume.

Format

```
EXEC CICS XCTL PROGRAM(name)
          [COMMAREA(data) [LENGTH(n)]]
          [RESP(var)] [RESP2(var)]
END-EXEC
```

LENGTH(n) on COMMAREA is the IBM-canonical input bound — caller's value caps the bytes passed to the receiving program. CHANNEL(...) (the CICS-TS container alternative to COMMAREA) is rejected with a short "not supported in bricks (use COMMAREA)" message — single-region bricks doesn't model channels.

Description

Transfer of control. The named program runs in place of the current one; the COMMAREA, if supplied, becomes its DFHCOMMAREA. The current task's chain continues with the new program; RETURN TRANSID semantics carry through unchanged.

Options

PROGRAM(name) — *required* The target program. Resolved through runtime/transactions.conf.

COMMAREA(data) Bytes to flow to the target.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Control transferred.
PGMIDERR	27	Target program not in transactions.conf.
INVREQ	16	PROGRAM missing.

Example

```
EXEC CICS XCTL PROGRAM('VALIDATE') COMMAREA(SCR) END-EXEC
```

LINK

Synchronously call a sub-program with bidirectional COMMAREA.

Format

```
EXEC CICS LINK PROGRAM(name)
          [COMMAREA(var | 'literal')]
END-EXEC
```

Description

The target program runs in a fresh frame with its DFHCOMMAREA preloaded from the caller's COMMAREA(var) (or literal). When the sub-program exits, its final DFHCOMMAREA is written back to the caller's variable.

Caller state — NextTransid, Commarea, LastResponse, LastMapName — is saved before the LINK and restored on return, so the LINK is transparent to the caller's pseudo-conversational context.

The per-transaction ACL is rechecked on the target so a low-privilege caller cannot escalate by linking into an admin program.

REXX and COBOL programs can LINK to each other freely; DFHCOMMAREA is marshalled as opaque bytes.

Options

PROGRAM(name) — *required* The sub-program to call. Resolved through runtime/transactions.conf.

COMMAREA(var | 'literal') Bidirectional data area. A bare variable is read on entry and written on exit; a literal is read-only.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Sub-program returned cleanly.
PGMIDERR	27	Target not in transactions.conf, sub-program errored, or the caller is not authorised for the target.

Example

```
SAV = ''                                /* working area */
EXEC CICS LINK PROGRAM('CUSV') COMMAREA(SAV) END-EXEC
IF EIBRESP <> 0 THEN SIGNAL ERR
PARSE VAR SAV STATUS '|' MSG
```

ABEND

Abnormally terminate the current task.

Format

```
EXEC CICS ABEND [ABCODE(code)]
                [NODUMP] [CANCEL]
                [RESP(var)] [RESP2(var)]
END-EXEC
```

CANCEL bypasses any installed HANDLE ABEND trap — the task terminates rather than recovering. NODUMP is accepted silently (bricks has no dump facility).

Description

Ends the task with the supplied 4-character abend code (default AAAA). Bricks logs the abend on the operator console and reflects it in the TxCB's status; pseudo-conversational chains are broken (no NextTransid is honoured after an ABEND).

Options

ABCODE(code) 4-character code to record. Defaults to AAAA.

NODUMP Accepted for parity with IBM CICS; no dumps are produced anyway.

Conditions

ABEND does not return; nothing follows it in the program flow.

Example

```
IF EIBRESP <> 0 THEN
  EXEC CICS ABEND ABCODE('FI01') END-EXEC
```

START

Schedule a transaction to fire on this same terminal after a delay (or at a wall-clock time), optionally passing a byte payload the new task picks up via RETRIEVE. This is bricks's implementation of CICS's classic deferred-work pattern: a task queues "do X in 60 seconds" and ends; the terminal returns to the blank prompt; when the timer elapses, bricks dispatches X automatically.

Format

```
EXEC CICS START TRANSID(name)
                [INTERVAL(hhmmss) | TIME(hhmmss) |
                AFTER HOURS(h) MINUTES(m) SECONDS(s) MILLISECS(ms) |
                AT    HOURS(h) MINUTES(m) SECONDS(s) MILLISECS(ms)]
                [FROM(area) [LENGTH(n)]]
                [TERMINID(tttt)]
                [REQID(handle)] [PROTECT]
                [RESP(var)] [RESP2(var)]
END-EXEC
```

REQID(handle) tags the START with a request-id handle. IBM uses this to disambiguate when cancelling or deferring; bricks's single-region scheduler accepts the option without action so program source ports cleanly. PROTECT declares the START as recoverable — bricks already enrolls every START in the task-end SYNCPOINT, so this is silent passthrough. SYSID and QUEUE — explicitly rejected with a short "not supported" message.

Options

TRANSID(name) (*required*) 4-character transaction id, must exist in runtime/transactions.conf. The dispatcher applies the normal per-transaction ACL at fire time, so the operator must still be authorised to invoke name.

INTERVAL(hhmmss) Delay before firing, expressed as up to 6 digits in HHMMSS form (left-padded with zeroes). INTERVAL(30) = 30 s; INTERVAL(000130) = 1 min 30 s; INTERVAL(010000) = 1 h. Omit (or pair with INTERVAL(0)) to fire on the next prompt-loop iteration.

TIME(hhmmss) Absolute wall-clock target, also HHMMSS. Respects the operator's time_zone= setting. If the target is earlier than now, bricks rolls forward to the same time tomorrow (real-CICS semantics). Mutually

exclusive with INTERVAL.

AFTER / AT HOURS(h) MINUTES(m) SECONDS(s) MILLISECS(ms) The IBM-canonical compositional alternatives to INTERVAL / TIME — supply any subset of the four components. AFTER is relative, AT is the next wall-clock instant with that time-of-day, rolling forward to tomorrow when it has already passed (same rule as TIME). The components are additive **counts**, not clock digits, so AFTER MINUTES(90) is an hour and a half; IBM's ranges are HOURS 0–99, MINUTES 0–5999, SECONDS 0–359999 and MILLISECS 0–360000000, with an AFTER total capped at 100 hours and an AT total under 24 hours. MILLISECS(ms) is the bricks sub-second extension and also works on its own (START TRANSID('X') MILLISECS(100)); components coded with neither AFTER nor AT are read as AFTER. Mutually exclusive with INTERVAL / TIME.

FROM(area) [LENGTH(n)] Payload bytes the new task will pull back via RETRIEVE. LENGTH truncates area to n bytes; if omitted, the entire string value of area is queued. The bytes are copied at START time, so the issuing task may mutate area afterwards without disturbing the queued payload.

TERMID(tttt) Must equal the issuing terminal's EIBTRMID in this version. Cross-terminal STARTs return RESP-INVREQ with "only same-terminal STARTs are supported."

Scope and limitations (v1)

- **Same-terminal only.** A START always fires on the issuing TCB. No-TERMID background tasks and TERMID(other) are rejected; both require infrastructure (synthetic TCBs, cross-session authorisation) that's deferred.
- **In-memory queue.** Pending STARTs do **not** survive a bricks restart. If you need recovery-protected scheduling you'll have to wait for START PROTECT (out of scope).
- **No SYSID or QUEUE options.** Each returns RESP-INVREQ with the option name in the error string so a port can be triaged. (REQID is recorded on the queued entry and honoured by CANCEL REQID; PROTECT is accepted as a no-op; HOURS / MINUTES / SECONDS / MILLISECS are implemented — see AFTER / AT above.)
- **One pending payload per task.** When the START fires, the FROM bytes go into tcb.RetrieveBuf. The new task's first RETRIEVE drains it; a second RETRIEVE returns RESP-ENDDATA (29).

Conditions

- RESP-NORMAL on successful enqueue.
- RESP-INVREQ for missing TRANSID, cross-terminal TERMID, unsupported option, or malformed INTERVAL/TIME.

Example

```
EXEC CICS START TRANSID('TIMC') INTERVAL('000030')
              FROM(MSG)
END-EXEC.
```

Schedules TIMC to fire on this terminal in 30 seconds with the contents of MSG queued as the payload. See runtime/cobol/timc.cob and runtime/rexx/timr.rexx for the worked examples.

RETRIEVE

Pull back the FROM(...) payload of the START that scheduled the *current* task. RETRIEVE is the receiving half of the START / RETRIEVE pair; the first thing a transaction fired by START typically does.

Format

```
EXEC CICS RETRIEVE INTO(var) [LENGTH(n)]  
                        [RESP(var)] [RESP2(var)]  
END-EXEC
```

LENGTH(n) is the IBM in/out parameter — caller's value caps the bytes returned, the post-call value is the actual byte count.

Options

INTO(var) (*required*) Variable / DATA area to receive the payload bytes. Must be a name reference, not a string literal.

LENGTH(lenvar) Optional. The variable named here is written with the actual byte length of the payload after a successful read, so the caller can size the consumer code correctly.

Behaviour

- Cold dispatch (operator typed the TRANSID at the prompt) → RESP-ENDDATA (29), nothing written to INTO.
- Scheduled re-entry → RESP-NORMAL, payload bytes copied into INTO, tcb.RetrieveBuf cleared.
- Second RETRIEVE in the same task → RESP-ENDDATA (the buffer is cleared on first read; v1 supports one pending payload per task).

Idiom

The canonical START / RETRIEVE program tests EIBRESP first to discover *which dispatch path* it's on, then branches:

```
MOVE SPACES TO BUF.  
EXEC CICS RETRIEVE INTO(BUF) END-EXEC.  
  
IF EIBRESP = RESP-NORMAL THEN  
    PERFORM SHOW-REMINDER      *> scheduled fire  
    EXEC CICS RETURN END-EXEC  
END-IF.  
  
*> Cold path: prompt the operator, schedule next fire.  
PERFORM PROMPT-AND-SCHEDULE.
```

REXX equivalent in runtime/rexx/timr.rexx; COBOL in runtime/cobol/timc.cob.

Chapter 6. System services

ASSIGN

Read EIB, session, and environment fields into the program.

Format

```
EXEC CICS ASSIGN <FIELD>(target) [<FIELD>(target) ...]  
END-EXEC
```

Description

Reads one or more system fields into program variables. Multiple options can be combined in a single ASSIGN call. Each <FIELD> is one of the keywords below; target is the variable that receives the value.

Options

Option	Returns
USERID(t)	Signed-on userid (empty until CSSN succeeds).
TERMINID(t) / EIBTRMID(t)	Unique 4-digit terminal id (T0001 ...).
EIBAID(t)	Single-byte AID character of the most recent SEND MAP / RECEIVE MAP. Compare with C2X(EIBAID) = 'F3' (REXX) or EIBAID = X'F3' (COBOL) to detect PF3, etc.
EIBCPASN(t)	1-based cursor position from the most recent map response.
EIBCALEN(t)	Length of DFHCOMMAREA flowed in from the caller.
CWALENG(t)	Byte length of the region Common Work Area — the wrkarea= value from bricks.cnf. 0 when no CWA is configured. See EXEC CICS ADDRESS .
TWALENG(t)	Byte length of this task's Transaction Work Area — the TWASIZE 6th field of the transaction's transactions.conf row. 0 when no TWA is configured. See EXEC CICS ADDRESS .
TCTUALENG(t)	Always 0 (bricks does not model a TCTUA).
SCREENHT(t) / SCREENWD(t)	Negotiated terminal rows / columns.
ALTSCRNHT(t) / ALTSCRNWD(t)	Same values; bricks does not distinguish primary and alternate sizes.
CONNECTED(t) / CONNTIME(t)	Wall-clock connect timestamp YYYY-MM-DD HH:MM:SS.
TLS(t)	yes when the session is on the TLS listener, no otherwise.
DATE(t)	Today as YYYYMMDD. (<i>bricks-specific</i>)
TIME(t)	Now as HHMMSS. (<i>bricks-specific</i>)
TODAYYR(t) / TODAYMO(t) / TODAYDY(t)	Today's year / month / day individually. (<i>bricks-specific</i>)
DAYCOUNT(t)	Days since 1900-01-01 (the IBM ABSTIME epoch). Subtract two values for an exact day delta. Behaviour change in round 2: previously returned Unix-epoch days; bricks now matches real CICS.
SYSID(t)	Local CICS region ID (BRKS in bricks).

Option	Returns
APPLID(t)	Region VTAM applid (BRICKS01 in bricks).
NETID(t)	Network ID of the connected terminal (TCPIP in bricks; placeholder for a future config knob).
ABCODE(t)	Most recent abend code (the value EIBABCODE carries). Empty before any ABEND.
EIBTASKN(t)	Numeric task number — the integer body of TxCB.ID.
STARTCODE(t)	How the task was started. T0 (terminal operator) is the bricks default; future work threads S (started by EXEC CICS START) through the scheduler.

The bricks-specific options exist primarily so the COBOL subset can do date math without REXX-style intrinsic functions; see `runtime/cobol/qagc.cob` for an example that computes age in years from a birthdate.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	All requested fields returned.
INVREQ	16	A field name is unknown, or its target is a literal rather than a variable.

Example

```
EXEC CICS ASSIGN USERID(USR)
                TERMID(TRM)
                SCREENHT(SCRH)
                SCREENWD(SCRW)
END-EXEC
```

ASKTIME

Refresh the EIB date/time fields and, optionally, return the current time as a 15-digit absolute timestamp (ABSTIME).

Format

```
EXEC CICS ASKTIME [ABSTIME(target)]
END-EXEC
```

Description

ASKTIME re-reads the system clock and writes the current date and time into EIBDATE and EIBTIME. Both fields follow the IBM CICS formats:

- EIBDATE — packed-style decimal 0CYDDDD, where C is (century – 19) and YYDDD is the two-digit year and Julian day-of-year. For 2026-05-12 (Julian day 132), EIBDATE = 1026132.
- EIBTIME — six-digit HHMMSS.

If ABSTIME(target) is given, target receives a 15-digit decimal string representing milliseconds since 1900-01-01 00:00:00.000. Pass that value to FORMATTIME to break it back into formatted date / time components.

ASKTIME is the only way to get a fresh ABSTIME; the bricks-specific ASSIGN DATE / TIME shortcuts return formatted strings only.

Options

ABSTIME(target) Variable that receives the 15-character absolute time. Optional; when omitted, only EIBDATE and EIBTIME are refreshed.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Always (clock reads cannot fail).

Example

```
EXEC CICS ASKTIME ABSTIME(NOW) END-EXEC
EXEC CICS FORMATTIME ABSTIME(NOW)
                YYYYMMDD(TODAY)
                TIME(CLOCK) TIMESEP(':')

END-EXEC
SAY 'Stamped at' TODAY CLOCK
```

FORMATTIME

Decode an ABSTIME value into formatted date and time fields.

Format

```
EXEC CICS FORMATTIME ABSTIME(source)
                [DATE(t)] [DATEFORM(fmt)] [DATESEP(c)]
                [YYYYMMDD(t)] [MMDDYYYY(t)] [DDMMYYYY(t)]
                [MMDDYY(t)] [DDMMYY(t)]
                [YYYYDDD(t)] [YYDDD(t)]
                [TIME(t)] [TIMESEP(c)]
                [TIMEZONE(name)]
                [YEAR(t)] [MONTHOFYEAR(t)] [DAYOFMONTH(t)]
                [DAYOFWEEK(t)] [DAYCOUNT(t)]
                [RESP(var)] [RESP2(var)]

END-EXEC
```

TIMEZONE(name) is the IBM-canonical per-call zone override — accepts IANA names (America/New_York, Europe/Paris) plus the UTC alias. Unknown zone → INVREQ with a short clean message. Without TIMEZONE, the handler's configured location (from the time_zone line of bricks.cnf) is used.

Description

FORMATTIME is the IBM-standard companion to ASKTIME. Given a 15-digit ABSTIME value (typically returned by ASKTIME ABSTIME(...), but any 15-digit decimal millisecond stamp is accepted), it writes the requested representations into the named target variables.

Any combination of output options may be requested in a single call; options that are omitted cost nothing.

Description

Consults the in-memory user record (loaded from `users.conf`) and the resource's group list (from the dispatch table). Writes one `DFHVALUE` integer into each of the four optional target host variables — the program then compares against the named constants in the `DFHVALUE` copybook.

The verb is stateless and has no side effects: nothing is logged, no lock is taken, no journal entry written. Calling it a thousand times in a tight loop is harmless.

Options

Option	Direction	Notes
RESOURCE (name)	input	TRANSID or program name being checked. Required.
RESCCLASS (class)	input	TCICSTRN (transactions, default) or PCICSPSB (programs). Both class names consult the same transaction-table lookup — bricks has no separate program table, so a TRANSID and the program it dispatches share one ACL row. Anything else → INVREQ QuerySecurityResp2BadResclass (RESP2=2).
READ (t)	output	DFHVALUE-READABLE (50) when the caller may read, DFHVALUE-NOTREADABLE (51) otherwise.
UPDATE (t)	output	DFHVALUE-UPDATABLE (52) / DFHVALUE-NOTUPDATABLE (53).
CONTROL (t)	output	DFHVALUE-CTRLABLE (54) / DFHVALUE-NOTCTRLABLE (55).
ALTER (t)	output	DFHVALUE-ALTERABLE (56) / DFHVALUE-NOTALTERABLE (57).

A literal-string argument on any output axis (`READ('F00')`) is silently skipped — matches the same "literal or empty → no-op" rule the `RESP(var)` / `RESP2(var)` clauses use everywhere.

The verb is a per-user inquiry, so the caller must be signed on. A `PUBLIC` resource is reachable at *dispatch* time without sign-on, but `QUERY SECURITY` answers it like any other resource: an unauthenticated terminal sees all four axes as the `NOT-*` code. After sign-on the verb delegates the `READ` axis to the same `txn.IsAllowed` rule the dispatch gate uses — including `PUBLIC` and `*` semantics — so a signed-on caller against a `PUBLIC` resource sees `READ=DFHVALUE-READABLE`.

Bricks-flavoured access mapping (explicit deviation)

Real CICS / RACF projects per-class permissions onto the four axes via a class-specific matrix. Bricks's existing ACL substrate is binary (you are in the resource's allow-list, or you are not), so the bricks projection narrows the IBM matrix:

Axis	Bricks rule
READ	Caller's userid shares any group with the resource's allow-list, OR the resource has no ACL (open), OR caller is in ADMIN.
UPDATE	Caller is in the ADMIN group.
CONTROL	Caller is in the ADMIN group.
ALTER	Caller is in the ADMIN group.

Document this rule in operator-facing notes alongside any program that gates a menu on `UPDATE` — "only members of `ADMIN` see the `D` key" reads more directly than the abstract `DFHVALUE` matrix.

Conditions

Condition	EIBRESP	EIBRESP2	Cause
NORMAL	0	0	Verb completed; up to four host variables written.
INVREQ	16	QuerySecurityResp2NoResource (1)	RESOURCE(. . .) is missing or resolves to the empty string.
INVREQ	16	QuerySecurityResp2BadResclass (2)	RESCCLASS(. . .) is neither TCICSTRN nor PCICSPSB.
INVREQ	16	QuerySecurityResp2NoChecker (3)	The dispatcher did not wire a SecurityChecker (operator misconfiguration).

Unknown resource → NORMAL with every requested axis set to the NOT-* code. No NOTFND is emitted: the verb is a yes/no inquiry, and "I don't know that resource" collapses cleanly to "all axes denied". Programs that need to distinguish "denied" from "absent" should pair QUERY SECURITY with a separate INQUIRE of the transaction table (not yet implemented; tracked on the deferred list).

NOTAUTH (70) is the IBM-canonical "the caller is not authorised" response code; QUERY SECURITY never emits it (a denied caller sees NORMAL plus the NOT* DFHVALUE on each axis). The VERIFY PASSWORD verb in the following section is the verb that *does* return NOTAUTH for a bad-credentials outcome.

The admin group name is hard-coded to ADMIN in main.go (securityAdminGroup). If a deployment renames it, the UPDATE / CONTROL / ALTER axes return the NOT-* code for every user (the READ axis is unaffected because it consults the per-resource Groups list, not the admin group). A future bricks.cnf admin_group=X knob would lift this hard-coding.

DFHVALUE constants

The runtime/cobolcopy/DFHVALUE.cpy copybook declares each constant twice per the cobol_copybook_dual_names convention — once under the bricks-style short name and once under the DFHVALUE- IBM-traditional alias. Either form compiles:

Bricks mnemonic	IBM alias	Value
READABLE	DFHVALUE-READABLE	50
NOTREADABLE	DFHVALUE-NOTREADABLE	51
UPDATABLE	DFHVALUE-UPDATABLE	52
NOTUPDATABLE	DFHVALUE-NOTUPDATABLE	53
CTRLABLE	DFHVALUE-CTRLABLE	54
NOTCTRLABLE	DFHVALUE-NOTCTRLABLE	55
ALTERABLE	DFHVALUE-ALTERABLE	56
NOTALTERABLE	DFHVALUE-NOTALTERABLE	57

Storage format — bricks decimal-text deviation. DFHVALUE constants ship as PIC S9(8) VALUE (DISPLAY) per the bricks decimal-text host-variable convention — see Chapter 8a for the same shape on RBA. Programs that declare target host vars as PIC S9(8) COMP and then IF WS-CACR-UP = DFHVALUE-UPDATABLE will silently mismatch: the comparison sees a 4-byte binary half-word on one side and the seven ASCII digits "0000052" on the other, and the IF always falls false. Declare the target as plain PIC S9(8) (DISPLAY) so the comparison is bit-equal to

the constant; the `DFHVALUE.cpy` header block carries the long-form rationale and points back here.

COBOL example – gating a Delete menu entry

The shipped `runtime/cobol/gust.cob` sample demonstrates the typical pattern. In its `DISPATCH` paragraph, before the `EVALUATE ACTION` switch, an admin gate sits in front of the `D=Delete` branch:

```

COPY DFHVALUE.
...
01 WS-CACR-UP PIC S9(8) VALUE 0.
...
IF VALID-ACTION = 'Y' AND ACTION = 'D' THEN
    EXEC CICS QUERY SECURITY RESOURCE('GUST')
                                UPDATE(WS-CACR-UP) END-EXEC
    IF WS-CACR-UP NOT = DFHVALUE-UPDATABLE THEN
        MOVE 'Delete not permitted for this user.' TO MSG
        MOVE 'N' TO VALID-ACTION
    END-IF
END-IF.

```

Non-admin callers see the message in MSG on the next menu SEND; admins fall through to the standard ACT-DELETE path.

REXX example

```
EXEC CICS QUERY SECURITY RESOURCE('GUST'),
      READ(WS_RD),
      UPDATE(WS_UP) END-EXEC

IF WS_UP = 52 THEN
  SHOW_DELETE = 'Y'
ELSE
  SHOW_DELETE = 'N'
```

REXX programs may compare against the bare integer (52 above) or import the same numeric values from a constants stem if the shop prefers symbolic names.

Password verification — VERIFY PASSWORD

Companion verb to QUERY SECURITY. Asks "is this plaintext password the right one for this userid?" against the same users.conf-backed authenticator CSSN uses for sign-on. Used by batch / SOAP entry points that need to re-authenticate before a privileged operation, and by terminal screens that want a "type your password again" gate in front of a destructive action — even though the session is already signed on.

Format

[illegible]

Description

Re-authenticates the (userid, password) pair against the running `users.conf` and returns `NORMAL (0)` on success, `NOTAUTH (70)` on any failure. **The password is compared exactly** — byte for byte, case included. (The verifier used to try the as-typed, upper-cased and lower-cased forms against the single stored hash, so up to three different inputs satisfied this verb for one account; `VERIFY PASSWORD` inherited that from the shared authenticator. Exactly one spelling is accepted now. Nothing about the stored hashes changed — each was always built from the exact enrolled password.) The userid itself is still matched case-insensitively. The verb itself is stateless — no journal entry, no lock — and reads the same in-memory user record `cssn` consults at sign-on. Calling it repeatedly is harmless aside from the `bcrypt` cost on each call (the `FileStore.Authenticate` path spends the same compute on unknown users as on known ones to defeat username-enumeration via response timing).

Options

Option	Direction	Notes
<code>USERID(uid)</code>	input	Userid to verify. Required. Trimmed of surrounding whitespace; empty → <code>INVREQ RESP2=1</code> .
<code>PASSWORD(pwd)</code>	input	Plaintext password. Required. Passed verbatim to the verifier (no trim — leading/trailing whitespace is part of the credential); empty or missing <code>PASSWORD</code> passes through to the verifier and folds onto the standard <code>NOTAUTH (70)</code> bad-credential branch. (A fast-reject on empty <code>PASSWORD</code> was withdrawn in review pass 3 — see the Conditions table below for the rationale.)
<code>CHANGETIME(t)</code>	output	<i>Parsed but not populated</i> — see deviation below.
<code>EXPIRYTIME(t)</code>	output	<i>Parsed but not populated</i> .
<code>DAYSLEFT(d)</code>	output	<i>Parsed but not populated</i> .
<code>INVALIDCOUNT(n)</code>	output	<i>Parsed but not populated</i> .
<code>LASTUSETIME(t)</code>	output	<i>Parsed but not populated</i> .

Bricks deviation — non-distinguishing failure mode (explicit)

Loud deviation from IBM RACF. Real CICS distinguishes "unknown userid" from "bad password" via separate `RESP2` sub-codes. **Bricks intentionally does not.** Every authentication failure — unknown userid, wrong password, future account lockout, rate-limit refusal — collapses onto the single `NOTAUTH (70)` result with `RESP2=0`. The motivation is straightforward: this verb is safe to expose behind public-facing entry points (a WAPI re-auth endpoint, a SOAP credential check), and any `RESP2`-level distinction would let a hostile caller enumerate valid userids by inspecting the sub-code. The bricks `FileStore.Authenticate` path already pays the `bcrypt` cost on unknown users for the same timing-resistance reason, so the bricks path is uniform end-to-end. The contract is enforced in code by the `cics.ErrPasswordVerifyBadCredentials` sentinel: every `PasswordVerifier` implementation must collapse its failure modes onto this single sentinel before returning.

Bricks deviation — password-policy outputs not populated

The optional outputs `CHANGETIME` / `EXPIRYTIME` / `DAYSLEFT` / `INVALIDCOUNT` / `LASTUSETIME` parse cleanly (so source ported from a real z/OS shop continues to compile) but bricks does not populate them. Bricks's `users.conf` records `bcrypt` hash + group list and nothing else — there is no password-policy metadata to

write back. If the named target is a host variable, it is left unchanged after the verb returns; if it is a literal, the option is silently dropped.

A program that relies on `DAYSLEFT` or `EXPIRYTIME` to drive an "expires in N days" warning must gate that branch on whether the variable is still at its pre-call value. The shipped `runtime/rexx/cssr.rexx` sample (below) does not invoke the optional outputs.

Conditions

Condition	EIBRESP	EIBRESP2	Cause
NORMAL	0	0	Credentials matched.
NOTAUTH	70	0	Bad credentials — wrong password OR unknown userid OR empty/missing PASSWORD OR any future failure mode. The sub-code is intentionally 0 to defeat enumeration (see deviation above).
INVREQ	16	VerifyPasswordResp2NoUserid (1)	USERID(. . .) missing or resolves to the empty string.
INVREQ	16	VerifyPasswordResp2NoVerifier (3)	The dispatcher did not wire a PasswordVerifier (operator misconfiguration).
IOERR	17	0	The verifier reported a transient error (file-system fault, future store unreachable, ...). The shipped bricks passwordAdapter folds every observed FileStore error onto NOTAUTH today, but custom verifiers may return arbitrary errors that surface as IOERR — a deployer who plugs in an LDAP or RACF-passthrough verifier should expect this path to fire.

No RESP2=2 row. Sub-code 2 was withdrawn in review pass 3: it formerly meant "PASSWORD missing or empty" and was a fast-reject before the verifier ran. The wall-clock gap between "PASSWORD absent — fast INVREQ" and "PASSWORD wrong — ran-bcrypt-then-NOTAUTH" let an attacker time the verb and learn whether the verifier was reached. Empty / missing PASSWORD now falls through to the verifier and lands on the standard NOTAUTH branch. The numeric gap at 2 is preserved deliberately so an operator reading a RESP2=3 log line from a deployed program still finds the same "no verifier configured" cause.

Refused in CECI. The CECI debugging shell refuses `EXEC CICS VERIFY PASSWORD` outright because the operator's typed PASSWORD value would echo on the input pane (and could land in the CECI audit log if `ceci_audit` is ever re-enabled). Programs that need a "type your password again" gate must run as their own TRANSID (see `runtime/rexx/cssr.rexx` for the canonical sample).

REXX example — re-authenticate before a privileged action

The shipped `runtime/rexx/cssr.rexx` sample demonstrates the canonical "type your password again" gate (CSSR = "Sign-on Re-verify"; TRANSID registered in `runtime/transactions.conf` with `USERS,ADMIN` groups). Usage: type CSSR at the bricks blank prompt; the CSSRPW map then prompts for the password in a HIDDEN (non-echoing) writable field. PF3 cancels without verifying.

```

ADDRESS CICS
EXEC CICS ASSIGN USERID(USR) END-EXEC
MAP.USERID    = USR          /* preload the signed-on user */
MAP.PASSWORD  = ''          /* clear any leftover value  */
EXEC CICS SEND    MAP('CSSRPW') FROM(MAP) ERASE END-EXEC
EXEC CICS RECEIVE MAP('CSSRPW') INTO(MAP) END-EXEC
IF EIBAID = '3' THEN DO      /* PF3 = cancel          */
    EXEC CICS SEND TEXT FROM('CSSR cancelled.') ERASE END-EXEC
    EXEC CICS RETURN END-EXEC
END
PWD = STRIP(MAP.PASSWORD)
EXEC CICS VERIFY PASSWORD USERID(USR) PASSWORD(PWD) RESP(RC) END-EXEC
IF RC = 0 THEN
    MSG = 'Password OK. Privileged step authorised.'
ELSE IF RC = 70 THEN
    MSG = 'Bad credentials. Privileged step refused.'
ELSE
    MSG = 'VERIFY PASSWORD failed RESP=' || RC || ' RESP2=' || EIBRESP2
EXEC CICS SEND TEXT FROM(MSG) ERASE END-EXEC
EXEC CICS RETURN END-EXEC

```

The RESP(RC) clause writes the result into a host variable inline so the IF RC = 70 branch reads cleanly instead of fishing through EIBRESP after the fact.

Why the map, not a command-tail. A prior draft of CSSR read the password from the initial-input buffer (CSSR <password>). That pattern echoes the password on the 3270 input pane, drops it into any CONS shell view, and lands it in CECI's input echo if an operator ever reissues the command from there. The shipped sample uses a CSSRPW map whose PASSWORD writable field carries the DARK (HIDDEN) attribute, so the operator's keystrokes never render on screen. Production re-auth flows must follow the same pattern — see the [3270 EBCDIC 037 memory](#) notes on writable-field attributes if you build your own map.

COBOL idiom — gating LINK with a typed-again password

The same pattern in COBOL. The WS-PWD field is the operator's re-typed password (collected on a RECEIVE MAP from a sensitive- action confirmation screen, not shown). On NOTAUTH the program falls through to a refuse-message path; on NORMAL it can LINK to the privileged sub-program.

```

77 WS-VP-RC PIC S9(8) VALUE 0.
77 WS-USR   PIC X(8).
77 WS-PWD   PIC X(32).
...
EXEC CICS ASSIGN USERID(WS-USR) END-EXEC
EXEC CICS VERIFY PASSWORD USERID(WS-USR)
                        PASSWORD(WS-PWD)
                        RESP(WS-VP-RC) END-EXEC

EVALUATE WS-VP-RC
    WHEN 0
        EXEC CICS LINK PROGRAM('ADMNDEL') END-EXEC
    WHEN 70
        MOVE 'Password does not match. Try again.' TO MSG
    WHEN OTHER
        MOVE 'VERIFY PASSWORD failed (config?).' TO MSG
END-EVALUATE.

```

The EXEC CICS ASSIGN USERID(WS-USR) step is the canonical way to read the signed-on userid into a host variable — bricks does **not** populate an EIBUSER field, so a port from a z/OS shop that used

USERID(EIBUSER) directly must be rewritten to go through ASSIGN first. The shipped runtime/rexx/cssr.rexx sample uses the REXX flavour of the same pattern.

A COBOL caller that intends to EVALUATE WS-VP-RC WHEN 70 must declare WS-VP-RC as plain PIC S9(8) (DISPLAY) to match the bricks decimal-text host-variable convention — same rule as the DFHVALUE constants documented under QUERY SECURITY above.

System inquiry — INQUIRE SYSTEM

Recently added in 2.7.x. INQUIRE SYSTEM is the entry point for SIT-level inquiries — values the operator set in bricks.cnf that a program needs at run time. Bricks v1 honours **only** the GMMTEXT option (the IBM CICS canonical "Good Morning" sign-on text); the dispatch infrastructure is in place so future SIT options (RELEASE, CICSSTATUS, MAXTASKS, JOBNAME, APPLID, SYSID) extend the same verb without touching the Handler / Dispatcher field set.

Format

```
EXEC CICS INQUIRE SYSTEM
      [GMMTEXT(var)]
      [RESP(rc)] [RESP2(rc2)]
END-EXEC
```

Description

Reads SIT-level configuration into the program. v1 honours only the GMMTEXT(var) option, which writes the configured gmtxt= value from bricks.cnf (default "Welcome to bricks") into the named host variable. A bare INQUIRE SYSTEM (no options, or only RESP/RESP2) is a no-op success matching real CICS.

Options

Option	Bricks v1 behaviour
GMMTEXT(var)	Writes the configured GMTEXT into var. Full 256-byte value is delivered verbatim — no truncation at verb time. (The bricks LogonPrompt renderer truncates to cols-1 at paint time, but the verb's host-variable surface preserves the full storage.)
Other (RELEASE, CICSSTATUS, MAXTASKS, JOBNAME, APPLID, SYSID, etc.)	Not yet supported in v1 — surfaces INVREQ with EIBRESP2 = 1.

Conditions

Condition	EIBRESP	EIBRESP2	Cause
NORMAL	0	0	Verb succeeded. Bare INQUIRE SYSTEM (no options) also returns NORMAL even when the dispatcher hasn't wired a SystemInfoProvider.
INVREQ	16	1	The operator named an option v1 doesn't honour yet. The diagnostic names the option.
INVREQ	16	2	h.SystemInfo == nil AND a substantive option was named — the dispatcher did not wire a provider. Mirrors the QUERY SECURITY / VERIFY PASSWORD RESP=3 loud-fail convention.

Reserved RESP2 sub-codes 3..9 are kept free for future INQUIRE form sub-errors (INQUIRE FILE, INQUIRE TASK, etc.). v1 programs that pin numeric RESP2 codes will continue to see only 1 or 2 for INQUIRE SYSTEM; future bricks releases will not silently re-bind these reservations.

Bricks deviations (loud)

Per the `bricks_cics_ibm_canonical` memory every deviation from IBM canonical INQUIRE SYSTEM is called out individually below.

Bricks deviation — GMMTEXT is the only honoured option in v1.

Of the 30+ IBM INQUIRE SYSTEM options (RELEASE, CICSSTATUS, MAXTASKS, JOBNAME, APPLID, SYSID, ...) bricks v1 honours only GMMTEXT. Any other named option surfaces INVREQ RESP2=1 with the option name in the diagnostic, so a port from a z/OS shop fails loudly rather than silently returning zero / empty. Adding any deferred option is a ~5-line extension to the `cics.SystemInfoProvider` interface (one new method) plus a single case in `cics/inquire.go`'s `inquireSystem` handler — no Handler / Dispatcher / Frame changes.

Bricks deviation — EBCDIC-037 fail-fast at startup.

The `gmtext=` value is validated against the EBCDIC-037 printable range (0x20..0x7E) minus the 3270_ebcdic_037_printables deny-set ([] { } ~ \ ` | ^) at config load. A `gmtext=` containing any non-printable byte (e.g. an embedded bell `\x07`, a control character, or a non-ASCII glyph) or any deny-set glyph is a startup error. Real CICS would silently substitute on the wire; bricks refuses to start so the operator notices.

Bricks deviation — empty gmtext= restores the default.

A bare `gmtext=` line (or one with only whitespace inside the quotes) DOES NOT blank the banner — it restores the "Welcome to bricks" default. Mirrors the `ntp_server=on` aliasing precedent and matches the bricks last-wins-per-key configuration policy. An operator who genuinely wants no banner cannot achieve that via `gmtext=`; the only way to suppress the sign-on banner entirely is to leave both `gmtext=` and the legacy welcome line out of the LogonPrompt path (currently not possible without a code change).

Bricks deviation — GMTEXT replaces the legacy welcomeLine on every welcome screen.

GMTEXT is painted at row 0 of the connect-time `ShowLogoSplash` AND at row 1 of the LogonPrompt (when `enforce_secure_login=yes`) AND on the CSSF LOGOFF confirmation splash. The legacy `Welcome to BRICKS HH:MM:SS` banner with timestamp only paints when `gmtext=` is empty or absent. `BlankPrompt` (the steady-state TRANSID cursor) is NOT a banner site — adding GMTEXT there would repaint on every return from a transaction, which surprises operators.

Real CICS paints both the GMTEXT and a "good morning" line on the sign-on prompt. Bricks stacks neither — the operator sees one centred banner instead of two, which the architect's C7 review flagged as the correct layout. When `gmtext=` is empty / whitespace (which the config load layer should have already restored to default, but the renderer composes safely either way) the legacy `welcomeLine` paints as fallback.

Sample REXX

The shipped `runtime/rexx/gmtq.rexx` (TRANSID GMTQ) demonstrates the recommended idiom — `RESP(RC)` for clean error branching, no comma-continuations inside the `EXEC` body (per the `bricks_rexx_exec_one_line` memory):

```
/* GMTQ -- read GMTEXT and echo it to the terminal. */
ADDRESS CICS

MSG = ''
EXEC CICS INQUIRE SYSTEM GMMTEXT(MSG) RESP(RC) END-EXEC

IF RC \= 0 THEN
  MSG = 'INQUIRE SYSTEM GMMTEXT failed; EIBRESP=' || RC || ' EIBRESP2=' || EIBRESP2

EXEC CICS SEND TEXT FROM(MSG) ERASE END-EXEC
EXEC CICS RETURN END-EXEC
```

A program that omits `RESP(RC)` still gets the Go-level error in `EIBRESP` via the existing `captureRespVars` pipeline (a `RESP` clause that names a host variable is the recommended shape for clean per-call error branching; the `EIBRESP` post-check shape also works).

SIT-knob — `gmtext=` in `bricks.cnf`

```
# bricks.cnf
gmtext="Sample bank -- production" # max 256 bytes, EBCDIC-037 printable only
```

Defaults to "Welcome to bricks" when the line is absent or empty. See [Configuration — bricks.cnf](#) in the README for the full knob reference.

ADDRESS — EXEC CICS ADDRESS

Storage addressing. `ADDRESS` hands a program a pointer to a CICS-owned storage area — the `COMMAREA`, the `EIB`, the region `CWA`, or the per-task `TWA`. It is the front half of the bricks COBOL pointer model: the handle `ADDRESS` writes into a `USAGE POINTER` item is later mapped onto a based `LINKAGE 01` with `SET ADDRESS OF`. See [Chapter 22 — pointers](#) for the COBOL surface.

Format

```
EXEC CICS ADDRESS
  [COMMAREA(ptr)] [EIB(ptr)] [CWA(ptr)] [TWA(ptr)]
  [TCTUA(ptr)] [ACEE(ptr)]
  [RESP(rc)] [RESP2(rc2)]
END-EXEC
```

Description

Each named option's `ptr` operand (a `USAGE POINTER` item in COBOL) is set to an opaque **handle** addressing the requested area. A COBOL program then maps a based `LINKAGE 01` onto those bytes with `SET ADDRESS OF L TO ptr`. The handles are not displayable machine addresses, but `SET ptr UP/DOWN BY n` arithmetic *is*

supported — it adjusts the offset **within the pointer's current target buffer** (see the [COBOL pointer model](#)).

In **REXX** ADDRESS is accept-and-no-op: every operand is set to 0 (NULL) and NORMAL is returned. REXX has no pointers — the data a REXX program needs (COMMAREA, EIB fields) is already injected as variables — so the verb exists only for source compatibility.

Options

Option	Handle addresses	Notes
COMMAREA(ptr)	DFHCOMMAREA	Always a non-NULL handle (the based DFHCOMMAREA default buffer always exists). EIBCALEN tells you whether a caller actually flowed bytes in.
EIB(ptr)	DFHEIB	Always non-NULL. Maps the real 01 DFHEIB block (children EIBRESP, EIBAID, EIBCALEN, ...).
CWA(ptr)	region Common Work Area	NULL (handle 0) when wrkarea=0 in bricks.cnf. One slice shared by every task; writes are visible region-wide (unsynchronised). Length via ASSIGN CWALENG.
TWA(ptr)	per-task Transaction Work Area	NULL when the transaction's TWASIZE (transactions.conf 6th field) is 0. Fresh per task. Length via ASSIGN TWALENG.
TCTUA(ptr)	terminal user area	Always NULL — not modelled in bricks.
ACEE(ptr)	security ACEE	Always NULL — not modelled in bricks.

CSA and TCTTE are **not** accepted (both obsolete as addressable areas in modern CICS); naming them — or any other option — returns INVREQ.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	The named areas were addressed. NULL is still NORMAL — a NULL handle for an unconfigured CWA / TWA (or for TCTUA / ACEE) is the documented outcome, not an error.
INVREQ	16	An option other than the accepted set was named (e.g. CSA, TCTTE). The diagnostic names the option.

Bricks deviations (loud)

Bricks deviation — TCTUA and ACEE are always NULL. Neither the per-terminal user area nor the security ACEE is modelled. ADDRESS TCTUA(p) / ADDRESS ACEE(p) return NORMAL with p set to 0; a program that dereferences either gets a clean NULL-pointer abend, not silent garbage.

Bricks deviation — ADDRESS CWA is NULL unless wrkarea>0. A region with no configured CWA (the default wrkarea=0) yields a NULL CWA handle — exactly as real CICS does on a region with no CWA. Set wrkarea= in bricks.cnf to allocate one.

The CWA is unsynchronized — serialize updates with ENQ/DEQ. The region CWA is a single byte slice shared by reference into every concurrent task; bricks does **not** auto-serialize access (neither does real CICS). Concurrent unsynchronized writes from different tasks are undefined — guard CWA updates with EXEC CICS ENQ / DEQ, as on real CICS.

Example (COBOL)

```
01 WS-CPTR  USAGE POINTER.
...
EXEC CICS ADDRESS COMMAREA(WS-CPTR) END-EXEC.
IF EIBCALEN = 0 THEN
    DISPLAY 'no commarea flowed in'
END-IF.
SET ADDRESS OF DFHCOMMAREA TO WS-CPTR.
```

The shipped runtime/cobol/gmpt.cob (TRANSID GMPT) exercises ADDRESS COMMAREA and ADDRESS EIB alongside GETMAIN.

GETMAIN / FREEMAIN — EXEC CICS GETMAIN, EXEC CICS FREEMAIN

Dynamic task storage. GETMAIN allocates a buffer from the task heap and writes a pointer to it into a USAGE POINTER item; FREEMAIN releases it. Together with SET ADDRESS OF they let a COBOL program build and address records whose size or count isn't known at compile time.

Format

```
EXEC CICS GETMAIN SET(ptr)
    {FLENGTH(n) | LENGTH(n)}
    [INITIMG(byte)] [SHARED] [NOSUSPEND]
    [USERDATAKEY | CICSDATAKEY]
    [RESP(rc)] [RESP2(rc2)]
END-EXEC

EXEC CICS FREEMAIN {DATA(area) | DATAPOINTER(ptr)}
    [RESP(rc)] [RESP2(rc2)]
END-EXEC
```

Description

GETMAIN allocates *n* bytes (FLENGTH preferred, LENGTH accepted as a synonym), fills every byte with INITIMG (first byte of the operand; default 0x00), and writes the buffer's handle into the SET pointer. $n \leq 0$ is LENGERR. A program then maps a based LINKAGE 01 onto the buffer with SET ADDRESS OF L TO ptr and reads / writes through the based item.

FREEMAIN resolves its operand to a GETMAIN handle and releases the buffer. DATAPOINTER(ptr) names a pointer variable holding the handle directly; DATA(area) names the based data area whose current address is the handle. Freeing an address that was **not** obtained by GETMAIN (e.g. ADDRESS OF a declared item, or a stale / zero handle) is INVREQ.

In **REXX** both verbs are accept-and-no-op: GETMAIN still validates the length (LENGERR on ≤ 0) and writes a non-zero dummy handle so control flow is preserved, but no real bytes back it; FREEMAIN returns NORMAL with nothing to free.

Options

Option	Verb	Effect
SET(ptr)	GETMAIN	Required. Receives the buffer handle.

Option	Verb	Effect
FLENGTH(n) / LENGTH(n)	GETMAIN	Allocation size in bytes; FLENGTH wins if both appear.
INITIMG(byte)	GETMAIN	Fill byte (first byte of the operand). Default 0x00.
SHARED	GETMAIN	Accepted for source compatibility — see the loud deviation below; NOT region-scoped in bricks.
NOSUSPEND	GETMAIN	Accepted no-op (bricks never suspends a GETMAIN).
USERDATAKEY / CICSSTACKKEY	GETMAIN	Accepted no-ops (bricks has no storage-key model).
DATA(area) / DATAPOINTER(ptr)	FREEMAIN	Identifies the buffer to release.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	GETMAIN allocated and wrote the handle; FREEMAIN released the buffer.
LENGERR	22	GETMAIN length ≤ 0 .
NOSTG	42	GETMAIN length exceeds the 16 MiB region storage limit — no buffer is allocated.
INVREQ	16	GETMAIN with no SET; or FREEMAIN of an address not obtained by GETMAIN (including a zero / NULL handle).

Bricks deviations (loud)

Bricks deviation — SHARED GETMAIN is NOT region-scoped. Real CICS SHARED storage survives the allocating task and is reachable region-wide. Bricks frees **every** task GETMAIN buffer (SHARED or not) when the task's storage is dropped at task end — there is no region-scoped storage pool yet. The SHARED flag is accepted so ports compile, but a program relying on cross-task SHARED storage will not see it persist. **Use a TS queue or the CWA (ADDRESS CWA, wrkarea=) for genuinely cross-task data.**

Bricks note — pointer arithmetic is per-buffer. GETMAIN handles are opaque interned (buffer, offset) pairs. SET ptr UP BY n / DOWN BY n adjusts the offset **within the same buffer** the pointer already addresses; it does not walk across separate allocations. A past-end offset is allowed until dereferenced (then a cleanabend); below 0 or on a NULL pointer the SET itself diagnoses cleanly. See the [COBOL pointer model](#).

Example (COBOL)

```

01 WS-PTR  USAGE POINTER.
LINKAGE SECTION.
01 GM-AREA BASED.
   05 GM-TEXT PIC X(16).
PROCEDURE DIVISION.
   EXEC CICS GETMAIN SET(WS-PTR) FLENGTH(16) INITIMG(' ') END-EXEC.
   SET ADDRESS OF GM-AREA TO WS-PTR.
   MOVE 'POINTER-OK' TO GM-TEXT.           *> writes into the GETMAIN buffer
   EXEC CICS FREEMAIN DATAPOINTER(WS-PTR) END-EXEC.
```

The shipped runtime/cobol/gmpt.cob (TRANSID GMPT) runs exactly this round-trip and paints the read-back via SEND TEXT.

INQUIRE / SET resources — FILE, TASK, TRANSACTION, TERMINAL, PROGRAM

Resource inquiry. Beyond INQUIRE SYSTEM (SIT values), bricks honours INQUIRE / SET for five resource classes, in both the direct and browse access shapes. Status attributes return the IBM happy-path defaults; sizes and counts bricks actually tracks are real.

Format

```
Direct: EXEC CICS INQUIRE <class>(name) <attr(var)>... [RESP(rc)] END-EXEC

Browse: EXEC CICS STARTBROWSE <class> [RESP(rc)] END-EXEC
        EXEC CICS INQUIRE <class>(area) NEXT <attr(var)>... RESP(rc) END-EXEC    (repeat)
        EXEC CICS ENDBROWSE <class> END-EXEC

Set:    EXEC CICS SET <class>(name) <attr(value)>... [RESP(rc)] END-EXEC
```

<class> is one of FILE, TASK, TRANSACTION, TERMINAL, PROGRAM.

Description

Direct reads one named resource's attributes into host variables. **Browse** freezes the resource-name list at STARTBROWSE and walks it one resource per INQUIRE ... NEXT: the class option (INQUIRE FILE(area) NEXT) receives the current resource's name and the attribute options receive its values; the cursor advances each call and returns END at exhaustion. One browse per class per task. **SET** validates the named resource exists and returns NORMAL — see the loud deviation: it does **not** persist a state change.

Attribute accept-lists per class

Class	INQUIRE attributes	Real vs. default
FILE(name)	ENABLESTATUS OPENSTATUS RECORDSIZE KEYLENGTH KEYPOSITION RECORDFORMAT ACCESSMETHOD BASEDSNAME	RECORDSIZE/KEYLENGTH/BASEDSNAME are real (store catalog). ENABLESTATUS=ENABLED, OPENSTATUS=OPEN, KEYPOSITION=0, RECORDFORMAT=VARIABLE, ACCESSMETHOD=VSAM are defaults.
TASK(id)	TRANSACTION FACILITY USERID RUNSTATUS PRIORITY	TRANSACTION/FACILITY/USERID are real . RUNSTATUS=RUNNING, PRIORITY=1 are defaults. id is the numeric EIBTASKN form.
TRANSACTION(name)	STATUS PROGRAM PRIORITY TWSIZE PROFILE	PROGRAM/TWSIZE are real . STATUS=ENABLED, PRIORITY=1, PROFILE=DFHCICST are defaults.
TERMINAL(id)	ACQSTATUS SERVSTATUS USERID NETNAME	USERID/NETNAME are real . ACQSTATUS=ACQUIRED, SERVSTATUS=INSERVICE are defaults.
PROGRAM(name)	STATUS LANGUAGE PROGTYPE RESCOUNT USECOUNT LENGTH	LANGUAGE (COBOL/REXX) and USECOUNT are real . STATUS=ENABLED, PROGTYPE=PROGRAM, RESCOUNT=0, LENGTH=0 are defaults.

SET <class>(name) accepts the same class option and the documented mutation attributes (e.g. SET FILE('x') OPENSTATUS(CLOSED), SET TRANSACTION('x') STATUS(DISABLED)).

Conditions

Condition	EIBRESP	Where
NORMAL	0	Direct INQUIRE / SET succeeded; INQUIRE ... NEXT returned a resource; STARTBROWSE / ENDBROWSE succeeded.
END	83	INQUIRE <class> NEXT after the last resource — the browse is exhausted.
NOTFND	13	Direct INQUIRE / SET named a resource that doesn't exist.
INVREQ	16	Unknown resource class; an unknown attribute keyword; a direct INQUIRE with no resource name and no NEXT; ENDBROWSE with no open browse.
ILLOGIC	21	INQUIRE <class> NEXT with no STARTBROWSE active for that class.

Bricks deviations (loud)

Bricks deviation — status attributes are IBM defaults, not live state. Bricks has no persistent resource-state model. Every status field (OPENSTATUS, ENABLESTATUS, RUNSTATUS, ACQSTATUS, SERVSTATUS, transaction / program STATUS) returns the IBM happy-path default. A program must not rely on INQUIRE to discover that a file is CLOSED or a transaction DISABLED — that state is never modelled. The attributes bricks genuinely tracks (RECORDSIZE, KEYLENGTH, LANGUAGE, USECOUNT, TWASIZE, PROGRAM, USERID, TRANSACTION, ...) are real.

Bricks deviation — SET is acknowledge-only. SET FILE(...) OPENSTATUS(CLOSED), SET TRANSACTION(...) STATUS(DISABLED), etc. validate that the named resource exists (NOTFND if not) and return NORMAL, but **do not persist** the change — the very next INQUIRE would otherwise contradict the SET. This is deliberate: bricks fails honestly rather than faking a state transition it can't keep.

Example (REXX) — direct + browse

```
ADDRESS CICS

/* Direct: this task's own attributes. */
EXEC CICS ASSIGN EIBTASKN(MYTASK) END-EXEC
EXEC CICS INQUIRE TASK(MYTASK) RUNSTATUS(RS) TRANSACTION(TR) RESP(RC) END-EXEC

/* Browse: walk every FILE. RC goes non-zero (END) at exhaustion. */
EXEC CICS STARTBROWSE FILE RESP(RC) END-EXEC
DO WHILE RC = 0
  FN = ''
  EXEC CICS INQUIRE FILE(FN) NEXT OPENSTATUS(OS) RECORDSIZE(RZ) RESP(RC) END-EXEC
  IF RC = 0 THEN SAY FN OS RZ
END
EXEC CICS ENDBROWSE FILE END-EXEC
```

The shipped runtime/rexx/inqr.rexx (TRANSID INQR) runs this direct-plus-browse pattern and paints the result with SEND TEXT.

Part 5. EXEC CICS file and queue commands

Part 5 begins here. This part covers the persistent and ephemeral data services — VSAM-style KSDS files and the temporary- storage / transient-data queue families. Both surfaces are per-task and bound by the same SYNCPOINT mechanism described in [Chapter 10](#).

Chapter 7. KSDS file commands

VSAM in Bricks. Bricks's KSDS file surface is the application- programmer's view of the persistent record store. Each FILE name in EXEC CICS READ FILE(name) corresponds to a B+tree bucket in the embedded data/files.boltddb database. A separate DEFINE CLUSTER step is **optional**: the first EXEC CICS WRITE creates the bucket if it isn't already there, and subsequent READ / REWRITE / DELETE / STARTBR / READNEXT / READPREV / RESETBR / ENDBR see it. Records are opaque byte payloads; the application chooses the layout (typically a pipe-delimited group of fields, the convention every sample uses).

Operators who want the classic mainframe lifecycle — declare a cluster with DEFINE CLUSTER, bulk-load it from a sequential file with REPRO, drop it with DELETE, enumerate the catalogue with LISTCAT, dump records with PRINT — should use the **IDCA** transaction described under *VSAM file management* below. Files created through IDCAMS and files auto-created by EXEC CICS WRITE are byte-identical: same bucket, same _catalog row, same CEMT INQUIRE FILE view.

These commands operate on a **single record**, identified by a key, in a CICS FILE. Each FILE is a bbolt bucket inside data/files.boltddb; record bodies are opaque bytes (the application chooses the layout). See *How file storage works* in the README for the on-disk model.

READ

Read a record from a KSDS by key.

Format

```
EXEC CICS READ FILE(name)
           {INTO(target) | SET(target)}
           RIDFLD(key)
           [UPDATE]
           [LENGTH(len)]
END-EXEC
```

Description

B+tree exact-key lookup, $O(\log n)$. The record bytes (whatever the application stored) come back into the target. UPDATE records a per-session lock on the key, gating a subsequent REWRITE on the same FILE.

When LENGTH(var) is a bare variable, the actual record length is written back.

Options

FILE(name) — *required* The CICS FILE name.

INTO(target) / SET(target) — *one required* Destination for the record body. INTO and SET are accepted interchangeably here.

RIDFLD(key) — *required* The key to look up.

UPDATE Record a per-session lock so a subsequent REWRITE can update this record.

LENGTH(len) Receives the record's actual length when the option is a bare variable.

Bricks deviation from IBM canonical

LENGTH(var) on READ FILE is OUTPUT-only, not input-as-max- buffer. Real CICS VSAM variable-length reads LENGTH(var) as a cap on the bytes returned and writes back the actual size; bricks ignores the input value and only writes back. The write-back-then- read-as-input shape would otherwise corrupt subsequent reads in a loop (LENGTH(LEN) on call 1 stores 4; call 2 silently caps a 13-byte record to 4). The bricks KSDS store is fixed-size-per- record so there is no operator-meaningful truncation here.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Record returned.
NOTFND	13	Key not present.
INVREQ	16	Missing required option, or invalid file name.
IOERR	17	Underlying store error.

Example

```
EXEC CICS READ FILE('CUSTOMERS')
          INTO(REC) RIDFLD(CKEY)
END-EXEC
IF EIBRESP = 13 THEN MSG = 'Customer not found'
```

WRITE

Insert a new record into a KSDS.

Format

```
EXEC CICS WRITE FILE(name)
          FROM(data) [LENGTH(n)]
          RIDFLD(key)
          [MASSINSERT]
          [RESP(var)] [RESP2(var)]
END-EXEC
```

LENGTH(n) is the IBM in/out parameter — caller's value caps the bytes written, returned value is the actual bytes stored. MASSINSERT is accepted for source compatibility; bricks's KSDS engine is uniformly fast so the flag is a silent hint.

Description

B+tree insert. The bucket for the FILE is created implicitly on first WRITE — there is no EXEC CICS DEFINE FILE step. DUPREC if a record with the same key already exists.

The write commits in a single bbolt update transaction with fsync on success.

Options

FILE(name) — *required* Target FILE. Created if absent.

FROM(data) — *required* The record bytes to store.

RIDFLD(key) — *required* The key under which to store the record.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Record written.
DUPREC	14	A record with this key already exists.
INVREQ	16	Missing required option, or invalid name.
IOERR	17	Underlying store error.

Example

```
REC = NM || '|' || AD || '|' || CY || '|' || PH
EXEC CICS WRITE FILE('CUSTOMERS') FROM(REC) RIDFLD(CKEY) END-EXEC
```

REWRITE

Replace the record locked by the most recent READ ... UPDATE.

Format

```
EXEC CICS REWRITE FILE(name)
                FROM(data)
END-EXEC
```

Description

Overwrites the value at the key locked by the most recent READ FILE ... UPDATE on the same FILE. Releases the per-FCB update lock at end of transaction.

Bricks deviation from IBM canonical: READ UPDATE + REWRITE do **not** survive the CECI PF5 boundary. Real conversational CICS programs hold the UPDATE lock across the operator's pause; CECI in bricks treats each PF5 press as an implicit SYNCPOINT, which releases non-HOLD locks. Recommended patterns: type both verbs on one PF5 input (assembled as a single command), or wrap the pair with ENQ resource HOLD ... DEQ resource so the lock persists across the press. Outside CECI (REXX or COBOL programs running as a normal task) the lock survives until the next SYNCPOINT or task end, matching real CICS.

Options

FILE(name) — *required* **FROM(data)** — *required*

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Record updated.
INVREQ	16	No prior READ ... UPDATE, or the lock has been released.
IOERR	17	Underlying store error.

Example

```
EXEC CICS READ FILE('CUSTOMERS') INTO(REC) RIDFLD(CKEY) UPDATE END-EXEC
PARSE VAR REC NM '|' AD '|' CY '|' PH
PH = NEWPHONE                               /* mutate */
REC = NM || '|' || AD || '|' || CY || '|' || PH
EXEC CICS REWRITE FILE('CUSTOMERS') FROM(REC) END-EXEC
```

DELETE

Remove a record from a KSDS.

Format

```
EXEC CICS DELETE FILE(name)
                [RIDFLD(key) [KEYLENGTH(n) GENERIC [NUMREC(var)]]]
                [RESP(var)] [RESP2(var)]
END-EXEC
```

Description

If RIDFLD is supplied, deletes that key. Otherwise deletes the key locked by the most recent READ ... UPDATE on the same FILE.

IBM-canonical GENERIC delete-by-prefix: with KEYLENGTH(n) GENERIC, every record whose key starts with the first n bytes of RIDFLD is deleted. The count goes back through NUMREC(var) when supplied; RC = NOTFND when no record matches.

Options

FILE(name) — *required*

RIDFLD(key) Optional. When omitted, the most recent READ-UPDATE key is used.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Record deleted.
NOTFND	13	Key not present.

Condition	EIBRESP	Cause
INVREQ	16	No RIDFLD and no prior READ-UPDATE; or invalid name.
IOERR	17	Underlying store error.

Example

```
EXEC CICS DELETE FILE('CUSTOMERS') RIDFLD(CKEY) END-EXEC
```

VSAM file management — the IDCA (IDCAMS) transaction

IDCA is the BRICKS equivalent of mainframe **IDCAMS / AMSERV** — an Access Method Services screen for declaring, loading, dropping, listing, and printing VSAM-style KSDS clusters without writing an application program. It is a built-in transaction (no entry in `runtime/transactions.conf` is required) and is **admin-only**: only signed-on users in the `admin` group see the screen.

Invocation

At the `>` prompt on any 3270 terminal, type:

```
IDCA
```

ENTER drops you into a `SYSIN / SYSPRINT` screen. Type one or more control statements in the `SYSIN` area, press ENTER to run them, and the `SYSPRINT` area shows classic `IDC0001I FUNCTION COMPLETED...`-style messages plus the highest condition code. PF5 clears the input; PF7 / PF8 scroll the listing; PF3 exits.

Supported commands

Command	Syntax
DEFINE CLUSTER	DEFINE CLUSTER (NAME(name) [RECORDSIZE(avg max)] [KEYS(len off)] [INDEXED] [REPLACE])
REPRO (load)	REPRO INFILE(seq-file) OUTFILE(cluster) [KEYS(len off)] [SKIP(n)] [COUNT(n)]
REPRO (unload)	REPRO INFILE(cluster) OUTFILE(seq-file) [COUNT(n)]
DELETE	DELETE name [CLUSTER]
LISTCAT	LISTCAT [LEVEL(prefix)] [NAMES ALL]
PRINT	PRINT INFILE(cluster) [CHARACTER HEX DUMP] [COUNT(n)] [FROMKEY(k)] [TOKEY(k)]

Tokenisation is case-insensitive for keywords; values inside parens (file names, keys) are taken verbatim. A trailing `-` on a line is a continuation marker (JCL convention). Lines starting with `*` and text bracketed by `/* ... */` are comments.

Sequential file convention

`REPRO INFILE(...)` and `REPRO OUTFILE(...)` name files **relative to** `runtime_dir` (`./runtime/` by default — see the `runtime_dir` knob in `bricks.cnf`). Subdirectories are allowed, so an operator can keep loads under `loads/`, dumps under `dumps/`, ad-hoc files under `tmp/`, etc.:

```

REPRO INFILE(tmp/orders.in)          OUTFILE(ORDERS)          KEYS(7 0)
REPRO INFILE(loads/jan/customers.in) OUTFILE(CUSTOMERS)        KEYS(7 0)
REPRO INFILE(ORDERS)                  OUTFILE(dumps/orders.dat)

```

Bricks treats a name with a `.` in it as a sequential file and a name without one as a cluster, so subdirectory-bearing paths still pick LOAD vs UNLOAD direction correctly as long as the seq side has an extension. Each line of a sequential file becomes one record. The `KEYS(len off)` clause carves the key out of those bytes (length first, then offset — matching real IDCAMS); omit it and Bricks defaults to the cluster's declared `KeyMax` from `DEFINE CLUSTER`.

Parent directories are not auto-created. Writing to `OUTFILE(dumps/cust.dat)` when `runtime/dumps/` doesn't exist fails cleanly — `mkdir -p runtime/dumps` on the host first.

Security

Operator-supplied paths are sandboxed to `runtime_dir`. The resolver applies four independent traversal-defence layers before any file is opened: a per-segment lexical filter (alphanumeric + `_.` only, rejects `..`, rejects leading-dot, rejects absolute paths and backslashes), a `filepath.Clean` round-trip check, an absolute-prefix invariant against `runtime_dir`, and a `filepath.EvalSymlinks` check that rejects symlinks whose target leaves the sandbox. Attempts like `INFILE(..etc/passwd)`, `INFILE(/etc/passwd)`, or `INFILE(symlink)` where the symlink points outside the root all return `IDC3009I` with condition code 8 and never reach `os.Open`. The IDCA transaction is also admin-gated at the screen level; the path sandbox is defence in depth on top of that.

Worked example

```

DEFINE CLUSTER (NAME(CUSTOMERS) RECORDSIZE(80 80) KEYS(7 0))
REPRO INFILE(tmp/customers.in) OUTFILE(CUSTOMERS) KEYS(7 0)
LISTCAT LEVEL(CUST)
PRINT INFILE(CUSTOMERS) CHARACTER COUNT(5)

```

`LISTCAT` produces a CEMT-style fixed-column listing:

```

CATALOG LISTING -- 1 ENTRY(S)
  NAME           RECORDS  KEYMAX  RECMAX  LAST-MOD
  CUSTOMERS           42       7      80  28May26 15:04
IDC0001I FUNCTION COMPLETED, CONDITION CODE IS 0

```

Condition codes

Same scale as classic IDCAMS:

Code	Meaning
0	success
4	warning (e.g. duplicate records bypassed on REPRO, empty LISTCAT)
8	error (file not found, file locked by another terminal, DEFINE without REPLACE on an existing cluster)
12	severe (parse error, I/O failure)

The highest code any single statement produced is reported on the final `IDC0001I FUNCTION COMPLETED, HIGHEST CONDITION CODE WAS nn` line.

Lock behaviour and concurrency

DELETE and REPRO (load mode) check the file's lock state before proceeding. A file held under EXEC CICS READ ... UPDATE by another terminal is reported as IDC3009I ** name LOCKED BY term=Txxxx, NOT DELETED with condition code 8 — the same refusal the CEDA VSAM purge screen produces. Locks held by the operator's own terminal are ignored (so a user can DELETE a file they themselves are mid-update on, if they really want to).

Every IDCAMS command produces one line in the audit log (log/<timestamp>.log):

```
idcams=DEFINE target=CUSTOMERS term=T0001 user=admin status=OK detail="keymax=7 recmax=80"
idcams=REPRO target=CUSTOMERS term=T0001 user=admin status=OK detail="source=customers.in reco
rds=42 skipped=0 dup=0"
idcams=DELETE target=CUSTOMERS term=T0001 user=admin status=OK
```

Relationship to CEDA VSAM

CEDA VSAM and IDCA are two views of the same catalogue:

- **CEDA VSAM** — interactive table of every cluster with one-keystroke purge (P selector + secondary type-the-name confirmation). Use when you want to browse and purge a handful of files visually.
- **IDCA** — command-driven AMS. Use when you want to script a reproducible cluster lifecycle, declare metadata up front, or bulk-load from a sequential file.

Both go through the same audit logging and same lock checks; both operate on the same data/files.boltdb buckets that EXEC CICS WRITE populates from REXX and COBOL.

Offline IDCAMS — the idcams CLI

idcams is the **standalone** counterpart to the IDCA transaction. Same engine, same control statements, same audit log format — but it opens data/files.boltdb directly from the shell, so the bricks region must be **stopped** while it runs. This is the classic mainframe batch-IDCAMS model: shut the region down, run a deck of DEFINE / REPRO / DELETE statements, bring the region back up.

```
idcams [-c bricks.cnf] [-d datadir] [-r runtime_dir] -sysin deck.idcams
idcams [-c bricks.cnf] -e "DEFINE CLUSTER (NAME(X))"
idcams                                     # reads SYSIN from stdin
```

Flags:

Flag	Purpose
-c <path>	path to bricks.cnf (default: ./bricks.cnf then ../bricks.cnf)
-d <dir>	data directory holding files.boltdb (overrides bricks.cnf)
-r <dir>	runtime_dir sandbox root for REPRO INFILE/OUTFILE (overrides bricks.cnf)
-sysin <file>	read IDCAMS control statements from this file
-e <stmt>	execute one inline statement and exit
-logdir <dir>	audit-log directory (default: bricks.cnf log_location)
-q	suppress the SYSIN echo on stdout (listing only)

Exit code is the highest IDCAMS condition code (0=OK, 4=warning, 8=error, 12=severe), so the CLI plugs into shell pipelines and CI scripts directly. A REPRO that bypassed duplicate records exits 4; a

LISTCAT on an empty catalogue exits 4; a DELETE of a missing file exits 8.

Worked example:

```
$ idcams -sysin reload.idcams
idcams - data=data tmp=runtime/tmp sysin=reload.idcams
-----
IDCAMS  SYSTEM SERVICES                                BRICKS / KICKS
                                           TIME: 14:21:08

      DEFINE CLUSTER (NAME(CUSTOMERS) RECORDSIZE(80 80) KEYS(7 0) REPLACE)
IDC0508I CLUSTER CUSTOMERS DEFINED  (KEYMAX=7, RECMAX=80)
IDC0001I FUNCTION COMPLETED, CONDITION CODE IS 0

      REPRO INFILE(customers.in) OUTFILE(CUSTOMERS) KEYS(7 0)
IDC0005I NUMBER OF RECORDS PROCESSED WAS 12450
IDC0001I FUNCTION COMPLETED, CONDITION CODE IS 0

IDC0001I FUNCTION COMPLETED, HIGHEST CONDITION CODE WAS 0
$ echo $?
0
```

If the bricks region is still running, the CLI detects the bbolt lock contention and exits with code 16 plus a hint to use the IDCA transaction instead — no hang, no partial write, no chance of corrupting an in-flight region's view of the catalogue.

Chapter 8. KSDS browse commands

A **browse** walks a CICS FILE in B+tree key order. The browse runs inside a bbolt MVCC read transaction, so the cursor sees a stable point-in-time snapshot — concurrent WRITE / REWRITE / DELETE on the same FILE do not disturb an in-progress browse.

The browse is **per-task** and tied to one FILE. A program may have multiple browses open on different FILEs at once; a second STARTBR on the same FILE replaces the first (STARTBR is implicitly idempotent in CICS).

The dispatcher releases any cursor the program forgot to ENDBR via a defer handler `CloseBrowses()` at task end.

Bricks deviation from IBM canonical: browse cursors survive SYNCPOINT. Real CICS closes VSAM RPL browses at every SYNCPOINT; bricks browses are bbolt MVCC snapshots and remain valid until the matching ENDBR (or the dispatcher's task-end defer). Closing them at SYNCPOINT would be theater — the MVCC snapshot is already isolated from concurrent writers. This is what lets a STARTBR on CECI PF5 #1 drive READNEXT on PF5 #2 (each press is a per-PF5 implicit SYNCPOINT).

STARTBR

Open a browse cursor on a KSDS file.

Format

```
EXEC CICS STARTBR FILE(name)
                        [RIDFLD(start)]
                        [GTEQ | EQUAL]
                        [GENERIC]
                        [KEYLENGTH(n)]
                        [REQID(handle)]
                        [RESP(var)] [RESP2(var)]

END-EXEC
```

REQID(handle) is the IBM-canonical concurrent-browse handle: a program may hold several independent browses against the same file by tagging each STARTBR / READNEXT / READPREV / ENDBR / RESETBR with a distinct integer handle. Without REQID the file's single default cursor is used — backward-compatible with every existing bricks program.

Description

Opens a B+tree browse cursor on the file. With no RIDFLD, positions on the first key. With RIDFLD and GTEQ (the IBM default and bricks default), positions on the first key $\geq$ start. With EQUAL, requires an exact match (NOTFND if absent). With GENERIC and KEYLENGTH(n), positions on (and walks only through) keys whose first n bytes match the first n bytes of RIDFLD.

Options

FILE(name) — *required*

RIDFLD(start) Starting key. Default: first key in the file.

GTEQ | EQUAL Comparison rule. Default GTEQ.

GENERIC Restrict the walk to keys whose prefix matches.

KEYLENGTH(n) With GENERIC, the prefix length to match.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Cursor opened.
NOTFND	13	EQUAL requested but no such key.
INVREQ	16	Missing required option, invalid name, etc.
IOERR	17	Underlying store error.

Example

```
EXEC CICS STARTBR FILE('CUSTOMERS')
                        RIDFLD('NY-')
                        GENERIC KEYLENGTH(3)

END-EXEC
```

READNEXT

Step the open browse cursor forward by one record.

Format

```
EXEC CICS READNEXT FILE(name)
                        {INTO(target) | SET(target)}
                        [RIDFLD(keyvar)]
                        [LENGTH(len)]
END-EXEC
```

Description

Advances the cursor and returns the next record (or the first record if this is the first READNEXT after STARTBR). When RIDFLD is a bare variable, the matching key is written back; when LENGTH is a bare variable, the actual record length is written back.

Returns ENDFILE past the last key (or past the end of the GENERIC prefix). Records that were deleted by a concurrent transaction between STARTBR and the read are skipped automatically (bounded forward loop, no goroutine-stack risk).

Options

FILE(name) — *required*

INTO(target) / SET(target) — *one required*

RIDFLD(keyvar) Receives the key of the record returned.

LENGTH(len) Receives the actual record length.

Bricks deviation from IBM canonical

LENGTH(var) on READNEXT (browse-family) is OUTPUT-only, not input-as-max-buffer. Real CICS reads the caller's LENGTH(var) as a cap on the bytes returned and writes back the actual size; bricks ignores the input value and only writes back. The write-back-then- read-as-input shape would otherwise corrupt subsequent reads in a loop (LENGTH(LEN) on call 1 stores 4; call 2 caps a 13-byte record to 4). The bricks KSDS store is fixed-size-per-record so there is no operator-meaningful truncation here. A LENGTH(20) literal on a 13-byte record returns the full 13 bytes (the literal is silently ignored).

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Record returned.
ENDFILE	20	Past the last key (or out of GENERIC prefix).
INVREQ	16	No STARTBR is open on this FILE.
IOERR	17	Underlying store error.

Example

See STARTBR example, then:

```
DO FOREVER
  EXEC CICS READNEXT FILE('CUSTOMERS') INTO(REC) RIDFLD(K) END-EXEC
  IF EIBRESP = 20 THEN LEAVE
  SAY K ':' REC
END
```

READPREV

Step the open browse cursor backward by one record.

Format

```
EXEC CICS READPREV FILE(name)
                        {INTO(target) | SET(target)}
                        [RIDFLD(keyvar)]
                        [LENGTH(len)]
END-EXEC
```

Description

Same writeback rules as READNEXT. Returns ENDFILE before the first key (or before the start of the GENERIC prefix). Useful for paginating backward through a key range.

Options

As for READNEXT.

Bricks deviation from IBM canonical

Same shape as READNEXT: LENGTH(var) is OUTPUT-only, not input-as-max-buffer. The write-back-then-read-as-input pattern would corrupt subsequent reads in a backward walk; bricks ignores the operator-supplied input value and only writes back the actual record size. Real VSAM variable-length permits LENGTH-as-max- buffer; the bricks KSDS-fixed model has no operator-meaningful truncation here.

Conditions

As for READNEXT, with ENDFILE meaning *before the first key*.

Example

```
EXEC CICS READPREV FILE('CUSTOMERS') INTO(REC) RIDFLD(K) END-EXEC
```

RESETBR

Reposition the cursor of an open browse without closing the underlying read transaction.

Format

```
EXEC CICS RESETBR FILE(name)
                        RIDFLD(start)
                        [GTEQ | EQUAL]
                        [GENERIC]
                        [KEYLENGTH(n)]
END-EXEC
```

Description

Cheaper than ENDBR + STARTBR when the program wants to jump within the same browse session — the read snapshot is preserved.

Options

As for STARTBR. RIDFLD is required.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Cursor repositioned.
INVREQ	16	No STARTBR is open on this FILE.
NOTFND	13	EQUAL requested but no such key.

Example

```
EXEC CICS RESETBR FILE('CUSTOMERS') RIDFLD('00500') END-EXEC
```

ENDBR

Close an open browse cursor.

Format

```
EXEC CICS ENDBR FILE(name)
END-EXEC
```

Description

Releases the cursor and the underlying bbolt read transaction. The dispatcher releases any cursor the program forgot to ENDBR at task end, but explicit ENDBR is recommended.

Options

FILE(name) — *required*

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Cursor closed.
INVREQ	16	No STARTBR is open on this FILE.

Example

```
EXEC CICS ENDBR FILE('CUSTOMERS') END-EXEC
```

Chapter 8a. File Control on tmp_dir Sequential Files

The browse verbs in Chapter 8 (STARTBR / READNEXT / READPREV / RESETBR / ENDBR) and the direct positional READ FILE ... RIDFLD also operate on **sequential files** living in the tmp_dir sandbox — the same files READQ TD / WRITEQ TD and the REXX LINEIN / LINEOUT family read and write. The routing is transparent and existence-based: if tmp_dir/<name> exists, the browse verbs use the line-record / RBA-keyed sequential backend; otherwise they fall through to the KSDS bbolt store (Chapter 8). When the same name is present in BOTH back-ends, the sequential file wins on STARTBR, and bricks logs a WARN line at startup so the operator knows the routing.

Use this chapter when the program needs **non-FIFO** access to a sequential file — random-access reads by byte offset, backward walks, or jumping to a known position with RESETBR. Pure FIFO consumption stays cheapest via READQ TD (Chapter 9).

RBA semantics

The Relative Byte Address (RBA) is the **decimal byte offset of the start of a record**, expressed as plain ASCII text — consistent with READQ TS NUMITEMS and the rest of the bricks numeric conventions. A three-line file "line1\nline2\nline3\n" (each line is 5 bytes plus the LF terminator) has records at RBA 0, 6, and 12.

Bricks deviation from IBM canonical: real CICS ESDS exposes RBA as a 4-byte fullword (PIC S9(8) COMP). Bricks renders it as decimal text so REXX programs and COBOL PIC X(n) host variables both see a value they can arithmetic on without datatype conversion. The deviation is called out here and in README.md so a porter from a z/OS shop isn't surprised.

STARTBR (and RESETBR) with an RBA that falls **mid-line** rounds backward to the prior LF+1 so the next READNEXT returns the **containing** record rather than a garbled split. Mid-line RIDFLD on a direct READ FILE ... RIDFLD(rba) rounds the same way. This matches the operator's mental model: "the byte I named lives inside record X; give me X."

Record boundary

Records terminate with a single LF (0x0A). The LF is excluded from the INT0(area) payload. A trailing line without LF is returned once; the next READNEXT returns ENDFILE. The byte-level encoding is the same as Chapter 9's TD queues — see "The tmp_dir sandbox" for the strict ASCII rule.

Bricks deviation: real ESDS uses an RDW (Record Descriptor Word) or fixed-length records, not LF-terminated lines. Bricks aligns the sequential format with the rest of the tmp_dir contract (REXX LINEIN,

COBOL READQ TD) so one file is portable across all three I/O surfaces.

Verb subset

Verb	Supported?	Notes
STARTBR	yes	RIDFLD optional — defaults to BOF (= RBA 0). EQUAL, GENERIC, KEYLENGTH rejected with INVREQ.
READNEXT	yes	Returns next line; RIDFLD writeback = decimal RBA.
READPREV	yes	Walks backward; same RBA conventions.
RESETBR	yes	RIDFLD parsed as decimal RBA; backward-rounds same as STARTBR.
ENDBR	yes	Releases the FD.
READ FILE ... RIDFLD(rba)	yes	One-shot positional read; LENGTH semantics same as READNEXT.
UNLOCK FILE	yes	No-op accepted (bricks has no record locks).
WRITE FILE	no — INVREQ "use WRITEQ TD"	
REWRITE FILE	no — INVREQ "use WRITEQ TD"	
DELETE FILE	no — INVREQ "use DELETEQ TD"	

LENGTH semantics (divergence from KSDS!)

On KSDS browse verbs (Chapter 8), LENGTH(var) is **output-only** — the store decides record size and writes it back to var. The recent KSDS fix made this explicit: treating LENGTH-on-READNEXT as an INPUT bound triggered the write-back-then-read-as-input bug that truncated record 2 to the length of record 1.

On tmp_dir sequential files, LENGTH(var) is **INPUT bound + OUTPUT actual** — the program's declared buffer capacity is honoured. If the record exceeds the buffer, the verb returns LENGERR, truncates the payload to the buffer size, and writes the **actual** (untruncated) record length back to var. The "actual not truncated" rule is the canonical real-CICS LENGERR semantic so a program can see how much was lost.

```
LEN = 80                                /* buffer cap */
EXEC CICS READNEXT FILE('orders.txt') INTO(REC) RIDFLD(K) LENGTH(LEN) END-EXEC
IF EIBRESP = 22 THEN                    /* LENGERR */
    SAY 'Record was' LEN 'bytes; lost' (LEN - 80) 'bytes after truncation'
```

RESP / RESP2 mapping

Condition	EIBRESP	Cause
NORMAL	0	Record returned (or ENDBR succeeded).
INVREQ	16	EQUAL / GENERIC / KEYLENGTH on sequential file; bad RBA (negative, parse failure); WRITE / REWRITE / DELETE against a tmp_dir name; missing STARTBR cursor.
LENGERR	22	Record exceeded LENGTH(var) buffer; payload truncated. Actual length writeback.
NOTFND	13	File not present in tmp_dir at READ FILE ... RIDFLD(rba), or STARTBR opened OK but the RBA landed past EOF on the next READNEXT. Real-CICS FILENOTFOUND (RESP=12) is not separately surfaced; programs already test for NOTFND.

Condition	EIBRESP	Cause
ENDFILE	20	READNEXT past end, or READPREV past beginning. (Both directions reuse one constant — matches real CICS.)
IOERR	17	Filesystem error (rare; the sandbox is local-only).

Interaction with `WRITEQ TD` mid-browse

Bricks's sequential backend is single-FD-per-cursor. When a concurrent task issues `WRITEQ TD` against the **same file** while this task holds a `STARTBR` open, the appended bytes ARE picked up: on each `READNEXT` whose internal position has reached the size we captured at `STARTBR`, the cursor re-Stats the file descriptor. If the file has grown, the new bytes flow into the next `READNEXT`; otherwise, `ENDFILE` is returned as normal.

This is a documented deviation — real CICS holds a transient view of the dataset at `STARTBR` time and does not see mid-browse writes. Bricks chose the "moving target" rule so a long-running report transaction can stream a file that's still being written by a feeder task.

Real-CICS-veteran gotchas

- **RBA-is-decimal-text** — not a 4-byte `COMP` fullword. A REXX variable receiving RBA is a plain string; arithmetic works directly.
- **Record boundary is LF**, not a fixed-length record or `RDW`. A binary-mode browse is out of scope today.
- **STARTBR without RIDFLD** = BOF. Real CICS requires `RIDFLD`.
- **Browses survive SYNCPOINT** — same deviation already documented for `KSDS` (Chapter 8). Released by `ENDBR` or by the dispatcher's defer at task end.
- **Mid-browse writes are visible** via Stat-refresh-on-EOF. Real CICS doesn't do this.

Chapter 9. Temporary storage and transient data commands

Bricks supports two related queue families:

- **Temporary Storage (TS)** — ordered, persistent sequences of opaque byte items inside the `bbolt` database. Each queue is a sub-bucket whose keys are 8-byte big-endian item numbers (1, 2, 3 ...) and whose values are the item payloads. Item-addressed; reads and writes are $O(\log N)$. Use for ephemeral state, scratch storage, producer/consumer chaining inside the database.
- **Transient Data extra-partition (TD)** — sequential text files in a sandboxed on-disk directory (`tmp_dir`). Line-oriented; read-once-and-advance / append-only semantics. Use for staging imports from text files into `VSAM` and exports out of it. The REXX `LINEIN` / `LINEOUT` / `STREAM` family hits the same backend, so a file written by `COBOL` is readable by REXX and vice-versa.

The encoding contract for `TD` files is strict: ASCII bytes only (`0x09` TAB, `0x20–0x7E` printable), no EBCDIC, no UTF; lines are terminated by a single LF (`0x0A`); CR (`0x0D`) is rejected on write with `INVREQ`. The sandbox enforces a flat namespace under `tmp_dir` — no sub-directories, no leading dot, no `..`, no slashes in the queue name. See [The `tmp\_dir` sandbox](#) below for the full rules.

READQ TS

Read an item from a temporary storage queue.

Format

```
EXEC CICS READQ TS QUEUE(name)
                        INTO(target)
                        [ITEM(n) | NEXT]
                        [LENGTH(len)]
                        [NUMITEMS(num)]
END-EXEC
```

QNAME is accepted as a synonym for QUEUE.

Description

READQ TS runs in one of two modes:

1. **Cursor mode** — advance the **per-task implicit cursor**: the first cursor-less READQ on a queue returns item 1, the second returns item 2, and so on. The cursor is keyed on the running TxCB and released when the task ends — a fresh invocation of the same TRANSID starts at item 1 again.
2. **Explicit-item mode** — ITEM(n) reads that exact item without touching the cursor. Useful for random-access reads where the program knows the item number in advance.

The mode is selected by IBM-canonical strict semantics. Cursor mode fires when NEXT is present OR when ITEM is absent. Explicit-item mode fires when ITEM(n) is present without NEXT; the variable's current value names the item to read and must be a positive integer. The four operator patterns are:

READQ TS QUEUE(q) INTO(d)	– cursor mode; item number not returned
READQ TS QUEUE(q) INTO(d) NEXT	– cursor mode; same as above
READQ TS QUEUE(q) INTO(d) NEXT ITEM(N)	– cursor mode; writes item number to N
READQ TS QUEUE(q) INTO(d) ITEM(N)	– explicit mode; N must be a positive integer input

Bricks deviation from IBM canonical: the per-task implicit cursor survives SYNCPOINT. Real CICS clears the read cursor on recoverable TSQs at every SYNCPOINT; bricks does not. The bricks design treats SYNCPOINT as a unit-of-work commit boundary that does not touch cursor state, browse cursors, TD handles, WEB sessions, or DOCUMENT tokens. This is what lets a CECI operator walk a queue across PF5 presses (each press is a per-PF5 implicit SYNCPOINT), and what lets a long-running REXX program issue periodic SYNCPOINTS without rewinding its read loop.

Options

QUEUE(name) / QNAME(name) — *required*

INTO(target) — *required*

ITEM(n) | NEXT NEXT (or ITEM absent) selects cursor mode. ITEM(n) without NEXT is **strict input**: the variable must hold a positive integer or INVREQ fires with a short error pointing the operator at NEXT. NEXT ITEM(var) is the cursor-loop form — the resolved item number is written back into var on success.

LENGTH(len) Receives the actual item length.

NUMITEMS(num) Receives the queue's current high-water item count.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Item returned.
ITEMERR	26	ITEM(n) out of range, or implicit cursor past the last item.
QIDERR	44	Queue does not exist.
INVREQ	16	Missing required option, invalid name, or ITEM(var) (no NEXT) when var is unset / non-numeric.

Example

```
DO FOREVER
  EXEC CICS READQ TS QUEUE(QNM) INTO(REC) NEXT END-EXEC
  IF EIBRESP = 26 THEN LEAVE          /* end of queue */
  SAY REC
END
```

WRITEQ TS

Append a new item to a queue, or rewrite an existing item.

Format

```
EXEC CICS WRITEQ TS QUEUE(name)
                        FROM(data)
                        [ITEM(n) REWRITE]
                        [NUMITEMS(var)]
                        [MAIN | AUXILIARY]
                        [RESP(var)] [RESP2(var)]
END-EXEC
```

QNAME is accepted as a synonym for QUEUE. NUMITEMS(var) receives the queue's current item count after the write — IBM canonical; lets programs drive paging without a follow-up READQ. MAIN and AUXILIARY are storage-class hints from real CICS; bricks uses a uniform bbolt-backed store so both are accepted silently.

Description

In **append** mode (no ITEM REWRITE), the item is added at the next sequence number; an in-memory high-water-mark counter assigns the number without scanning. When ITEM(var) is a bare variable, the assigned item number is written back.

In **rewrite** mode (both ITEM(n) and REWRITE present), item n is replaced.

The write commits in a single bbolt transaction so a crash mid-write leaves either the prior state or the new one — never a partial.

Options

QUEUE(name) / QNAME(name) — *required* **FROM(data)** — *required* **ITEM(n) REWRITE** Both required for rewrite mode; either alone is rejected.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Item written.
ITEMERR	26	ITEM(n) REWRITE and item n does not exist.
INVREQ	16	Missing required option, invalid name, or ITEM/REWRITE not paired.
IOERR	17	Underlying store error.

Example

```
EXEC CICS WRITEQ TS QUEUE('AUDIT') FROM(MSG) ITEM(I) END-EXEC
SAY 'Wrote item' I
```

DELETEQ TS

Delete an entire temporary storage queue.

Format

```
EXEC CICS DELETEQ TS QUEUE(name) END-EXEC
```

Description

Drops the queue's sub-bucket and resets the in-memory counters and cursors. Subsequent WRITEQ will recreate it starting at item 1.

Options

QUEUE(name) / QNAME(name) — *required*

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Queue deleted.
QIDERR	44	Queue does not exist.

Example

```
EXEC CICS DELETEQ TS QUEUE('AUDIT') END-EXEC
```

READQ TD

Read the next line from a sequential text file in tmp_dir.

Need non-FIFO access? READQ TD is the cheapest read for plain forward-sequential consumption. When the program needs random access by byte offset, backward walks, or repositioning by RBA on the same file, switch to the browse verbs against the same tmp_dir

name — see [Chapter 8a. File Control on tmp_dir Sequential Files](#). Both verbs read the same on-disk bytes, so a file written via WRITEQ TD is immediately browsable via STARTBR.

Format

```
EXEC CICS READQ TD QUEUE(name)
                        INTO(target)
                        [LENGTH(len)]
END-EXEC
```

QNAME is accepted as a synonym for QUEUE. name is a flat filename under tmp_dir; sub-directories and traversal are rejected with INVREQ.

Description

Reads the next LF-terminated line from tmp_dir/name, strips the terminating LF, and copies the payload into target. The file is auto-opened in read mode on the first call per task and the read cursor advances item-by-item across subsequent calls. A program that WRITES then READs the same queue causes the handler to close the write handle and reopen for read; the read cursor restarts at line

1. The handle is closed automatically at task end.

Options

QUEUE(name) / QNAME(name) — *required*

INTO(target) — *required*

LENGTH(len) Receives the line's byte count (after the LF strip).

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Line returned.
QZERO	12	End-of-file. target is untouched; bricks closes the handle so a later READQ on the same queue rewinds to line 1.
QIDERR	44	File does not exist in tmp_dir.
INVREQ	16	Missing INTO / invalid name / sandbox rejection.
IOERR	17	Underlying file-system error.

Example

```
PERFORM IMPORT-ONE UNTIL DONE-FLAG = 'Y'.
...
IMPORT-ONE.
  MOVE SPACES TO REC.
  EXEC CICS READQ TD QUEUE('orders.sample.txt') INTO(REC) END-EXEC.
  IF EIBRESP = 12 THEN
    MOVE 'Y' TO DONE-FLAG
  END-IF.
  IF EIBRESP = 0 THEN
    PERFORM HANDLE-RECORD
  END-IF.
```

WRITEQ TD

Append a line to a sequential text file in tmp_dir.

Format

```
EXEC CICS WRITEQ TD QUEUE(name)
                        FROM(data)
                        [LENGTH(len)]
END-EXEC
```

QNAME is accepted as a synonym for QUEUE.

Description

Appends data plus a single LF to tmp_dir/name. The file is auto-opened in append mode (creating it if absent) on the first call per task. Trailing ASCII spaces on data are right-stripped before write — COBOL PIC X(n) values arrive padded to the right and the canonical text-file convention is to strip them. The payload is validated byte-by-byte: anything outside 0x09 (TAB), 0x0A (LF), or 0x20–0x7E (printable ASCII) causes the write to fail with INVREQ and a byte-offset diagnostic. CR (0x0D) is explicitly rejected — bricks emits Unix-style LF-only files.

Options

QUEUE(name) / QNAME(name) — *required*

FROM(data) — *required*

LENGTH(len) Accepted for syntactic compatibility; the actual byte count is LENGTH(data) after rstrip.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Line appended.
INVREQ	16	Missing FROM / invalid name / sandbox rejection / non-ASCII byte / CR in payload.
IOERR	17	Underlying file-system error.

Example

```
DO I = 1 TO REC.0
  EXEC CICS WRITEQ TD QUEUE('export.txt') FROM(REC.I) END-EXEC
END
```

DELETEQ TD

Delete a sequential text file in `tmp_dir`.

Format

```
EXEC CICS DELETEQ TD QUEUE(name) END-EXEC
```

Description

Closes any handle the task has open on `name`, then unlinks the file. Already-gone is treated as success (matches the `DELETE FILE` forgiving semantics in bricks). Use to reset an export staging file before a fresh batch.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	File deleted (or already absent).
INVREQ	16	Invalid name / sandbox rejection.
IOERR	17	Underlying file-system error.

Example

```
EXEC CICS DELETEQ TD QUEUE('export.txt') END-EXEC.
```

The `tmp_dir` sandbox

`tmp_dir` is configured in `bricks.cnf`:

```
tmp_dir = runtime/tmp
```

Defaults to `runtime/tmp` (under `runtime_dir`) when the line is omitted. The directory is created at startup if missing.

Name rules. A queue name must match `[A-Za-z0-9._-]{1,255}` — no leading dot, no `..`, no slash, no backslash, no NUL. Bricks also runs `filepath.Rel` against every resolved path as defense-in-depth, so symlink shenanigans bounce too.

Encoding rules. ASCII only — no EBCDIC, no UTF-8, no UTF-16. The write path validates every byte and rejects with `INVREQ` on the first violation, naming the offending offset and hex value. Lines terminate with LF only; CR is rejected on write. The read path splits on LF and preserves any CR bytes it finds verbatim — bricks does **not** silently strip them. If you need to import a file authored on Windows, pre-strip the CRs:

```
tr -d '\r' < windows.csv > runtime/tmp/clean.csv
```

Cross-language interop. Both COBOL (READQ TD / WRITEQ TD) and REXX (LINEIN / LINEOUT / STREAM) hit the same `*cics.TmpStore` backend. A file produced by one language is readable by the other; the only constraint is that both sides agree on the column format inside each line (canonical bricks convention: pipe-delimited).

Task-end cleanup. Every handle the program opens is closed automatically when the task ends. A program that forgets to call `STREAM CLOSE` or `DELETEQ TD` does not leak descriptors.

Chapter 10. Recovery and condition handling

This chapter covers two related families:

- **Unit-of-work commands** — `SYNCPPOINT` and `SYNCPPOINT ROLLBACK` group multiple file / TS mutations into one atomic step that can be committed or undone.
- **Non-local control commands** — `HANDLE CONDITION`, `IGNORE CONDITION`, `HANDLE AID`, and `HANDLE ABEND` arm program labels that the dispatcher branches to on a non-NORMAL response, on a particular AID key, or when the task abends. Together with the per-call `EIBRESP` test ([Chapter 12](#)) and `REXX SIGNAL ON ERROR` ([Chapter 18](#)) they cover the full range of CICS error-handling idioms.

How the unit of work works

Each task owns a **journal** of pending undo entries that lives on its `TxCB`. Every successful `WRITE` / `REWRITE` / `DELETE` and every `WRITEQ TS` / `DELETEQ TS` appends an inverse operation to the journal *immediately after* the underlying `bbolt` write commits.

- `SYNCPPOINT` clears the journal — the work becomes irrevocably committed, and a new unit of work begins.
- `SYNCPPOINT ROLLBACK` walks the journal in reverse and applies each inverse op (re-writing pre-images, re-creating deleted records, restoring queue contents, etc.), then clears the journal.
- `RETURN` performs an **implicit** `SYNCPPOINT` before the task ends.
- `ABEND` performs an **implicit** `SYNCPPOINT ROLLBACK` *unless* a `HANDLE ABEND` exit catches it; the exit may then choose whether to commit or roll back explicitly.

In-flight uncommitted writes are visible to concurrent tasks (this matches IBM CICS under `READ NO UPDATE`); the rollback model exists to recover *the current task* from its own partially-applied work, not to provide multi-task isolation.

How the condition / AID / abend traps work

After every `EXEC CICS` command, the dispatcher writes `EIBRESP` and then consults three optional per-task tables:

Table	Populated by	Consulted	Branches when
Condition map	<code>HANDLE CONDITION</code> / <code>IGNORE CONDITION</code>	After every command	<code>EIBRESP</code> $\neq$ <code>NORMAL</code> and a label is armed for that condition (or for <code>ERROR</code>).
AID map	<code>HANDLE AID</code>	After every command that updates <code>EIBAID</code> (<code>SEND MAP</code> , <code>RECEIVE MAP</code> , <code>RECEIVE</code>)	The new <code>EIBAID</code> matches an armed key (<code>PF1–PF24</code> , <code>PA1–PA3</code> , <code>ENTER</code> , <code>CLEAR</code> , or the catch-all <code>ANYKEY</code>).

Table	Populated by	Consulted	Branches when
Abend exit	HANDLE ABEND	When the task abends	An exit is armed and not cancelled.

When a trap fires, control transfers to the named label as if the program had `GO TO`'d (COBOL) or `SIGNAL`'d (REXX) it directly. The trap *stays armed* across the branch — re-arming after each fire is not required (and the `IGNORE` / bare-name disarm forms exist for when the program wants the trap to stop firing).

SYNCPPOINT

Commit the current unit of work and begin a new one.

Format

```
EXEC CICS SYNCPPOINT END-EXEC
```

Description

Clears the task's undo journal. All file / TS writes performed since the last `SYNCPPOINT` (or since the task began) are made permanent; they can no longer be rolled back. A fresh, empty journal is started and subsequent mutations accumulate against it.

`SYNCPPOINT` is implicit on `EXEC CICS RETURN` — programs that run to completion never need to call it explicitly. The explicit form is useful in long-running tasks that want to checkpoint partial progress, or in programs that mix recoverable phases with phases where rollback is no longer meaningful.

Options

None.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Journal cleared.

`SYNCPPOINT` cannot fail in bricks because the underlying `bbolt` writes have already committed by the time the journal entry was appended.

Example

```
* Phase 1: build the audit row, commit immediately so a later
* validation failure does not undo the audit trail.
EXEC CICS WRITEQ TS QUEUE('AUDIT') FROM(Ws-LOG) END-EXEC.
EXEC CICS SYNCPPOINT END-EXEC.

* Phase 2: real business work; rolled back if validation fails.
EXEC CICS REWRITE FILE('CUSTOMERS') FROM(Ws-CUST) END-EXEC.
IF Ws-INVALID
  EXEC CICS SYNCPPOINT ROLLBACK END-EXEC
  EXEC CICS ABEND ABCODE('VAL1') END-EXEC
END-IF.
```

SYNCPPOINT ROLLBACK

Discard the current unit of work and begin a new one.

Format

```
EXEC CICS SYNCPPOINT ROLLBACK END-EXEC
```

Description

Walks the task's undo journal in reverse and applies each inverse operation:

- WRITE is undone by deleting the new key.
- REWRITE is undone by re-writing the captured pre-image.
- DELETE is undone by re-writing the deleted record.
- WRITEQ TS (append) is undone by deleting the appended item.
- WRITEQ TS ... REWRITE is undone by re-writing the pre-image.
- DELETEQ TS is undone by restoring every prior item from a snapshot taken before the delete.

After the walk the journal is cleared and a new unit of work begins.

ROLLBACK is implicit when EXEC CICS ABEND runs **without** a HANDLE ABEND exit. When an exit *does* intercept the abend, the exit code decides whether to call SYNCPPOINT ROLLBACK (typical) or SYNCPPOINT (commit partial work and continue).

Options

ROLLBACK is the only option — it is what distinguishes this form from plain SYNCPPOINT.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	All journal entries successfully reverted.
ROLLEDBACK	82	One or more inverse ops failed (the underlying bbolt error is logged). The journal is still cleared; some writes may remain partially undone.

Example

```
ADDRESS CICS

EXEC CICS WRITE FILE('ACCT') RIDFLD(DEBIT_ID) FROM(DEBIT_REC) END-EXEC
EXEC CICS WRITE FILE('ACCT') RIDFLD(CREDIT_ID) FROM(CREDIT_REC) END-EXEC

IF DEBIT_TOTAL <> CREDIT_TOTAL THEN DO
  EXEC CICS SYNCPPOINT ROLLBACK END-EXEC      /* both writes undone */
  EXEC CICS ABEND ABCODE('IMBL') END-EXEC
END

EXEC CICS SYNCPPOINT END-EXEC                  /* commit both writes */
```

HANDLE CONDITION

Arm program labels to receive control on named EXEC CICS conditions.

Format

```
EXEC CICS HANDLE CONDITION cond1[(label)] cond2[(label)] ...
END-EXEC
```

Description

Builds a per-task map from condition name to label. After every EXEC CICS command, when EIBRESP ≠ NORMAL, the dispatcher consults the map:

- 1. If the *exact* condition name is armed, control branches to the named label.
- 2. Otherwise, if the catch-all ERROR is armed, control branches to that label.
- 3. Otherwise, the response code is left in EIBRESP for the program to test, exactly as if HANDLE CONDITION had never been issued.

The (label) form arms the trap; the bare condition name (no parens) **disarms** it. Arming a previously-armed condition replaces the old label.

HANDLE CONDITION is the legacy CICS error-handling style. Modern programs that use RESP(rc) style (or REXX SIGNAL ON ERROR) do not need it; the two styles can coexist within the same program because both ultimately read EIBRESP.

Options

condN(label) Arm condN (one of the names in [Chapter 12. Response codes](#), e.g. NOTFND, DUPREC, MAPFAIL, INVREQ, IOERR, LENGERR, QIDERR, ITEMERR, ENDFILE, PGMIDERR, EOC) so it branches to label.

condN (no parens) Disarm condN.

ERROR(label) Arm a catch-all that fires for any non-NORMAL response not matched by a more specific arm.

A single HANDLE CONDITION may combine any number of arm and disarm options.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Map updated.
INVREQ	16	Unknown condition name.

Example

```
ADDRESS CICS

EXEC CICS HANDLE CONDITION NOTFND(NOREC)
                        DUPREC(DUP)
                        ERROR(OOPS)

END-EXEC

EXEC CICS READ FILE('CUSTOMERS') INTO(REC) RIDFLD(CKEY) END-EXEC
SAY 'Found:' REC
EXIT

NOREC: SAY 'No customer with key' CKEY; EXIT 4
DUP:   SAY 'Duplicate key on insert'; EXIT 8
OOPS:  SAY 'Other CICS error, RC=' EIBRESP; EXIT 12
```

IGNORE CONDITION

Suppress the `HANDLE CONDITION` trap for one or more conditions.

Format

```
EXEC CICS IGNORE CONDITION cond1 cond2 ...
END-EXEC
```

Description

Marks each named condition as *ignored*: even if `HANDLE CONDITION` previously armed it, the trap will no longer fire and the program must test `EIBRESP` itself. `IGNORE CONDITION` stays in effect until a later `HANDLE CONDITION cond(label)` re-arms the same condition.

The distinction between *unarmed* (never armed, or disarmed by the bare-name form) and *ignored* matters only when `ERROR(label)` is armed: an unarmed condition still falls through to `ERROR`, but an ignored one does not.

Options

Each operand is a bare condition name — no parentheses, no label. Multiple conditions may be combined in a single call.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Map updated.
INVREQ	16	Unknown condition name.

Example

```
EXEC CICS HANDLE CONDITION ERROR(BADIO) END-EXEC.

* We *expect* NOTFND on this READ – don't trip the ERROR trap.
EXEC CICS IGNORE CONDITION NOTFND END-EXEC.

EXEC CICS READ FILE('CUSTOMERS') INTO(WS-REC) RIDFLD(WS-KEY) END-EXEC.
IF EIBRESP = 13
    PERFORM CREATE-NEW-CUSTOMER
END-IF.
```

HANDLE AID

Arm program labels to receive control on a particular attention key.

Format

```
EXEC CICS HANDLE AID key1[(label)] key2[(label)] ...
END-EXEC
```

Description

Builds a per-task map from AID byte to label. After any command that updates EIBAID (SEND MAP, RECEIVE MAP, RECEIVE), the dispatcher consults the map and branches if the new EIBAID matches an armed key.

The bare key name (no parens) disarms a previously-armed key. ANYKEY is a catch-all matched by any AID for which no specific arm exists.

Options

keyN(label) / keyN keyN is one of ENTER, CLEAR, PA1–PA3, PF1–PF24, or ANYKEY. Parenthesised arms the trap; bare disarms it.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Map updated.
INVREQ	16	Unknown AID name.

Example

```
ADDRESS CICS

EXEC CICS HANDLE AID PF3(QUIT) PF12(QUIT) CLEAR(QUIT) END-EXEC

DO FOREVER
  EXEC CICS RECEIVE MAP('CUST1') INTO(SCR.) END-EXEC
  ... process ...
  EXEC CICS SEND MAP('CUST1') FROM(SCR.) ERASE END-EXEC
END

QUIT:
  EXEC CICS RETURN END-EXEC
```

HANDLE ABEND

Arm a program-level abend exit.

Format

```
EXEC CICS HANDLE ABEND { LABEL(label) | PROGRAM(name) | CANCEL | RESET }
END-EXEC
```

Description

Registers an exit that runs when the task abends (either via `EXEC CICS ABEND` or via an unhandled runtime error). The exit *replaces* the implicit `SYNCPPOINT ROLLBACK` that an uncaught abend would perform — it is now the exit's responsibility to call `SYNCPPOINT` or `SYNCPPOINT ROLLBACK` explicitly.

`EIBABCODE` is set to the abend code (default `AAAA`) before the exit receives control.

Options

LABEL(label) Branch to `label` within the current program when the task abends. This is the form REXX programs almost always use.

PROGRAM(name) `XCTL` to the named program when the task abends. Currently accepted by the parser; the runtime treats `PROGRAM(name)` as equivalent to `LABEL(name)` (i.e. it branches inside the current program rather than transferring control).

CANCEL Disarm the current exit and restore the previous one (one level of nesting is preserved).

RESET Re-enable an exit previously suspended by `CANCEL`.

Exactly one of `LABEL`, `PROGRAM`, `CANCEL`, `RESET` must be given.

Conditions

Condition	EIBRESP	Cause
NORMAL	0	Exit registered or restored.
INVREQ	16	Conflicting / missing options, or RESET with no prior exit.

Example

```
ADDRESS CICS

EXEC CICS HANDLE ABEND LABEL(CLEANUP) END-EXEC

EXEC CICS WRITE FILE('LEDGER') RIDFLD(K) FROM(REC) END-EXEC
EXEC CICS WRITE FILE('AUDIT') RIDFLD(K) FROM(REC) END-EXEC

/* something later may ABEND */
EXIT

CLEANUP:
  SAY 'Caught abend' EIBABCODE
  EXEC CICS SYNCPOINT ROLLBACK END-EXEC      /* undo both writes */
  EXEC CICS RETURN END-EXEC
```

Chapter 11. The Execute Interface Block (EIB)

The EIB is a per-task scratch area populated by the dispatcher before the program runs and by each EXEC CICS command on return. In bricks the EIB is exposed as a set of well-known variable names in the program's frame; both REXX and COBOL auto-inject the names that aren't already declared.

Field	Set by	Meaning
EIBAID	SEND MAP / RECEIVE MAP	Single-byte AID character of the most recent map response. Use C2X(EIBAID) (REXX) or compare with X'F3' (COBOL) to detect PF/PA/CLEAR/ENTER.
EIBCPOSN	SEND MAP / RECEIVE MAP	1-based cursor position of the most recent map response.
EIBCALEN	dispatcher (per- task entry)	Length of DFHCOMMAREA flowed in from the prior task or LINK caller. Zero when no COMMAREA was passed.
EIBTRMID	dispatcher	This terminal's id. Same as ASSIGN TERMID(...).
EIBRESP	every EXEC CICS	The response code of the most recent command (see Chapter 12).
EIBRESP2	every EXEC CICS	The secondary response code of the most recent command (currently always 0).
RC	every EXEC CICS	Mirror of EIBRESP, for REXX IF RC <> 0 style.
DFHCOMMAREA	dispatcher / LINK / RETURN	The COMMAREA bytes; in COBOL, auto-injected as PIC X(2000) if not declared.

REXX

EIBAID and friends are ordinary REXX variables that the handlers write through the `cics.Frame` interface. A REXX program tests them exactly like any other variable:

```
IF EIBRESP <> 0 THEN SAY 'Command failed, RC=' EIBRESP
IF C2X(EIBAID) = 'F3' THEN EXEC CICS RETURN END-EXEC
```

SIGNAL ON ERROR ([Chapter 18](#)) is the alternative to per-call IF EIBRESP <>.

COBOL

The same names are auto-injected as PIC items if the program doesn't declare them (`cobol.ensureSystemItems`). Test them with ordinary IF clauses:

```

IF EIBRESP NOT = 0
  DISPLAY 'Command failed'
END-IF.

IF EIBAID = X'F3'
  EXEC CICS RETURN END-EXEC
END-IF.

```

There is no COBOL equivalent of `SIGNAL ON ERROR` yet.

Chapter 12. Response codes

After every `EXEC CICS` command the handler writes `EIBRESP`, `EIBRESP2`, and `RC` into the program's frame. The full set of constants is in `cics/resp.go`; the values a typical bricks program tests for are:

Constant	Value	Meaning
NORMAL	0	Success.
ERROR	1	Generic error.
EOC	6	End-of-chain (second <code>RECEIVE</code> in same task).
NOTFND	13	Record / key not found.
DUPREC	14	<code>WRITE</code> collided with an existing key.
DUPKEY	15	Duplicate key on a non-unique index (reserved).
INVREQ	16	Invalid request: bad option combination, missing data store, invalid name.
IOERR	17	Underlying store / network IO failed.
NOSPACE	18	Out of storage (reserved).
NOTOPEN	19	File not open (reserved).
ENDFILE	20	Browse walked past the last key (or out of <code>GENERIC</code> prefix).
LENGERR	22	Length mismatch.
QZERO	23	Queue is empty (reserved).
ITEMERR	26	TS item out of range; typical at end-of-queue.
PGMIDERR	27	<code>LINK / XCTL</code> target not in <code>transactions.conf</code> , or sub-program errored, or caller is not authorised for the target.
MAPFAIL	36	<code>SEND MAP / RECEIVE MAP</code> could not find or render the named map.
QIDERR	44	TS queue id invalid or unknown.

Idiomatic error handling

REXX, per-call test:

```

EXEC CICS READ FILE('CUSTOMERS') INTO(REC) RIDFLD(K) END-EXEC
SELECT
  WHEN EIBRESP = 0 THEN NOP
  WHEN EIBRESP = 13 THEN MSG = 'Customer' K 'not found'
  OTHERWISE          MSG = 'I/O error, RC=' EIBRESP
END

```

REXX, signal-on:

```

SIGNAL ON ERROR NAME CICSERR
EXEC CICS READ FILE('CUSTOMERS') INTO(REC) RIDFLD(K) END-EXEC
...
EXIT
CICSERR:
  SAY 'CICS error at line' SIGL ', RC=' RC
EXIT

```

COBOL:

```

EXEC CICS READ FILE('CUSTOMERS') INTO(REC) RIDFLD(CKEY) END-EXEC
EVALUATE EIBRESP
  WHEN 0    CONTINUE
  WHEN 13   MOVE 'Customer not found' TO MSG
  WHEN OTHER MOVE 'I/O error' TO MSG
END-EVALUATE.

```

Chapter 13. Commands not implemented

The following commands are not implemented and the parser does not recognize them at all:

Command	Use
GETMAIN, FREEMAIN	Dynamic storage. REXX has dynamic variables; COBOL has WORKING-STORAGE.
LOAD, RELEASE	Explicit program-fetch lifecycle. Bricks caches program text on first dispatch.
WAIT EVENT, POST	ECB-driven inter-task signalling — bricks's scheduler model does not expose ECBs.

Recently added in 2.7.x. QUERY SECURITY and VERIFY PASSWORD both graduated from this deferred list — programs can now ask the four-axis authorisation question and re-authenticate a userid / password pair through the verbs described in [Chapter 6 — Security inquiry](#).

DELAY

```

EXEC CICS DELAY
  {INTERVAL(hhmmss) | TIME(hhmmss) |
   FOR HOURS(h) MINUTES(m) SECONDS(s) MILLISECS(ms) |
   UNTIL HOURS(h) MINUTES(m) SECONDS(s)}
  [RESP(rc) RESP2(rc2)]
END-EXEC.

```

Suspends the issuing task until the given duration elapses (INTERVAL / FOR) or until the named time-of-day arrives (TIME / UNTIL). A bare DELAY with no clauses is a NORMAL yield. INTERVAL and TIME accept the same HHMMSS digit forms as START. FOR / UNTIL are the IBM-canonical compositional alternatives — supply any subset of HOURS, MINUTES, SECONDS, MILLISECS; the millisecond clause is the only sub-second form available.

Bricks sleeps the calling goroutine; PA1 break-out continues to work because the connection's prompt loop is unaffected. A long DELAY on a terminal that disconnects mid-sleep ends when the conn read deadline poisons the next dispatch — the task is not held alive past the connection.

CANCEL

```
EXEC CICS CANCEL
      {REQID(handle) | TRANSID(name) [TERMID(t)]}
      [RESP(rc) RESP2(rc2)]
END-EXEC.
```

Drops a pending START from the in-region scheduler. REQID(handle) matches against the 1–8-character handle the originating START supplied; bricks records the handle on the queued entry. TRANSID(name) drops the soonest pending START to that transid on TERMID (default = the issuing terminal). RESP=NOTFND (13) when no match.

The scheduler is in-memory; a bricks restart drops every pending entry, after which any CANCEL returns NOTFND.

ENQ / DEQ

```
EXEC CICS ENQ RESOURCE(name) [LENGTH(n)] [HOLD] [NOSUSPEND]
      [RESP(rc) RESP2(rc2)]
END-EXEC.

EXEC CICS DEQ RESOURCE(name) [LENGTH(n)]
      [RESP(rc) RESP2(rc2)]
END-EXEC.
```

Named-resource locking with region scope. ENQ blocks until the resource is acquired; with NOSUSPEND, an already-held resource returns RESP=ENQBUSY (55) immediately. A re-entrant ENQ by the same task on the same resource is a NORMAL no-op. The default lock scope is the next SYNCPOINT (commit or rollback) or task end; HOLD extends the scope past SYNCPOINT so the lock persists until an explicit DEQ or task end.

LENGTH(n) is accepted for source compatibility but bricks uses the entire resolved RESOURCE value as the key (uppercased, trimmed).

A non-NOSUSPEND ENQ that has waited five minutes for the holder to release returns ENQBUSY rather than blocking indefinitely. That cap converts a dropped-owner edge case into a clean recoverable status; legitimate same-region holders are never close to that limit.

Operator visibility — CEMT INQUIRE ENQ

CEMT INQUIRE ENQ (letter E on the INQUIRE menu) pages every enqueue currently held and every task suspended behind one. It answers the question a hung terminal raises: who holds the lock, and for how long.

Column	Meaning
RESOURCE	The uppercased, trimmed RESOURCE key. Names longer than 24 characters are truncated with a trailing +; the head is kept so the value still matches the ENQ RESOURCE(...) literal in your source.
OWNER	The owning task — the task number, else TERM:<termid> for a lock taken outside a task, else ANON.
TRANID / TERM	Joined from the live task table. A lock that outlives its task shows – rather than a stale name.
STATE	OWNED or WAITING. One OWNED row heads each resource, followed by its waiters, longest-waiting first.
HOLD	YES when the lock was taken with HOLD and so survives SYNCPOINT — usually the answer to "why is this lock still here?". – on a waiting row.

Column	Meaning
AGE	How long this task has held the lock, or has been waiting for it. A re-entrant ENQ does not reset it.

The panel is read-only and open to every operator, not just ADMIN — an enqueue is a resource, not internal control-block detail, and every column is already visible elsewhere via INQUIRE TRANSACTION and INQUIRE TERMINAL. ENTER re-snapshots, PF7 / PF8 page, PF3 exits.

On an idle region the panel shows No enqueues held. — and that screen is live too: ENTER re-snapshots it, so an operator can sit on it and watch for a lock to appear, at which point the table takes over. PF3 / CLEAR backs out.

IBM's ENQTYPE and ENQSCOPE columns are deliberately omitted rather than filled with a constant: bricks has only EXECENQ enqueues, and its locks are process-scoped, so a blank ENQSCOPE (which means region scope in real CICS) would assert something bricks cannot guarantee. RETAINED likewise never appears as a STATE, because there are no shunted units of work to retain a lock.

Because the snapshot is retaken on every ENTER, a refresh landing between a holder's DEQ and the woken waiter's acquisition can show a lone WAITING row with no OWNED row above it. That window is real and brief, not a rendering fault.

WRITE OPERATOR

```
EXEC CICS WRITE OPERATOR
      TEXT(area) [LENGTH(n)]
      [ROUTECODES(codes)] [NUMROUTES(n)]
      [EVENTUAL | ACTION | CRITICAL | IMMEDIATE]
      [CONSNAME(name)]
      [RESP(rc) RESP2(rc2)]
END-EXEC.
```

Writes TEXT(area) to the bricks console with the SYS tag. The issuing TERMID is included so the operator can correlate the message with the task that wrote it. LENGTH(n) bounds the bytes written. ROUTECODES, NUMROUTES, CONSNAME, EVENTUAL, ACTION, CRITICAL, and IMMEDIATE are accepted for source compatibility but fold into a single console write — bricks has no multi-console fan-out.

The REPLY form (REPLY(var) MAXLENGTH(n) REPLYLENGTH(var) [TIMEOUT(n)]) is rejected with RESP=INVREQ. Bricks has no operator-console input port to source the reply from, so accepting it silently would hang the task. A program that needs operator acknowledgement must reach the operator through SEND MAP on a configured operator terminal.

SEND CONTROL

```
EXEC CICS SEND CONTROL [ERASE | ERASEAUP]
      [FREEKB] [ALARM] [FRSET]
      [CURSOR(pos)] [PRINT] [LDC(name)] [STRFIELD]
      [RESP(rc) RESP2(rc2)]
END-EXEC.
```

3270 control order without a data payload — used between a SEND MAP DATAONLY pre-paint and the next operator interaction to clear the screen, unlock the keyboard, or reset MDTs without re-painting any field content. ERASE and ERASEAUP clear the screen; the other flags are accepted but their wire-bits are not tunnelled by the underlying go3270 layer (FREEKB and FRSET are part of every paint's fixed WCC, so they

have no separate effect; ALARM / CURSOR / PRINT / STRFIELD / LDC are silently accepted so program source compiles). A flag-less SEND CONTROL is a NORMAL no-op.

Recently added

- CONVERSE — see [Chapter 4](#). Syntactic sugar for SEND MAP + RECEIVE MAP.
 - START / RETRIEVE — see Chapter 5. Same-terminal scheduling with payload pass-through. Cross-terminal and no-TERMID (headless) STARTs are still deferred; the v1 surface rejects them with RESP-INVREQ and a clear message.
 - DELAY, CANCEL, ENQ / DEQ, WRITE OPERATOR, SEND CONTROL — round-4 verbs documented above.
 - WEB * (server-side, Phase 1) — see "EXEC CICS WEB" below. The client-side surface (WEB OPEN / SEND / RECEIVE / CONVERSE / CLOSE) and the DOCUMENT API land in Phase 2; URIMAP / TLS / mTLS in Phase 3.
-
-

Part 7. EXEC CICS WEB command reference

Part 7 begins here. Bricks ships a complete EXEC CICS WEB family — both the **inbound** (server-side, the WAPI listener) and the **outbound** (client-side, programs calling external HTTP services). The DOCUMENT API for assembling response bodies from many fragments is documented here too, alongside the URIMAP catalogue, inbound mTLS, the CONVERSE WEB alias, and the dedicated copybooks.

EXEC CICS WEB — server side (Phase 1)

bricks ships an inbound HTTP listener that turns each matched request into a transaction dispatch via the EXEC CICS WEB * verbs. Enable it in bricks.cnf with enable_wapi=yes. Define routes in runtime/web_routes.conf — the URIMAP layer; the matched TRANSID is looked up in transactions.conf exactly as a 3270 operator's typed dispatch would. Both files reload on mtime change, so adding a new route requires no restart.

#	method	path-pattern	transid	[groups]	[response_timeout]
	GET	/api/customer/{id}	WAPI	public	
	POST	/api/customer	WAPC	admin	
	GET	/api/orders/{id}	OAPI	users	
	DELETE	/admin/cache/{key}	CACR	admin	5s

The matching transactions.conf rows route TRANSID → program (REXX or COBOL) — same format used by 3270 dispatch:

```
WAPI:rexx:wapi.rexx:public,users,admin
WAPC:cobol:wapic.cob:admin
OAPI:rexx:orders.rexx:users,admin
CACR:cobol:cacheclr.cob:admin
```

A single program is therefore reachable from **both** the 3270 prompt (the operator types its TRANSID) and the HTTP endpoint (a client hits the matching URL). The transaction distinguishes the front door by which family of verbs it calls — EXEC CICS WEB * returns INVREQ on a 3270 task, and EXEC CICS SEND MAP is meaningless on a web task — but most programs naturally use one set or the other.

A {name} placeholder in the path is exposed to the transaction via WEB READ QUERYPARM('name'), sharing one namespace with the inbound ?name=... query string (path captures win on conflict).

A path capture is percent-decoded before the program sees it. The request path is matched in its *escaped* form, so a percent-encoded / can never split a segment and fool a single-segment placeholder — but the captured value is then decoded, because it shares a namespace with the already-decoded query string. GET /api/customer/Smith%20John against /api/customer/{id} therefore gives WEB READ QUERYPARM('id') the value Smith John, the same text ?id=Smith%20John would have produced. (Before this, the capture arrived raw, so one program could see id="Smith%20John" next to q="Smith John" for the very same name and the record lookup found nothing.) A capture whose escape sequence is malformed makes the route not match, and the request 404s.

Authentication — the same `users.conf` that CSSN uses

A web client identifies itself using **HTTP Basic authentication**, and bricks verifies the credentials against the exact same `users.conf` and bcrypt hash that the CSSN sign-on transaction uses on the 3270. The credential carrier is different (an HTTP header instead of a 3270 password field); everything downstream is identical.

Layer	3270 / CSSN	WAPI / HTTP
Credential store	<code>runtime/users.conf</code>	<code>runtime/users.conf</code> (same file)
Hashing	bcrypt	bcrypt
Verification call	<code>auth.FileStore.Authenticate(user, pass)</code>	<code>auth.FileStore.Authenticate(user, pass)</code> (same call)
Credential carrier	Hidden PASSWORD field on the CSSN map, posted by SEND→RECEIVE	Authorization: Basic <code>base64(user:pass)</code> header
Operator-visible state	<code>sess.Authenticated=true</code> ; UCB attached for the whole 3270 session until CSSF LOGOFF	One-shot per request — <code>sess.Authenticated=true</code> for the dispatch only
What the transaction sees	EXEC CICS ASSIGN USERID(...) returns the operator	Same — EXEC CICS ASSIGN USERID(...) returns the bcrypt-verified user

The dispatch flow per request:

1. **No Authorization header** — if the route's groups column contains the literal token `public`, the request is accepted as anonymous (`UserID=""`, `Authenticated=false`); otherwise the response is **401 Unauthorized** with `WWW-Authenticate: Basic realm="bricks"` so a browser shows the standard credential prompt.
2. `Authorization: Basic base64(user:pass)` — the credential is base64-decoded, split on `:`, then `auth.FileStore.Authenticate` runs (same path CSSN uses), comparing the password **exactly** — byte for byte, case included. On success the user's `users.conf` groups are matched against the route's groups: **any one** shared group admits the caller.
3. **Authenticated AND a group matches** — dispatch proceeds with `sess.Authenticated=true`, `sess.UserID=<username>`, `sess.Groups=<users.conf groups>`. The dispatcher's `IsAllowed` gate runs again against the TRANSID's `transactions.conf` groups, so **both** layers must permit the caller before the program runs.
4. **Authenticated but no group match** — **403 Forbidden**.
5. **Authentication fails** (wrong password, unknown user, malformed header) — **401 Unauthorized**.

ACL policy is **opt-in**: a `web_routes.conf` row with no groups column denies every request, log warning at boot. Add the literal `public` to permit anonymous access; add specific group names (`users`, `admin`, etc.) to require Basic-Auth against `users.conf` for a user belonging to that group.

Two properties of that column are easy to get wrong:

- **The list is a UNION, not an intersection.** `admin,users` admits anyone in *either* group, and `public` anywhere in the list makes the whole route unauthenticated — `admin,public` is exactly as open as `public` alone. There is **no way to spell "public AND admin"**; if the route is meant to be restricted, leave `public` out. bricks warns at load when a row mixes them.
- **A row has at most 5 fields, and a 6th is a hard error.** It used to be silently discarded, so `GET /admin ADMIN public 5s admin` loaded as `Groups=[PUBLIC]` and served the `admin` route to

unauthenticated callers with a 200. Such a row is now refused (and the route 404s, which is the safe direction to fail in). Every group goes in **one** comma-separated field with no spaces.

A user `alice` listed in `users.conf` as `alice:<bcrypt-hash>:users,admin` can run the **same** TRANSID from both front doors with the same credential:

```
curl -u alice:alice-password http://localhost:8080/api/orders/42
# → 200 + JSON body, sess.UserID=alice

curl -u alice:wrong-password http://localhost:8080/api/orders/42
# → 401 Unauthorized

curl http://localhost:8080/api/orders/42
# → 401 Unauthorized (no Authorization header)
```

...and at a 3270 emulator she signs on with CSSN → `alice / alice-password` and types the same TRANSID. The program sees the same EIBTRMID / USERID / EIBCALEN set and serves both callers identically. Most programs don't even need to know which front door they came from.

What's intentionally NOT yet supported:

- **Bearer tokens / cookies / sessions** — every request is stateless and re-authenticates against `users.conf`. There is no equivalent of the 3270 "stay signed on until CSSF LOGOFF" state across HTTP requests. The token-mint endpoint and `Authorization: Bearer ...` flow land in Phase 3.
- **OAuth / JWT / OIDC** — Phase 3+ work.
- **Per-request authentication via the CSSN transaction itself** — CSSN is a 3270-only screen artefact; pointing a route at it doesn't produce the browser UX a credential prompt should have. The Phase 3 token endpoint will be the right shape for that.

Server-side verb set

The Phase 1 verb set:

Verb	Use
WEB EXTRACT METHOD(var) PATH(var) ...	Request metadata: METHOD / SCHEME / HOST / PORT / PATH / QUERYSTRING / HTTPVERSION / CLIENTADDR / SERVERADDR. Each option is the name of a target variable.
WEB READ HTTPHEADER(name) VALUE(var) [LENGTH(var)]	Read one inbound header. NOTFND when absent.
WEB STARTBROWSE HTTPHEADER / WEB READNEXT HTTPHEADER NAME(var) VALUE(var) [NAMELENGTH(var)] [VALUELENGTH(var)] / WEB ENDBROWSE HTTPHEADER	Walk every inbound header in sorted order. ENDFILE when exhausted.
WEB READ QUERYPARM(name) VALUE(var)	Read query parameter OR routing-table {name} capture. Both are percent-decoded — <code>%20</code> reaches the program as a space. NOTFND when absent.
WEB STARTBROWSE QUERYPARM / WEB READNEXT QUERYPARM ... / WEB ENDBROWSE QUERYPARM	Iterate every parameter / capture.
WEB READ FORMFIELD(name) VALUE(var)	Read one application/x-www-form-urlencoded field. Body parsed lazily on first call. A body that is not valid form encoding (a bad escape such as <code>%zz</code> , or a <code>;</code> separator) raises INVREQ, on every call for that request.

Verb	Use
WEB STARTBROWSE FORMFIELD / WEB READNEXT FORMFIELD ... / WEB ENDBROWSE FORMFIELD	Iterate form fields. STARTBROWSE raises INVREQ on a malformed body, as READ FORMFIELD does.
WEB RECEIVE INTO(var) [MAXLENGTH(n)] [LENGTH(var)] [TYPE(var)] [MEDIATYPE(var)]	Read the raw request body. LENGERR when over MAXLENGTH.
WEB WRITE HTTPHEADER(name) VALUE(value)	Set / append an outbound response header.
WEB SEND FROM(buf) [MEDIATYPE(s)] [STATUSCODE(n)]	Emit response status + body. First SEND wins for status/MEDIATYPE; subsequent SEND calls in the same task append to the body (matches CICS).
WEB PARSE URL URL(s) SCHEMENAME(var) HOST(var) PORT(var) PATH(var) QUERYSTRING(var) [HOSTLENGTH(var)] ...	Direction-agnostic URL splitter.
WEB CONVERTTIME DATESTRING(s) ABSTIME(var)	RFC 1123 / 850 / asctime → milliseconds since 1970-01-01 UTC.

Malformed form bodies are an ordinary condition. INVREQ from WEB READ FORMFIELD / WEB STARTBROWSE FORMFIELD on a body that will not decode is a plain condition: code RESP(...) and answer the client with a 400. (Earlier releases ended the task with an internal error even when RESP was coded.)

```
EXEC CICS WEB READ FORMFIELD('title') VALUE(WS-TITLE)
      RESP(WS-RESP) END-EXEC
IF WS-RESP = DFHRESP(INVREQ)
  PERFORM SEND-400
END-IF
```

The full REXX sample is `runtime/rexx/wapi.rexx`; in COBOL the helpful copybooks are DFHWBSC (status-code constants), DFHWBUH (common header-name literals), DFHWBMT (MIME-type literals), and DFHWBMETH (HTTP method literals).

COBOL server-side sample — `todo.cob` (transaction TODO)

`runtime/cobol/todo.cob` is the first server-side COBOL WAPI sample: a multi-user, multi-list to-do REST API. One program serves every `/api/todo/...` route in `runtime/web_routes.conf` (transaction row `TODO:cobol:todo.cob:PUBLIC`; needs `enable_wapi=yes`) and works out the request from the method plus the shape of the path (`EVALUATE WS-SHAPE ALSO WS-METHOD`). What it shows:

- **VSAM storage** in three KSDS files created on first WRITE — `TOD0USR` (key user), `TOD0LST` (user + list), `TOD0ITM` (user + list + item). Each user only ever sees their own data.
- **JSON both ways.** Every response object is built by `JSON GENERATE` from a small flat record (lower-case names via `NAME OF`, the item's done flag via `CONVERTING ... TO JSON BOOLEAN`); arrays are joined with `STRING ... WITH POINTER`. A JSON request body is read with `JSON PARSE`; form-encoded bodies still work through `WEB READ FORMFIELD`.
- **Units of work.** All input is validated before the first file update; an unexpected RESP after an update does `SYNCPPOINT ROLLBACK` and answers 500. A record lock collision answers 409.
- `EXIT PARAGRAPH` for early returns from validation paragraphs.
- **Auth** by `X-Todo-User` / `X-Todo-Key` headers (accounts are made with `POST /api/todo/users`); the routes themselves are public.

A browser front end lives in `runtime/web/todo/`. When the static web server is on (`webserver_port > 0`), the WAPI listener also serves the static doc root for paths no route matches, so the app opens at `http://host:`

<wapi_port>/todo/ on the same origin as the API. `scripts/todo_api_test.bash` [BASE] runs an end-to-end test against a running instance (needs `curl` and `jq`).

EXEC CICS WEB — client side (Phase 2a)

The same WEB verb family also drives **outbound** HTTP — a bricks transaction calling a remote HTTP service. No `bricks.cnf` toggle is required; bricks builds one shared `*http.Client` at startup (tuned by `web_client_timeout` / `web_client_max_idle_conns` / `web_client_tls_skip_verify`) and every task reuses it. The host's CA bundle validates outbound HTTPS certificates by default — point at any normal public HTTPS endpoint and it just works.

The client uses a per-task **session token**: WEB OPEN returns a 32-character hex token; subsequent WEB SEND / WEB RECEIVE / WEB CONVERSE / WEB CLOSE all thread that token via `SESSTOKEN(...)`. A token still open at task end is auto-closed (dispatcher defer, mirrors `CloseBrowsers` / `CloseAllTD`).

The Phase 2a verb set:

Verb	Use
WEB OPEN HOST(s) [PORT(n)] [SCHEME(s)] SESSTOKEN(var)	Open a logical client session. PORT defaults to 80 (HTTP) / 443 (HTTPS). SCHEME defaults to <code>http</code> . The token is written into the variable named by <code>SESSTOKEN</code> .
WEB CONVERSE SESSTOKEN(t) METHOD(s) PATH(s) [FROM(buf)] [LENGTH(n)] [QUERYSTRING(s)] [MEDIATYPE(s)] INTO(var) [STATUSCODE(var)] [STATUSTEXT(var)] [MEDIATYPE(var)] [TYPE(var)] [MAXLENGTH(n)] [LENGTH(var)]	One-shot send + receive. Builds the request, fires it through the shared client, writes the response body into <code>INTO</code> ; <code>LENGERR</code> when the body exceeds <code>MAXLENGTH</code> (default 1 MiB).
WEB SEND SESSTOKEN(t) METHOD(s) PATH(s) [FROM(buf)] [QUERYSTRING(s)] [MEDIATYPE(s)]	Stage a request on the session (no round-trip yet).
WEB RECEIVE SESSTOKEN(t) INTO(var) [STATUSCODE(var)] [STATUSTEXT(var)] [MEDIATYPE(var)] [MAXLENGTH(n)]	Fire the staged request; populate the named output variables from the response.
WEB CLOSE SESSTOKEN(t)	Release the session and the underlying response body.
WEB WRITE HTTPHEADER(name) VALUE(v) SESSTOKEN(t)	Set an outbound <i>request</i> header on the session. Repeated <code>WRITE</code> calls for the same name produce multi-valued headers (matches <code>http.Header.Add</code>). Must precede the next <code>SEND</code> / <code>CONVERSE</code> -- headers committed at request build time.
WEB READ HTTPHEADER(name) VALUE(var) [LENGTH(var)] SESSTOKEN(t)	Read one response header by name from the most-recent <code>SEND/CONVERSE/RECEIVE</code> on the session. Case-insensitive name match. <code>NOTFND</code> when absent; <code>INVREQ</code> when no response has landed yet.
WEB STARTBROWSE HTTPHEADER SESSTOKEN(t) / WEB READNEXT HTTPHEADER NAME(var) VALUE(var) [NAMELENGTH(var)] [VALUELENGTH(var)] SESSTOKEN(t) / WEB ENDBROWSE HTTPHEADER SESSTOKEN(t)	Iterate every response header in sorted order. <code>ENDFILE</code> when exhausted; <code>INVREQ</code> when no <code>STARTBROWSE</code> has run on the session.
WEB EXTRACT SESSTOKEN(t) [SCHEME(var)] [HOST(var)] [PORT(var)] [PATH(var)] [HOSTLENGTH(var)] [PORTNUMBER(var)]	Inspect the session's resolved endpoint <i>after</i> <code>OPEN</code> . Distinct from <code>WEB EXTRACT URIMAP(name)</code> (no token) and from the server-side request-meta form (no token, different option set).

Bricks issues Accept: */* and User-Agent: bricks-cics/1.0 on every outbound request by default; both can be overridden with `WEB WRITE HTTPHEADER('User-Agent') VALUE(...) SESSTOKEN(t)` before the next `SEND / CONVERSE` on that session. The `runtime/cobol/whdr.cob` sample (TRANSID WHDR) walks through the WRITE-then-READ flow end-to-end.

RESP codes: NORMAL, NOTFND (unknown SESSTOKEN), INVREQ (missing required option, web client not configured), IOERR (network / TLS / timeout), LENGERR (body over MAXLENGTH).

COBOL copybook for the client surface

`runtime/cobolcopy/DFHWBSI.cpy` carries the standard session-token + body buffer templates:

```
01 DFH-WB-SESS    PIC X(36).
01 DFH-WB-BODY    PIC X(8192).
01 DFH-WB-URL     PIC X(2048).
```

Pair it with the existing `DFHWPSC` (status codes) and `DFHWPMBTH` (HTTP method literals) so a CICS-style program reads idiomatically:

```
COPY DFHRESP.
COPY DFHWPSC.
COPY DFHWPMBTH.
COPY DFHWBSI.

EXEC CICS WEB OPEN HOST('api.github.com') PORT(443)
              SCHEME('HTTPS')
              SESSTOKEN(DFH-WB-SESS) END-EXEC.

EXEC CICS WEB CONVERSE SESSTOKEN(DFH-WB-SESS)
                      METHOD(WB-GET)
                      PATH('/zen')
                      INTO(DFH-WB-BODY)
                      STATUSCODE(STAT) END-EXEC.

EXEC CICS WEB CLOSE SESSTOKEN(DFH-WB-SESS) END-EXEC.

IF STAT = DFHRESP-WB-OK ...
```

URIMAP — symbolic outbound endpoints (Phase 3a)

A URIMAP entry in `runtime/web_routes.conf` defines a named upstream service that outbound transactions reference by symbol. Both layers of `web_routes.conf` (routes for inbound, URIMAPs for outbound) live in the same file and hot-reload together.

```
# Format: URIMAP NAME scheme://host[:port][/path-prefix]
URIMAP  GITHUB  https://api.github.com
URIMAP  WEATHER https://api.openweathermap.org/data/2.5
URIMAP  INTRANET https://api.internal:8443/v1
```

Names are 1..8 uppercase characters (CICS resource-name limit), drawn from A-Z, 0-9 and the national characters \$ @ # plus _. The charset is enforced at parse and at the CEDA URIMAP write path, so a name such as A;B — which used to load and then never match what a program asked for — is refused with URIMAP name: only A-Z 0-9 \$ @ # allowed. The path-prefix is optional; when set, it becomes the leading segment of every PATH the program later supplies on `WEB SEND / WEB CONVERSE`. So `URIMAP WEATHER ...data/2.5` paired with `WEB CONVERSE PATH('/weather?city=NYC')` issues against `https://api.openweathermap.org/data/2.5/weather?city=NYC`.

The URIMAP-aware verb shapes:

Verb	Use
WEB OPEN URIMAP(name) SESSTOKEN(var)	Open a client session against the named endpoint. SCHEME / HOST / PORT come from the URIMAP row; explicit HOST(. . .) / PORT(. . .) / SCHEME(. . .) on the same command override per-field (matches IBM CICS). NOTFND when the name is unknown.
WEB SEND URIMAP(name) METHOD(s) PATH(s) [FROM(buf)] [QUERYSTRING(s)] [MEDIATYPE(s)] INTO(var) [STATUSCODE(var)] [STATUSTEXT(var)] [LENGTH(var)] [MAXLENGTH(n)]	One-shot OPEN + CONVERSE + CLOSE: no SESSTOKEN required. Resolves the URIMAP, opens a fresh session, sends the request, populates INTO and the named output variables from the response, and releases the session. Use for fire-and-forget outbound calls where session-token plumbing has no value. (Phase 3b.)
WEB EXTRACT URIMAP(name) [SCHEME(var)] [HOST(var)] [PORT(var)] [PATH(var)] [HOSTLENGTH(var)] [PORTNUMBER(var)]	Inspect a URIMAP definition without opening a session. Used by programs that branch on the resolved endpoint or need to display it to an operator. NOTFND when the name is unknown.

URIMAP rows load even when enable_wapi=no — outbound-only deployments are supported without exposing any inbound listener.

The runtime/cobolcopy/DFHWEBUM.cpy copybook ships the conventional WORKING-STORAGE variables:

```
01 DFH-WB-UM-NAME      PIC X(8).
01 DFH-WB-UM-SCHEME    PIC X(8).
01 DFH-WB-UM-HOST      PIC X(64).
01 DFH-WB-UM-PORT      PIC X(5).
01 DFH-WB-UM-PATH      PIC X(128).
```

Outbound TLS verification (Phase 3a)

By default the outbound HTTPS client verifies upstream certificates against the host system's trust store — the OS's /etc/ssl/certs/... bundle on Linux, the Keychain on macOS, Windows trust store on Windows. Three bricks.cnf knobs adjust:

Knob	Use
web_client_ca_bundle	Path to a PEM file of trusted CAs. When set, bricks uses ONLY these for verification — the system trust store is ignored. Useful for internal CAs that aren't installed system-wide.
web_client_cert + web_client_key	PEM client-cert and key for outbound mTLS. Presented to upstream services that require the caller to authenticate with a certificate. Both must be set together; one without the other is a startup error.
web_client_tls_skip_verify	yes accepts ANY upstream certificate — test-only escape hatch for self-signed endpoints. Logs a loud startup WARNING; never use in production.

Operator visibility (Phase 3a)

- CEMT INQUIRE URIMAP — pages the URIMAP rows loaded from web_routes.conf (NAME / SCHEME / HOST / PORT / PATH-PREFIX).
- CEMT INQUIRE WEB — pages the inbound HTTP requests currently in dispatch (METHOD / URI / TERM / TRANSID / USER / AGE). When no requests are in flight, the screen still shows a one-row summary with the lifetime request / error totals.

Phase 2a sample — wzen.cob

runtime/cobol/wzen.cob (TRANSID WZEN) is the shipped end-to-end demonstration. It fetches GitHub's /zen endpoint over HTTPS — a single-line plain-text response, no JSON parsing needed — and renders the result on the WZEN1 3270 map. The program demonstrates the full client lifecycle, branches cleanly on each EIBRESP, and surfaces network / HTTP failures with a distinct error path so the operator sees what went wrong on the 3270 screen rather than getting an opaque abend.

Drive it from a 3270 emulator: sign on (CSSN → admin), type WZEN, press ENTER. The first line under the header shows the actual URL fetched (GET https://api.github.com/zen) so the demo is self-documenting; the body appears on row 8.

Phase 2b sample — whdr.cob (custom request + response headers)

runtime/cobol/whdr.cob (TRANSID WHDR) is the canonical end-to-end demonstration of WEB WRITE HTTPHEADER and WEB READ HTTPHEADER on a client session. It opens the GITHUB URIMAP, sets two outbound headers (User-Agent: bricks-whdr/1.0, Accept: text/plain), issues GET /zen, then reads three response headers — two that always land on a GitHub response (Content-Type, Server) and one deliberately-absent header to exercise the NOTFND branch. The 3270 screen shows the request headers sent, the response headers received, and the body excerpt.

Drive it the same way as WZEN: sign on with CSSN, type WHDR, press ENTER. Pair the screen with CEMT INQUIRE WEB in a second 3270 session to watch the per-route counters move each time you rerun the transaction.

EXEC CICS DOCUMENT — chunked body builder

The DOCUMENT API is the buffer-builder that pairs with WEB SEND DOCTOKEN(t) for programs assembling large responses from many small fragments. Classic CICS pattern for emitting HTML / XML / JSON in pieces; the bricks implementation also serves as the chunk store for WEB RETRIEVE DOCTOKEN(var) when programs want the inbound body as a document handle instead of a flat INTO(buf).

Verb	Use
DOCUMENT CREATE DOCTOKEN(var) [SYMBOLLIST(s)] [LISTLENGTH(n)] [DELIMITER(s)]	Allocate a fresh empty document; return its token in the named variable. Optional SYMBOLLIST seeds a name=value&... symbol table that subsequent INSERT SYMBOL(...) substitutes from.
DOCUMENT INSERT DOCTOKEN(t) [FROM(buf) [LENGTH(n)] TEXT(buf) BINARY(buf) SYMBOL(name) DOCUMENT(other-token)] [AT(position)]	Append (or splice with AT(position)) one chunk to the document. SYMBOL(name) substitutes from the bound symbol list (NOTFND when unknown); DOCUMENT(other-token) splices another document verbatim (NOTFND when unknown); TEXT and BINARY are codepage-aware in real CICS but collapse to FROM in bricks (single-codepage). LENGERR when the resulting buffer would exceed 4 MiB.
DOCUMENT SET DOCTOKEN(t) SYMBOLLIST(s) [LISTLENGTH(n)] [DELIMITER(d)] [UNESCAPED]	(Re-)bind the symbol list a subsequent INSERT SYMBOL substitutes from. Pass an empty SYMBOLLIST to clear.
DOCUMENT RETRIEVE DOCTOKEN(t) INTO(buf) LENGTH(var) [CHARACTERSET(s)] [DATAONLY HOSTCODEPAGE]	Read the assembled body back into INTO; LENGTH receives the byte count. CHARACTERSET / DATAONLY / HOSTCODEPAGE are accepted but ignored (bricks is single-codepage ASCII).
DOCUMENT DELETE DOCTOKEN(t)	Release the document. NOTFND when the token isn't known. Forgotten documents are released at task end automatically.

Verb	Use
WEB SEND DOCTOKEN(t) [STATUSCODE(n)] [MEDIATYPE(s)]	Emit the named document as the outbound response body. The DOCUMENT INSERT chain has already built it; SEND just splices the bytes onto BodyOut. NOTFND when the token isn't known.
WEB RETRIEVE DOCTOKEN(var)	Inbound-body-as-document. Allocates a fresh document, copies the request body into it, and returns the token. Alternative to WEB RECEIVE INTO(buf) for programs that prefer the document API.

The standard COBOL declaration:

```

COPY DFHDCDOC.    *> DFH-DOC-TOK (PIC X(36)) + delimiter constants

EXEC CICS DOCUMENT CREATE DOCTOKEN(DFH-DOC-TOK) END-EXEC.
EXEC CICS DOCUMENT INSERT DOCTOKEN(DFH-DOC-TOK)
                        FROM ('<html><body>') END-EXEC.
EXEC CICS DOCUMENT INSERT DOCTOKEN(DFH-DOC-TOK)
                        FROM (BODY-CHUNK) END-EXEC.
EXEC CICS DOCUMENT INSERT DOCTOKEN(DFH-DOC-TOK)
                        FROM ('</body></html>') END-EXEC.
EXEC CICS WEB SEND DOCTOKEN(DFH-DOC-TOK)
                        MEDIATYPE('text/html')
                        STATUSCODE(200) END-EXEC.

```

RESP codes: NORMAL, NOTFND (unknown DOCTOKEN, unknown SYMBOL, unknown spliced DOCUMENT), INVREQ (missing required option), LENGERR (INSERT would exceed 4 MiB), I0ERR (token generation failed).

EXEC CICS WEB — Phase 3b additions

Phase 3b closes out the master Phase-3 surface. Each addition slots onto the foundations laid by 3a (URIMAP catalogue, active- request tracking, outbound TLS extras) without breaking existing verb shapes.

CONVERSE WEB — **alias for** WEB CONVERSE

Some shops write the keywords in either order; bricks normalises both at the parser, so the two forms below dispatch identical:

```

EXEC CICS WEB CONVERSE SESSTOKEN(T) METHOD('GET')
                        PATH('/x') INTO(B) END-EXEC.
EXEC CICS CONVERSE WEB SESSTOKEN(T) METHOD('GET')
                        PATH('/x') INTO(B) END-EXEC.

```

WEB EXTRACT TCIPSERVICE — **inspect the receiving listener**

When WEB EXTRACT carries the TCIPSERVICE keyword (or the legacy EXTRACT TCPIP form without WEB), the verb reports the listener that received the inbound request:

Output	Source
TCIPSERVICE(var)	Symbolic listener name — WAPI for the plain port, WAPITLS for the TLS port.
PORTNUMBER(var)	Numeric port the listener is bound to.
IPADDRESS(var)	The server's bind address (mirrors bricks.cnf::dns_name).
CLIENT(var)	Caller's IP address (mirrors WEB EXTRACT CLIENTADDR).
AUTHENTICATE(var)	BASIC when the request carried Authorization: Basic ..., blank otherwise.

The legacy form drops the `WEB` keyword entirely — the parser recognises `EXTRACT TCPIP CLIENT(...)` `AUTHENTICATE(...)` and routes it through the same handler. Useful for programs ported from pre-CICS-TS web-services code.

Inbound mTLS + `WEB EXTRACT CERTIFICATE`

Set `web_inbound_client_ca=ca.pem` to require every TLS client to present a certificate signed by `ca.pem`. The TLS handshake rejects unauthorised clients before any verb runs — no transaction is dispatched at all. Inside the dispatched program, `WEB EXTRACT CERTIFICATE` reads the verified peer cert:

Output	Source on the peer cert
<code>COMMONNAME(var)</code>	Subject CN.
<code>ORGANISATION(var)</code>	Subject O.
<code>COUNTRY(var)</code>	Subject C.
<code>SERIALNUM(var)</code>	Hex serial number.
<code>ISSUER(var)</code>	Issuer CN.
<code>ISSUERORG(var)</code>	Issuer O.

`EIBRESP = NOTFND` on plain (non-TLS) requests and on TLS requests that carried no client cert — so one `TRANSID` can serve both authenticated and anonymous callers and branch on the resp.

```
COPY DFHRESP.
COPY DFHWBCC.    *> CN / ORG / CO / SERIAL / ISSUER buffers
EXEC CICS WEB EXTRACT CERTIFICATE
                COMMONNAME(DFH-WB-CN)
                ORGANISATION(DFH-WB-ORG)
                COUNTRY(DFH-WB-CO)
                SERIALNUM(DFH-WB-SERIAL)
                ISSUER(DFH-WB-ISSUER)

END-EXEC.
EVALUATE EIBRESP
  WHEN DFHRESP(NORMAL)
    *> Verified cert -- use the fields.
  WHEN DFHRESP(NOTFND)
    *> Plain HTTP or TLS without a client cert.
END-EVALUATE.
```

HTTP/2 on the inbound TLS listener

Go's `net/http` server auto-negotiates HTTP/2 over the TLS listener via ALPN whenever `tls.Config.NextProtos` is left unset — which bricks does on the WAPI TLS port. No knob is required; clients that advertise `h2` in their ALPN list get HTTP/2, clients that don't fall back to HTTP/1.1. Verify with `curl --http2 -v https://localhost:443/...` and look for ALPN, server accepted to use `h2` in the trace output. Cleartext HTTP/2 (`h2c`) is not enabled — the plain port stays on HTTP/1.1 to avoid a transitive dependency on golang.org/x/net/http2/h2c.

CEDA URIMAP — live edit the URIMAP catalogue

The Phase 3a `CEMT INQUIRE URIMAP` browser becomes operator-editable in 3b. Type `CEDA URIMAP` (or `CEDA M`) from the 3270 prompt for the list view:

S	NAME	SCHEME	HOST	PORT	PATH
—	GITHUB	https	api.github.com	443	—
—	WEATHER	https	api.openweathermap.org	443	/data/2.5
—	INTRANET	https	api.internal	8443	/v1

- Type **A** in the selector column for **Alter** — opens a form pre-loaded with the row's endpoint URL so a one-character edit is enough to repoint the URIMAP.
- Type **D** in the selector column for **Delete** — a Y/N confirmation overlay opens; the row is removed only when the operator types **Y**.
- Press **PF6** for **Define new** — opens an empty form for the NAME (1..8 chars) and ENDPOINT URL (scheme://host[:port] [/path-prefix]).

Every committed change is written atomically through a temp + rename to `web_routes.conf`, with an mtime guard that refuses the write if the file was modified externally between when the screen loaded and when the operator pressed PF5. Each successful mutation logs an audit line in the bricks log:

```
ceda=URIMAP op=DEFINE user=ADMIN term=T0001 target=NEWAPI detail=https://internal.svc:8443/v2
```

so the audit trail of who changed which row when is greppable.

The new in-process surface that backs the screen lives on `*web.Table: AddURIMap, AlterURIMap, DeleteURIMap`, and `Mtime()` for the optimistic-concurrency guard. Programs in REXX or COBOL **do not call these directly** — they're the CEDA implementation, not part of the EXEC CICS surface.

Part 3. The REXX language

Part 3 begins here. This part describes the REXX dialect bricks accepts and the small preprocessor that rewrites EXEC CICS blocks into the same parsed command shape COBOL produces. Every command in **Part 4** (EXEC CICS), **Part 6** (EXEC SQL) and **Part 7** (EXEC CICS WEB) is available verbatim from REXX with identical semantics.

Chapter 14. REXX program structure

A bricks REXX program is a flat sequence of statements with optional labels and procedures. Execution begins at the first statement; the program ends when execution falls off the bottom, or hits EXIT. A program file is plain UTF-8 (ASCII printable + tabs + LF only is recommended for terminal output). There is no fixed column discipline — newlines and whitespace are token separators.

```
/* HEL0 — minimum REXX program */
ADDRESS CICS
EXEC CICS SEND MAP('HEL01') ERASE END-EXEC
EXEC CICS RETURN END-EXEC
```

Comments

Block comments /* ... */ may span multiple lines; nesting is not supported. A REXX program traditionally opens with a one-line header comment so the operator gets a description in CEDA / CEMT when applicable.

Statements

Statements are separated by newlines or ;. The parser does not require a terminator after the last statement of a block. Labels have the form NAME: and may appear at any statement position.

String literals

A string is written in single or double quotes ('...' / "..."); a doubled quote inside is an escaped quote ('it''s'). A quoted string immediately followed by an X (hexadecimal) or B (binary) suffix — with no intervening blank — is a **hex / binary literal** whose value is the decoded bytes:

Literal	Value (bytes)
'F3'X	one byte 0xF3
'F1F2'X	0xF1 0xF2
'1'X	0x01 (an odd hex-digit count is left-padded with 0)
'0110'B	0x06 (left-padded to a full byte, then bit-packed)
''X / ''B	the empty string

The suffix is recognised only when the X / B is a standalone letter: '03'XY is the string '03' followed by the symbol XY (juxtaposition), *not* a hex literal. Invalid digits ('GG'X, '2'B) raise a SYNTAX error. Decoded literals are ordinary byte strings — C2X('F3'X) is 'F3', and X2C / D2C / B2X round-trip them.

Procedures

A label followed by `PROCEDURE [EXPOSE list]` defines a procedure with its own variable scope. `EXPOSE` re-routes named variables to the caller's frame recursively. Both simple names and stem roots (trailing dot) are accepted in the `EXPOSE` list. A procedure is dispatched by name from both the `CALL` statement and the `NAME(args)` function form; a `PROCEDURE` whose name matches a built-in shadows it (internal routines are resolved first, per standard REXX).

```
CALL GREET 'Alice'
EXIT

GREET: PROCEDURE EXPOSE LANG.
  PARSE ARG NAME
  SAY LANG.HELLO NAME
  RETURN
```

A procedure returns to its caller on `RETURN [expr]` or by falling off the bottom. The optional `expr` is assigned to the caller's special variable `RESULT`; a bare `RETURN` or a fall-off-end drops `RESULT` so the caller's existing value is not silently masked. Recursion is permitted (the only ceiling is the interpreter's `MaxSteps` runaway guard).

A procedure may be invoked either as a `CALL` statement or as a **function** in an expression — `Y = DOUBLE(21)`. In the function form the routine **MUST** `RETURN` a value; a bare `RETURN` or a fall-off-end raises a `SYNTAX` error ("function did not return data"). `RETURN ''` counts as data. (As a `CALL` statement, a value-less return is fine and simply leaves `RESULT` unset.)

The `NUMERIC` settings, the `ADDRESS` environment, and the `SIGNAL ON / CALL ON` condition traps are saved when a procedure is entered and restored on `RETURN`, so a routine may change them locally without leaking the change back to its caller.

```
SAY DOUBLE(21)           /* 42 – internal routine called as a function */
EXIT

DOUBLE: PROCEDURE
  PARSE ARG N
  RETURN N * 2
```

ADDRESS

`ADDRESS <env>` switches the active command handler. Bare strings inside an `ADDRESS` scope are commands routed to that handler; bricks ships a `CICS` handler. `ADDRESS` with no argument clears the active environment, after which a bare string is a syntax error.

```
ADDRESS CICS
"SEND MAP('CUST1') FROM(SCR.) ERASE"
"RETURN"
```

A command's return code is assigned to the special variable `RC`. A non-zero `RC` fires `SIGNAL ON ERROR` (if armed) — see [Chapter 18](#).

Special variables

Name	Set by	Cleared by
RC	Every command (bare string) executed under an ADDRESS env, every EXEC CICS / EXEC SQL.	Survives until the next command updates it.
RESULT	CALL fn ... when the called procedure does RETURN expr.	A bare RETURN, a fall-off-end, or DROP RESULT.
SIGL	Any direct SIGNAL and every condition trap (set to the source line of the failing statement).	Survives the trap; persists across normal flow.
SQLCODE / SQLSTATE / SQLERRMC	Every EXEC SQL verb.	Next EXEC SQL.

These names are otherwise ordinary writable variables; nothing stops a program from assigning them, though doing so is bad hygiene.

The EXEC CICS / EXEC SQL preprocessor

Bricks parses EXEC CICS verb operand-list END-EXEC (and the EXEC SQL form) as a single command at the same statement position a bare string would occupy under ADDRESS CICS. The two forms compile to the same dispatch path, so

```
EXEC CICS SEND MAP('HEL01') ERASE END-EXEC
```

and

```
ADDRESS CICS  
"SEND MAP('HEL01') ERASE"
```

are equivalent. The EXEC form is preferred in shipped programs because it doesn't depend on the current ADDRESS environment and survives an explicit ADDRESS switch.

Chapter 15. Variables and stems

Naming rules

- Names are case-insensitive; internally everything is uppercased before lookup.
- A name begins with a letter or `_` and may continue with letters, digits, `_`, or `!`, `?`, `#`, `@`, `$`. Stem-tail segments are separated by `.`.
- The literal RC, SIGL, RESULT, SQLCODE, SQLSTATE, SQLERRMC are reserved for runtime-managed values (see [Chapter 14 — Special variables](#)).

Simple variables and NOVALUE

An unset simple variable resolves to its own uppercased name — the classic REXX NOVALUE convention. SAY X when X has never been assigned prints X. If SIGNAL ON NOVALUE is armed (see [Chapter 18](#)) the reference traps instead of resolving to the name.

Assignment is `var = expr`; the right-hand side is evaluated once and stored verbatim (as a string — REXX has only one runtime type).

Stems

A stem is a variable name that ends in `..`. A bare assignment to the stem (`STEM. = value`) installs a default that any tail inherits when it has no explicit value. Per-tail values override the default:

```
STEM. = 'unset'
STEM.42 = 'forty-two'
SAY STEM.1    /* "unset"      -- default      */
SAY STEM.42   /* "forty-two" -- explicit tail */
```

Compound-symbol tail substitution

Non-numeric tail segments are resolved at every reference. With `J = 3`, the symbol `A.J` reads or writes `A.3`. The rules:

Tail segment shape	Behaviour
STEM.42 (numeric)	Literal — the digit run becomes the tail directly.
STEM.NAME where NAME is set	Substituted — value of NAME becomes the tail.
STEM.NAME where NAME is not set (NOVALUE)	Literal — the bare uppercased name becomes the tail.
STEM.I.J (multi-segment)	Each segment substituted independently — with <code>I=1</code> , <code>J=2</code> , references <code>STEM.1.2</code> .

The NOVALUE-stays-literal rule is the source of bricks's most common porting bug; see [Appendix B](#) for the standard pitfall.

DROP

`DROP name [name ...]` removes one or more variables from the current frame. A trailing dot on a stem name drops the default AND every tail — useful for resetting an accumulator between paginated reads:

```
DROP RECS.          /* clear the whole stem */
DROP A B C          /* drop three simples */
DROP STEM.42        /* drop just one tail */
```

PROCEDURE **and** EXPOSE

Inside a PROCEDURE, the parent frame's variables are invisible unless named on the EXPOSE list. EXPOSE accepts both simple names and stem roots:

```
GREET: PROCEDURE EXPOSE LANG. SUFFIX
...
```

`LANG.` exposes the parent's entire `LANG.` stem (default + all tails); `SUFFIX` exposes a single simple variable. Exposure is transitive — a procedure that exposes a name from a caller which also exposed it from its caller sees the outermost binding.

Chapter 16. Control flow

IF / THEN / ELSE

```
IF expr THEN stmt-or-block [ELSE stmt-or-block]
```

The THEN keyword is mandatory. Both arms accept either a single statement or a DO ... END block. Strict boolean evaluation: the condition must evaluate to 0 or 1; anything else raises SYNTAX (so IF X where X = 'foo' is an error, not a truthy test).

```
IF AID = 'F3' THEN DO
  EXEC CICS RETURN END-EXEC
  EXIT
END
ELSE
  CALL REFRESH
```

SELECT / WHEN / OTHERWISE / END

```
SELECT
  WHEN expr THEN stmt
  WHEN expr THEN stmt
  ...
  OTHERWISE stmt-or-block
END
```

WHEN arms are tested in order; the first match's body runs and the rest are skipped. OTHERWISE is optional but if no WHEN matched and OTHERWISE is absent, the runtime raises SYNTAX. Use OTHERWISE NOP for an explicit no-op default.

DO family

Form	Behaviour
DO ... END	Group several statements as one body (most useful inside IF / WHEN).
DO N ... END	Repeat the body N times. N is any expression that evaluates to a non-negative integer.
DO ctrl = start TO end [BY step] ... END	Counted loop. BY defaults to 1. The loop runs while ctrl is between start and end (inclusive). DO ctrl=1 TO 5 BY 0 is rejected at runtime (infinite-loop guard).
DO WHILE expr ... END	Pre-test loop — body runs while expr is truthy.
DO UNTIL expr ... END	Post-test loop — body runs at least once, then re-tests.
DO FOREVER ... END	Unbounded; exit only via LEAVE / RETURN / EXIT / SIGNAL.
DO ctrl OVER stem. ... END	Iterate ctrl over every set tail of stem.. Numeric tails are visited in numeric order, then non-numeric tails in lexicographic order. The default value (stem. = ...) is not visited.

The control variable in DO ctrl = / DO ctrl OVER is set on each iteration. LEAVE [ctrl] and ITERATE [ctrl] target either the innermost loop (no argument) or the loop whose control variable matches the named argument:

```
DO I = 1 TO 100
  DO J = 1 TO 100
    IF GRID.I.J = '*' THEN ITERATE I /* skip rest of inner */
  END
END
```

CALL / RETURN / EXIT

- **CALL** *name* [*arg1*, *arg2*, ...] — call a procedure or a built-in function in statement position. Arguments are expressions, evaluated left-to-right and made visible to the callee via `ARG(n)`. A **CALL** to a built-in throws away the returned value; a user procedure's **RETURN** *expr* populates `RESULT`. The same routine can also be called as a function in expression position (`x = name(args)`) — see [Chapter 14](#).
- **RETURN** [*expr*] — return from a procedure (or terminate a top-level program). When *expr* is present and the caller was **CALL**, the value lands in `RESULT`; without *expr*, `RESULT` is dropped.
- **EXIT** [*expr*] — terminate the entire program immediately. The optional *expr* becomes the process exit code (clamped to a byte by the host OS); without it the program exits 0.

SIGNAL

- **SIGNAL** *label* — non-local jump to a label. Active **DO** / **SELECT** / **PROCEDURE** frames are unwound first. `SIGL` is set to the source line of the **SIGNAL** statement.
- **SIGNAL ON** *cond* [**NAME** *label*] — arm a condition trap (see [Chapter 18](#)).
- **SIGNAL OFF** *cond* — disarm.

INTERPRET

INTERPRET *expr* evaluates *expr* to a string and executes that string as REXX source in the current frame. Labels and procedures inside the snippet are scoped to that one evaluation (they do not leak into the surrounding program).

NUMERIC

Form	Effect
NUMERIC DIGITS <i>n</i>	Sets the precision used to round arithmetic results. Default 9.
NUMERIC FUZZ <i>n</i>	Sets the tolerance applied to numeric comparisons. Default 0.
NUMERIC FORM SCIENTIFIC / NUMERIC FORM ENGINEERING	Accepted syntactically but bricks always renders numbers in plain decimal — neither form changes output today.

Arithmetic is `float64` internally, then rounded to **NUMERIC DIGITS** significant figures before storing. The `%` and `//` operators (integer divide and modulo) bypass that rounding.

NOP, PUSH, PULL, QUEUE, OPTIONS, TRACE

- **NOP** — explicit no-op (clearer than an empty `DO ... END` in a **WHEN**).
- **TRACE** [*opt*] — parsed and accepted but ignored; bricks does not run an interpretive tracer.
- **OPTIONS** ... — accepted and ignored (no recognised options).
- **PUSH** / **PULL** / **QUEUE** / **PARSE PULL** — accepted; the associated terminal data queue is always empty in bricks. Read forms return the empty string; write forms are silent no-ops. Use `EXEC CICS RECEIVE` for terminal input.

Bare strings

A bare string at statement position is a command sent to the current ADDRESS environment (see Chapter 14). Outside an active environment it is a syntax error.

```
ADDRESS CICS
"SEND MAP('CUST1') FROM(SCR.) ERASE"
"RETURN"
```

Chapter 17. PARSE templates

```
PARSE [UPPER] {VAR var | VALUE expr WITH | ARG | PULL} template
```

The source decides where the input string comes from; the template decides how it's chopped up and which pieces land in which variables.

Sources

Source	Input string
VAR var	Current value of the named variable. A compound name resolves its tail — <code>PARSE VAR airports.CODE ...</code> parses the value of <code>airports.<CODE></code> , not the literal name <code>AIRPORTS.CODE</code> .
VALUE expr WITH	The result of evaluating <code>expr</code> . The literal <code>WITH</code> keyword separates the source expression from the template.
ARG [, ARG, ...]	The arguments the program (or procedure) was called with. Each comma-separated segment of the template is one argument, in order — <code>PARSE ARG H, W</code> peels <code>ARG(1)</code> into the first sub-template and <code>ARG(2)</code> into the second.
PULL	The terminal-input queue, which in bricks is always empty. Use <code>EXEC CICS RECEIVE</code> for real input.

`PARSE UPPER` uppercases the source string before splitting it across the template — convenient for case-insensitive command parsing.

Template elements

Element	Effect
Bare variable name	In a <i>run</i> of bare variables, all but the last get one blank-delimited word each (leading blanks skipped); the last variable gets the verbatim remainder, preserving internal and trailing blanks. A <i>single</i> variable with no following anchor receives the whole segment verbatim (leading blanks included). The word separator is the space character only. A dynamic compound target (A.I) resolves its tail at assignment time.
. (period)	Placeholder — consumes one token without binding it. Useful as a "skip" in a variable run.
'literal'	String anchor. Matches the next occurrence of <code>literal</code> in the source and splits there. The literal itself is discarded.
n (positive integer)	Absolute column marker — jumps to column <code>n</code> (1-based) in the source. Variables before the marker get everything from the previous column up to column <code>n-1</code> .
+n / -n	Relative column marker — moves the cursor <code>n</code> characters forward or backward from the current position.

Worked examples

```
/* Whitespace tokens. TID = first word, CKEY = second word,
 * '.' discards the rest of the line. */
PARSE VAR LINE TID CKEY .

/* Pipe-delimited fields. */
PARSE VAR REC NM '|' AD '|' CY '|' PH

/* Fixed columns. */
PARSE VAR ROW 1 NAME 21 ADDR 51 PHONE

/* Multi-arg PARSE – peels two call arguments. */
GREET: PROCEDURE
  PARSE ARG NAME, GREETING
  SAY GREETING NAME
  RETURN

/* PARSE UPPER on user command, case-insensitive switch. */
PARSE UPPER VAR LINE VERB REST
SELECT
  WHEN VERB = 'LIST' THEN CALL LIST_HANDLER REST
  WHEN VERB = 'ADD' THEN CALL ADD_HANDLER REST
  OTHERWISE SAY 'unknown verb' VERB
END

/* Compound tail resolves on both sides. */
I = 2
A.2 = 'hello world'
PARSE VAR A.I W1 W2          /* W1='hello', W2='world' */

/* Whitespace preserved: the final variable keeps the run of
/* blanks between arg1 and arg2 (older PARSE collapsed them). */
LINE = 'CMD  arg1  arg2'
PARSE VAR LINE VERB REST    /* VERB='CMD', REST='arg1  arg2' */
```

Not supported

Feature	Status
PARSE LOWER	Not parsed (only UPPER is). Lowercase the source explicitly with LOWER(VAR) then PARSE VAR.
PARSE SOURCE	Not wired — returns empty.
PARSE VERSION	Not wired — returns empty.
PARSE NUMERIC	Not wired.
Regex / pattern matching in templates	REXX has none; use POS / LASTPOS + SUBSTR for ad-hoc splits.

Chapter 18. Conditions and SIGNAL ON

```
SIGNAL ON {ERROR | NOVALUE | SYNTAX | HALT} [NAME label]
SIGNAL OFF cond
SIGNAL label
```

SIGNAL ON cond [NAME label] arms a trap. The optional NAME clause picks an explicit handler label; without it, bricks defaults the handler label to the condition name (so SIGNAL ON ERROR jumps to a label named

ERROR:).

When a trapped condition fires:

- SIGL is set to the source line of the failing statement.
- Control transfers to the handler label, unwinding active DO / SELECT / PROCEDURE frames first.
- The trap remains armed for the next occurrence.

SIGNAL OFF cond disarms.

Recognised conditions

Condition	Fires on
ERROR	An EXEC CICS or EXEC SQL command returns a non-zero RC (i.e. the bare-string command path under ADDRESS CICS).
NOVALUE	A reference to a simple or compound variable that has never been assigned.
SYNTAX	Any other interpreter error — bad numeric, divide by zero, unknown function, out-of-range subscript, ...
HALT	External halt request. Parsed and reserved; bricks has no fan-in for an external halt today, so the condition never fires.

FAILURE, NOTREADY, and LOSTDIGITS from standard REXX are not recognised — SIGNAL ON FAILURE parses but never fires.

Example

```
SIGNAL ON ERROR NAME CICSERR
SIGNAL ON SYNTAX NAME OOPS

ADDRESS CICS
EXEC CICS READ FILE('CUSTOMERS') INTO(REC) RIDFLD(CKEY) END-EXEC
EXIT

CICSERR:
  SAY 'EXEC CICS failed at line' SIGL ', RC=' RC
  EXIT 12

OOPS:
  SAY 'Interpreter error at line' SIGL
  EXIT 16
```

Without an armed trap, the legacy "test EIBRESP after every verb" pattern still works. For per-condition (rather than blanket) traps, use EXEC CICS HANDLE CONDITION ([Chapter 10](#)); the two styles can coexist.

Chapter 19. Built-in functions

The bricks REXX interpreter ships with the functions listed below. A CALL to any of these in statement position throws away the return value; the same name in expression position returns the value. Built-in lookups are case-insensitive. Per standard REXX, a user-defined PROCEDURE of the same spelling **takes precedence** over the built-in (the internal routine shadows it) — for both the CALL statement and the NAME(args) function form.

Length / index

Signature	Returns
LENGTH(s)	Byte length of s.
POS(needle, hay [, start])	1-based byte position of the first occurrence of needle in hay starting at start (default 1), or 0 if not found.
LASTPOS(needle, hay [, end])	1-based byte position of the last occurrence at or before end (default = end-of-string), or 0.
WORDS(s)	Number of whitespace-delimited words in s.
WORDPOS(phrase, s [, start])	1-based word index where phrase first appears as a sub-sequence of words, or 0.
WORDINDEX(s, n)	1-based character position of the start of word n, or 0 if n is past the end.
WORDLENGTH(s, n)	Length of word n, or 0.
COUNTSTR(needle, hay)	Non-overlapping occurrence count.

Substring / words

Signature	Returns
SUBSTR(s, start [, length [, pad]])	Substring starting at 1-based start. With length, the result is exactly length characters (right-padded with pad, default space). Negative length raises SYNTAX.
LEFT(s, n [, pad])	Leftmost n characters, padded on the right with pad (default space).
RIGHT(s, n [, pad])	Rightmost n characters, padded on the left with pad.
SUBWORD(s, n [, length])	Words n through n+length-1 (default: to end). Single space between words in the result.
WORD(s, n)	Word n, or empty string.
DELSTR(s, n [, length])	s with length characters from position n removed (default: through end).
DELWORD(s, n [, length])	s with length words from word n removed.
INSERT(new, target [, n [, length [, pad]]])	Insert new into target after position n (default 0 = prepend), padded to length first.
OVERLAY(new, target [, n [, length [, pad]]])	Overwrite target at position n with new, extending target if needed.
CHANGESTR(needle, hay, repl [, count])	Replace every occurrence (or the first count) of needle with repl.

Whitespace / case

Signature	Returns
STRIP(s [, opt [, ch]])	Trim characters ch (default space) from L (left), T (right), or B (both, default).
SPACE(s [, n [, pad]])	Collapse internal whitespace; insert n copies of pad (default 1 space) between words. n=0 removes whitespace entirely.
CENTER(s, n [, pad]) / CENTRE(s, n [, pad])	Centre s in a field of width n; pad with pad (default space). Truncates equally from both ends when s is longer than n.
JUSTIFY(s, n [, pad])	Right-justify the last word in width n, padding earlier words.
UPPER(s) / LOWER(s)	ASCII case conversion.

Signature	Returns
REVERSE(s)	Byte-reverse s.
COPIES(s, n)	s repeated n times.
ABBREV(s, prefix [, minlen])	Returns 1 if prefix is a prefix of s with at least minlen characters; else 0.

Translation / bits

Signature	Returns
TRANSLATE(s [, out [, in [, pad]]])	Per-character map from in to out. With no out/in, uppercases. With only out, in defaults to X RANGE () (the full byte range), giving a per-position translation table.
VERIFY(s, ref [, opt [, start]])	1-based position of the first character in s (from start) that is not in ref (opt= 'N', default) or that IS in ref (opt= 'M'); 0 if none.
COMPARE(s1, s2 [, pad])	1-based position of the first differing byte after pad-equalising lengths with pad (default space); 0 if equal.
X RANGE([from [, to]])	All bytes from from (default '\x00') to to (default '\xFF') inclusive.
BITAND(s1 [, s2 [, pad]]) / BITOR(...) / BITXOR(...)	Bitwise logical operation; lengths equalised with pad.

Conversion

Signature	Returns
C2X(s)	Uppercase hex of every byte. Canonical way to compare EIBAID to a PF-key code: IF C2X(EIBAID) = 'F7' THEN ... (PF7).
X2C(hex)	Decode hex (spaces ignored) to a byte string.
D2X(n [, width])	Decimal integer to uppercase hex, optionally zero-padded to width.
X2D(hex [, n])	Hex to decimal. With n, treats the value as an n-digit signed integer (two's complement when the top bit is set).
C2D(s [, n])	Byte string to decimal. With n, signed two's complement over n bytes.
D2C(n [, len])	Integer to byte string, optionally len bytes wide.
B2X(bits)	Bit string (0/1 chars) to hex.
X2B(hex)	Hex to bit string.

Numeric

Signature	Returns
ABS(n)	Absolute value.
MAX(n, ...)	Maximum of all arguments.
MIN(n, ...)	Minimum of all arguments.
INT(n)	Integer part (truncate toward zero).
TRUNC(n [, places])	Truncate to places decimal positions (default 0).
MOD(a, b)	Modulo. Same value as a // b.
SIGN(n)	-1 if n<0, 0 if n=0, +1 if n>0.
FORMAT(n [, before [, after [, expx [, expt]]])	Numeric formatting with before characters in the integer part and after decimal places. expx / expt accepted syntactically but ignored.

Signature	Returns
DIGITS() / FUZZ() / FORM()	Return the current NUMERIC DIGITS / FUZZ / FORM settings.

Type / data

Signature	Returns
DATATYPE(s)	'NUM' if s parses as a number; 'CHAR' otherwise.
DATATYPE(s, 'N')	1 / 0 — is s numeric?
DATATYPE(s, 'W')	1 / 0 — is s a whole number?
DATATYPE(s, 'A')	1 / 0 — is s alphanumeric (letters, digits, no spaces)?

Date / time

Signature	Returns
DATE() / DATE('N')	Today as dd Mmm yyyy.
DATE('S')	Today as yyyyymmdd.
DATE('E')	Today as dd/mm/yy.
DATE('U')	Today as mm/dd/yy.
DATE('O')	Today as yy/mm/dd.
DATE('B')	Days since 0001-01-01.
DATE('B', 'yyyyymmdd', 'S')	Basedate of a given date — subtract two basedates for an exact day delta.
TIME() / TIME('N')	hh:mm:ss.
TIME('L')	hh:mm:ss.uuuuuu with microseconds.
TIME('S')	Seconds since midnight.

Variables / arguments

Signature	Returns
VALUE(name)	Read the value of the variable named at runtime — useful for indirect reads (VALUE('SCR.ROW'
VALUE(name, newval)	Set the variable named at runtime; returns the prior value. The two-argument form usually appears under `CALL VALUE 'SCR.ROW'
ARG()	Number of arguments the current procedure was called with.
ARG(n)	The nth argument (1-based), or empty if out of range.
SYMBOL(name)	Classifies a symbol: 'VAR' if it names a variable that has a value, 'LIT' if it is a valid symbol with no value assigned (or a constant symbol such as 77, .5 or 3D, which can never be assigned), 'BAD' if it is not a valid REXX symbol at all. Evaluation is two-step and stops there: the argument expression is evaluated once and the RESULT is the symbol name, so with J = 3 SYMBOL('J') tests J while SYMBOL(J) tests the name 3. Compound tails are derived exactly as an ordinary reference derives them, so SYMBOL('A.J') tests A.3. The answer is always three uppercase characters.
ADDRESS()	The name of the current host-command environment, uppercased — CICS unless an ADDRESS instruction changed it, and restored to the caller's value when a PROCEDURE returns. Note that ADDRESS() written alone as a whole statement is the ADDRESS <i>instruction</i> , not this function; use it inside an expression (SAY ADDRESS()).
RANDOM([min [, max]])	Random integer in [min, max] (default [0, 999]).

Signature	Returns
ERRORTEXT (code)	Echoes the numeric code as a string (placeholder — bricks does not maintain the standard REXX error-text table).

Stream I/O

The stream functions read and write text files in the `tmp_dir` sandbox — the same backend that powers COBOL's EXEC CICS READQ TD / WRITEQ TD / DELETEQ TD (see [Chapter 9](#)). The sandbox is strict: ASCII only, LF-terminated, no traversal.

Function	Form	Behaviour
LINEIN(name)	next line	Auto-opens name for read on first call; returns the next line (LF stripped). Empty string at EOF.
LINEIN(name, 1)	rewind	Closes and reopens name, then returns line 1. Only 1 is honoured; any other line number is treated as the "next line" form.
LINEOUT(name)	close	Closes the named stream. Returns 0.
LINEOUT(name, line)	append	Auto-opens name for append on first call; appends `line`
LINES(name)	EOF probe	Returns 1 while more data is available, 0 at EOF.
CHARIN(name, start, n)	byte read	Reads n bytes (default 1). start = 1 rewinds first; other start values are ignored.
CHAROUT(name, s)	byte write	Appends raw bytes. Same ASCII/LF validation as LINEOUT.
CHARS(name)	bytes remaining	Returns the file size minus the current read position.
STREAM(name, 'S')	state	Returns 'READY', 'NOTREADY', or 'ERROR'.
STREAM(name, 'D')	description	Returns a short status string ('OK' / parse-time reason).
STREAM(name, 'C', cmd)	command	OPEN READ / OPEN WRITE / OPEN APPEND / CLOSE / QUERY EXISTS / QUERY SIZE / DELETE.

```
/* Append a row, then read every row back. */
CALL LINEOUT 'export.txt', 'C0000099|WIDGET-Z|7|9.95'
CALL LINEOUT 'export.txt' /* close the writer */
DO WHILE LINES('export.txt') > 0
  SAY LINEIN('export.txt')
END
CALL STREAM 'export.txt', 'C', 'CLOSE'
```

Every handle a program opens is closed automatically at task end; an interpreter that forgets to call CLOSE does not leak descriptors. A sandbox violation ('../etc/passwd', 'sub/x', leading-dot) causes the next LINEIN / LINEOUT to raise an error and STREAM('S') to report 'ERROR'.

JSON

For programs that consume JSON — typically a REST response read with EXEC CICS WEB CONVERSE — two functions extract values by a dotted path. They parse the document on each call (fine for the small payloads a 3270 program renders).

Signature	Returns
JSONGET(json, path)	The scalar at path as a string. path is dot-separated; an integer segment indexes a JSON array (0-based). Returns '' for a missing key, an out-of-range index, a JSON null, a value that is itself

Signature	Returns
	an object/array, or a parse error. JSON numbers render like REXX numbers (671, not 671.0); booleans become 1/0.
JSONCOUNT(<i>json</i> , <i>path</i>)	The number of elements in the array at path (the empty path '' walks the whole document). '0' for a missing path, a non-array value, or a parse error.

```

/* RESP holds the WEB CONVERSE body */
N = JSONCOUNT(RESP, 'data')           /* array length           */
DO I = 0 TO N - 1
  P  = 'data.' || I
  FLT = JSONGET(RESP, P || '.flight.iata') /* "BA178"                */
  DEP = JSONGET(RESP, P || '.departure.iata')
  ARR = JSONGET(RESP, P || '.arrival.iata') /* path-aware: dep != arr */
END
MSG = JSONGET(RESP, 'error.message')    /* '' if no error object  */

```

Limitations (acceptable for typical REST payloads): a key containing a literal . is unreachable; array indices are non-negative only (no wildcards, no negative indices); JSON integers beyond 2^53 lose precision (they decode through a float). There is no JSON *builder* — assemble output JSON with string concatenation, as the WAPI sample does. (COBOL programs have JSON GENERATE / JSON PARSE instead — see [Chapter 22](#).)

Worked example. The **SABR** reservation transaction (`runtime/rexx/book.rexx`) uses JSONGET / JSONCOUNT to parse live aviationstack flight data fetched with EXEC CICS WEB CONVERSE and paint departures / arrivals boards on the 3270 — see its D0_FLIFO handler, and the **Live flight info (FLIFO)** section of [SABRE.md](#) for the operator-facing D0<apt>/D commands it drives.

Operators

Listed lowest to highest precedence. All operators are left-associative except ** which is right-associative.

Precedence	Operators	Notes
1 (lowest)		Logical OR. Strict boolean: both sides must evaluate to 0 or 1. Short-circuit.
2	& / &&	Logical AND. Same strict-boolean rules. && is the strict-AND alias. Short-circuit.
3	=== \= \== <> > < < > <= >=	Comparison. The strict variants (==, \==) compare byte-for-byte; the plain forms compare as numbers when both sides parse as numeric (subject to NUMERIC FUZZ) and as trimmed strings otherwise. \= and <> and >< are all "not equal".
4	and juxtaposition	Concatenation. A B joins with no separator; A B (whitespace between operands) joins with a single space. There is no zero-separator juxtaposition — write explicitly.
5	+ −	Addition / subtraction.
6	* / % //	Multiplication, real division, integer division (% — truncates toward zero), modulo (// — sign of the dividend).
7	unary + − \	Arithmetic sign and logical NOT.
8 (highest)	**	Exponentiation. Integer exponents are exact (via repeated squaring); fractional exponents use float64.

Comparison precedence note: REXX has no separate equality vs. ordering precedence, so A < B = C parses as (A < B) = C — comparing the boolean result of A < B against C.

Restrictions / unsupported

Feature	Status
TRACE (interpretive tracer)	Parsed and accepted; bricks does not emit trace output.
OPTIONS settings (NOVALUE, etc.)	Parsed and accepted; have no effect.
External function libraries / RXFUNCADD	No dynamic linking — every routine must be a PROCEDURE in the program or a built-in.
ADDRESS COMMAND / ADDRESS SYSTEM / ADDRESS TSO	Not wired — only ADDRESS CICS is recognised.
EXECIO and other host-environment I/O	Not provided — use the stream functions (next section) or EXEC CICS READQ TD / WRITEQ TD for file I/O.
PARSE PULL / terminal queue (PUSH / QUEUE)	Statements are parsed; the queue is always empty. Use EXEC CICS RECEIVE for terminal input.
NUMERIC FORM SCIENTIFIC / ENGINEERING	Statement is parsed; rendering is always plain decimal.
Strict IEEE arithmetic	Backed by float64 with NUMERIC DIGITS rounding, so very large factorials and the like will round before they overflow.

Part 2. The COBOL language

Part 2 begins here. Bricks ships a free-form COBOL interpreter (`cobol/`) that sits beside REXX as a second front end on the same EXEC CICS / EXEC SQL / EXEC CICS WEB surface. The chapters that follow describe COBOL source format, the DATA DIVISION, PROCEDURE DIVISION, the EIB binding, the COPY directive (including the standard DFHAID, DFHRESP, SQLCA, and DFHWB* copybooks), and the restrictions. Command-level behaviour is identical between COBOL and REXX — see **Parts 4, 6, and 7** for the verbs.

Chapter 20. COBOL source format

Bricks ships a free-form COBOL interpreter (`cobol/`) that sits beside REXX as a second front end on the same EXEC CICS surface. The language-specific layer is `cobol/frame.go`, which adapts COBOL's group-item world to the REXX-style STEM.TAIL lookup the CICS handlers use.

Source rules

- **Free-form.** No column 1-72 ruling; no Area A / Area B.
- **Hyphens** are allowed in identifiers (CUST-RECORD).
- **Comments** use modern `*>` anywhere on a line, or legacy `*` in column 1.
- **Strings** use `'...'` or `"..."` with quote-doubling for embedded quotes (`'don't', "say ""hi"""`).
- **Numeric literals** are decimal digit runs with an optional decimal point (42, 1.5, 0.050) and an optional sign written against the digits (-1, +5, -1.5). They work anywhere a number is accepted: VALUE, MOVE, COMPUTE, the arithmetic verbs, and conditions (IF B / 10 = 4.5). A fractional literal is placed by the receiving PIC's V (see [VALUE on a scaled item](#)).
- **Hex literals** like `X'F3'` and `X"7C"` decode to their byte value — used for IF EIBAID = X'F3' to check PF3 / PF12 / etc. without C2X(EIBAID) = 'F3' round-trips.
- **Identifiers** start with a letter or `_`, continue with letters, digits, `_`, and `-`. Names are case-insensitive (uppercased at parse time).
- **Periods** are scope terminators, not per-statement. The parser requires a period after each division header, section header, paragraph header, and top-level data item; periods inside an explicit END-IF / END-EVALUATE / etc. body are tolerated as inner punctuation, not as scope terminators.

Divisions

The four divisions parse in their canonical order:

```

IDENTIFICATION DIVISION.
PROGRAM-ID. HELC.

ENVIRONMENT DIVISION.          *> optional, must be empty

DATA DIVISION.
WORKING-STORAGE SECTION.
01 SCR.
    05 INFOLINE PIC X(78).
    05 GREETING PIC X(20).

PROCEDURE DIVISION.
MAIN.
    MOVE 'Hello' TO GREETING.
    EXEC CICS SEND MAP('HEL01') FROM(SCR) ERASE END-EXEC.
    EXEC CICS RETURN END-EXEC.
    STOP RUN.

```

An optional LINKAGE SECTION may follow WORKING-STORAGE (see [LINKAGE SECTION](#)); DFHCOMMAREA is auto-injected as PIC X(2000) if the program doesn't declare it (see `cobol.ensureSystemItems`), so a sub-program can MOVE DFHCOMMAREA TO key immediately. The dispatcher strips trailing space when reading the COBOL frame's DFHCOMMAREA back out, so a fixed-width buffer round-trips cleanly across an inter-language EXEC CICS LINK.

Chapter 21. DATA DIVISION

The DATA DIVISION declares every data item the program touches. Bricks expects a single WORKING-STORAGE SECTION followed by an optional LINKAGE SECTION (DFHCOMMAREA is auto-injected when the program omits it).

Level numbers

Level	Meaning
01	Top-level data item. Owns a self-contained byte buffer.
02 – 49	Subordinate items inside a group. Children share their parent's buffer at a computed offset.
77	Standalone elementary item (no children). Functionally equivalent to a flat 01.
88	Condition-name attached to the preceding non-88 item (see below).

The level numbers 66 (RENAMES) and 78 (constant) are **not** parsed.

PIC clauses — unedited

Unedited PICs hold raw data: digits, letters, or padded bytes.

Form	Class	Notes
X(n) or XXXX...	Alphanumeric	n bytes; left-justified, space-padded on MOVE. Up to 1 MiB.
A(n) or AAA...	Alphabetic	Same storage as X(n); runtime class-check rejects non-alpha on INSPECT-style tests.
9(n) or 999...	Numeric (unsigned)	n decimal digits; right-justified, zero-padded. Capped at 18 digits (int64 limit).

Form	Class	Notes
S9(n)	Numeric (signed)	Byte 0 holds the sign character (' ' for positive, – for negative); bytes 1..n-1 hold the digits. Cap 18 digits.
9(n)V9(m)	Numeric with implicit decimal	The V is positional — no byte is consumed. Total n+m digits stored; arithmetic uses the fixed-point decimal type so scale survives MOVE / COMPUTE / DISPLAY.
S9(n)V9(m)	Signed scaled numeric	Same as above plus a sign byte.

(n) is a repetition count; 9(3) and 999 are equivalent.

PIC clauses — edited numeric

Edited PICs hold a formatted representation of a numeric value suitable for display. Five edit characters are recognised:

Edit char	Role
Z	Replace a leading zero with a space (zero suppression).
.	Insert a literal decimal point at that column.
,	Insert a thousands separator at that column.
\$	Insert a \$. A <i>single</i> \$ is fixed (always emitted, must be the leftmost character). Two or more in a row <i>float</i> — the leftmost is the symbol and slides right to the position immediately before the first significant digit.
*	Check protection. Replaces a leading zero with * instead of a space — used on cheques to prevent forgery.

Constraints:

- At most one suppression family per PIC: Z, floating \$, and * are mutually exclusive (a single fixed \$ may coexist with Z — e.g. \$ZZ9.99).
- Z and * are forbidden in the fractional half (after the . insertion).
- The literal width of the mask is the PIC length; total digit positions (Z + floating-\$ + * + 9) are capped at 18.

Worked examples:

Source	PIC mask	Stored bytes
42	ZZ9	" 42"
0	ZZ9	" 0"
1.5	ZZZ.99	" 1.50"
1234.56	Z,ZZ9.99	"1,234.56"
12345.67	\$\$\$,\$\$9.99	"\$12,345.67"
0.45	\$\$\$,\$\$\$9.99	" \$.45"
12.34	**, **9.99	"****12.34"
42	\$9999	"\$0042"

When a scaled numeric source is moved through an edited receiver the source's V-scale is honoured: MOVE N TO N-Z where N is PIC 9(5)V99 VALUE 12345.67 and N-Z is PIC ZZ,ZZ9.99 stores "12,345.67".

The remaining ANSI edit characters — +, −, CR, DB, B (blank), 0 (zero insertion), / (slash), and the BLANK WHEN ZERO clause — are not yet supported.

VALUE

VALUE declares the initial content of an elementary data item. Accepted forms:

Form	Example
String literal	VALUE 'ENTER ID'. or VALUE "OK".
Numeric literal	VALUE 42. or VALUE 0. (signed +/- allowed).
Figurative	VALUE SPACES., VALUE ZEROS., VALUE HIGH-VALUES., VALUE LOW-VALUES., VALUE QUOTES.

The figurative aliases SPACE / ZERO / ZEROS / HIGH-VALUE / LOW-VALUE / QUOTE are normalised to their plural form at parse time. VALUE on a group item is rejected — initialise the children individually.

VALUE on a scaled (V) item

A **numeric** VALUE carries the value, and the PIC's v places it — exactly as MOVE and COMPUTE do with the same literal:

```
01 A PIC S9(4)V99 VALUE 1.5.      *> 1.50, stored 000150
01 B PIC S9(4)V99 VALUE 2.        *> 2.00, stored 000200
01 C PIC S9(4)  VALUE 1.5.        *> 1, truncated toward zero
```

Do **not** write the stored digit run and expect it to be read as scaled: PIC S9(4)V99 VALUE 150 is one hundred fifty, not 1.50.

A **quoted** VALUE is byte data and takes the alphanumeric path, so VALUE '150' on that same item is 1.50 — the digits are laid into the field's bytes rather than scaled. This mirrors MOVE '150', so a literal means the same thing wherever it is written.

Group items

A 01 (or any subordinate level >01) item with one or more children becomes a group. The group's storage is the concatenation of its children's storage; the group itself owns no extra bytes beyond what its children occupy.

```
01 PARENT.
  05 CHILD PIC X(8).
  05 OTHER PIC X(4).
```

MOVE source TO PARENT fans out to every elementary child; EXEC CICS SEND MAP FROM(PARENT) walks the children by name to look up map-field values.

OCCURS

OCCURS n [TIMES] declares an array of n copies of an item. The cap is 4096 occurrences per item.

Reference one slot with a 1-based subscript expression: K(I), K(5), K(I + 1). Subscript bounds are checked at runtime. OCCURS works on elementaries and on group items; nested OCCURS (an array of arrays) is **rejected at parse time**.

```

01 TABLE-AREA.
   05 ROW OCCURS 15 TIMES.
       10 K PIC X(8).
       10 N PIC X(28).
01 I PIC 9(4).
...
MOVE 'KEY-A' TO K(1).
MOVE 'KEY-B' TO K(I).

```

REDEFINES

REDEFINES overlays a second description on storage that an earlier item already owns. No new bytes are allocated — the redefining item and the redefined item are two names for the same buffer slice, and a MOVE through either name is visible through the other immediately. The PIC of whichever name the statement references decides how the bytes are interpreted.

```
level data-name-1 REDEFINES data-name-2 [PIC ...] [USAGE ...].
```

The clause must be the first thing after the data name. Rules bricks enforces at parse time (each mirrors the IBM Enterprise COBOL rule unless flagged):

- data-name-1 and data-name-2 must carry the **same level number** (01–49 or 77; never 88).
- data-name-2 must be the **immediately preceding** item at that level under the same parent. Only other redefinitions of the same item (with their subordinates and 88s) may sit between them.
- Multiple redefinitions all name the **original** item: B REDEFINES A then C REDEFINES A — never C REDEFINES B.
- The redefining item must **fit**: its storage may not exceed the redefined item's. This holds at level 01 too — stricter than IBM, which lets a non-EXTERNAL 01 redefiner grow the area (deferred; see Chapter 27). A *shorter* redefiner is fine and only overlays the leading bytes.
- Neither item may carry OCCURS on its own entry, and neither may live inside an OCCURS group. **Subordinates of the redefining item may carry OCCURS** — that subordinate-table shape is the main reason the clause exists.
- No VALUE anywhere in the redefining item's subtree (88-level condition VALUEs are fine). A VALUE on the *redefined* item is visible through the overlay at startup.
- FILLER may redefine (05 FILLER REDEFINES D8.); nothing can redefine a FILLER (it has no referenceable name). Likewise the auto-injected system items (EIB*, RC, SQLCA*, JSON-CODE, JSON-STATUS, DFHCOMMAREA) cannot be redefinition targets — they don't exist yet when the DATA DIVISION parses.

Dual view — one set of bytes, two shapes:

```

01 WS-DATE PIC 9(8) VALUE 20260611.
01 WS-DATE-R REDEFINES WS-DATE.
   05 WD-YYYY PIC 9(4).
   05 WD-MM PIC 9(2).
   05 WD-DD PIC 9(2).
...
DISPLAY WD-MM.           *> 06
MOVE 12 TO WD-MM.
DISPLAY WS-DATE.         *> 20261211

```

Row table — name the rows for the map, subscript them for the loop. This replaces the unrolled N-branch EVALUATE fan-out that older sample programs carry (see esdc.cob):

```

01 SCR.
  05 INFOLINE PIC X(78).
  05 ROWS-AREA.
    10 ROW1 PIC X(76).
    10 ROW2 PIC X(76).
    10 ROW3 PIC X(76).
  05 R-TAB REDEFINES ROWS-AREA.
    10 ROWV PIC X(76) OCCURS 3.
...
MOVE WS-LINE TO ROWV(WS-SLOT).
EXEC CICS SEND MAP('ESDS') FROM(SCR) ERASE END-EXEC.

```

SEND MAP FROM(SCR) still finds ROW1-ROW3: map routing resolves qualified names through storage, and the redefinition shares their bytes, so whatever the loop wrote through ROWV(n) is what the map paints.

Storage caveat — signed DISPLAY fields. bricks stores PIC S9(n) as **n+1 bytes**: byte 0 is a leading separate sign (' ' positive, '-' negative), bytes 1..n the digits — in effect SIGN IS LEADING SEPARATE, with space standing in for +. IBM's default is n bytes with the sign overpunched into the trailing digit. A redefinition sees bricks's layout, not IBM's: the full alphanumeric view of PIC S9(4) is X(5), and it reads -0123, sign first. Ported IBM copybooks that redefine signed DISPLAY fields must be widened by one byte. PIC 9(n), COMP (big-endian two's complement, 2/4/8 bytes) and COMP-3 (BCD, C/D/F sign nibble) overlays are byte-for-byte IBM-faithful — though a PIC X view of COMP bytes shows raw binary, exactly as it would on the mainframe.

FILLER

FILLER is a reserved name for anonymous storage. Multiple FILLER children under the same parent are allowed (they don't collide on the duplicate-name check) and they're never registered in the lookup tables — they can only be addressed through the parent.

USAGE — DISPLAY vs. COMP vs. COMP-3

Numeric items have three storage layouts:

USAGE	Synonyms	Storage
DISPLAY (default)	—	One byte per digit; signed PICs reserve byte 0 for the sign character (' ' / '-'). The historical bricks layout.
COMP	COMPUTATIONAL, COMP-4, BINARY	Big-endian two's-complement signed integer packed into 2, 4, or 8 bytes based on the total digit count.
COMP-3	COMPUTATIONAL-3, PACKED-DECIMAL	Packed decimal: two BCD digits per byte (high nibble first), with the last byte's low nibble carrying the sign.
POINTER	USAGE IS POINTER	An opaque 4-byte pointer handle, declared without a PIC clause: 01 WS-PTR USAGE POINTER.. Holds a CICS storage handle from EXEC CICS ADDRESS / GETMAIN or an ADDRESS 0F value. See the COBOL pointer model in Chapter 22.

The USAGE keyword is optional — PIC S9(8) COMP. and PIC S9(8) USAGE IS COMP. both work. USAGE may appear before or after the PIC clause within the same item.

COMP byte sizing follows IBM's halfword / fullword / doubleword convention:

Total digits (PIC 9(n) or S9(n)V9(m) → n+m)	Bytes
1 – 4	2 (halfword)

Total digits (PIC 9(n) or S9(n)V9(m) → n+m)	Bytes
5 – 9	4 (fullword)
10 – 18	8 (doubleword)

USAGE COMP requires a numeric PIC (9 or S9); declaring it on PIC X(n) is rejected at parse time. The signed/unsigned flag does *not* change the byte count — both PIC 9(4) COMP and PIC S9(4) COMP occupy two bytes. Bricks always stores the value as a two's-complement signed int internally.

Scale survives the format transition: a value moved between a PIC 9(5)V99 DISPLAY field and a PIC S9(8)V99 COMP field preserves the implicit decimal point. Within the interpreter every numeric load goes through the fixed-point decimal engine; DISPLAY decodes digits, COMP decodes binary bytes, but both produce the same decimal{val, scale} tuple.

COMP-3 byte sizing uses the packed-decimal formula `ceil((digits + 1) / 2)` — digits plus the trailing sign nibble, rounded up:

Total digits	Bytes
1	1
4	3
9	5
15	8
18	10

The sign nibble follows IBM Enterprise COBOL convention: 0xC for signed positive, 0xD for signed negative, 0xF for unsigned. On read, sign nibbles 0xB and 0xD are treated as negative; everything else is positive. Like COMP, COMP-3 requires a numeric PIC and the digit cap is 18 (the int64 fixed-point engine's ceiling) — PIC S9(19) and wider are parser-rejected.

The canonical mainframe shape PIC S9(15) COMP-3 is the standard CICS EIBABSTIME carrier:

```
01 SYS-DT PIC S9(15) COMP-3.
EXEC CICS ASKTIME ABSTIME(SYS-DT) END-EXEC.
EXEC CICS FORMATTIME ABSTIME(SYS-DT) YYYYMMDD(YYYYMMDD) END-EXEC.
```

ASKTIME writes the 15-digit ms-since-1900 decimal string into SYS-DT; bricks packs it into 8 BCD bytes transparently. FORMATTIME reads the packed bytes back into the same decimal string and parses it.

Other USAGE variants — COMP-1 (single-precision float) and COMP-2 (double-precision float) are explicitly rejected at parse time with a "not yet supported" message. SYNC, INDEXED BY, and OCCURS DEPENDING ON are also unsupported. (REDEFINES IS supported — see the [REDEFINES](#) section above.)

Level 88 condition-names

A level-88 line declares a boolean condition tied to the **preceding non-88 data item**. The condition is true when the parent's current value matches one of the listed values or falls within a THRU range. Reference as a bare name in any PROCEDURE DIVISION condition (see [Chapter 22 — 88-level condition names](#) for the runtime behaviour).

```

01 RESP-CODE PIC 9(3).
   88 RESP-OK      VALUE 0.
   88 RESP-MISSING  VALUE 13.
   88 RESP-RECOVER  VALUES 12, 26, 80.
   88 RESP-FATAL    VALUE 100 THRU 999.

```

VALUES (plural) and VALUE (singular) are interchangeable. Multiple values are separated by spaces or commas. THRU (synonym THROUGH) declares an inclusive range; ranges and single values can be mixed: 88 OK VALUES 0, 1 THRU 5, 9.

A condition-name lives in its own namespace, separate from data items. Declaring an 88 whose name collides with an existing data item is rejected at parse time.

LINKAGE SECTION

The LINKAGE SECTION declares records that are addressed through a pointer rather than allocated their own storage. Bricks supports both the classic auto-injected DFHCOMMAREA and operator-declared **based** records redirected with SET ADDRESS OF.

DFHCOMMAREA. Bricks auto-injects DFHCOMMAREA PIC X(2000) when the program doesn't declare it, so a sub-program can MOVE DFHCOMMAREA TO key immediately. The dispatcher strips trailing space when reading the COBOL frame's DFHCOMMAREA back out, so a fixed-width buffer round-trips cleanly across an inter-language EXEC CICS LINK. DFHCOMMAREA is itself a **based** 01 with a default buffer: a program can MOVE ... TO DFHCOMMAREA out of the box, *and* redirect it onto another area's bytes with SET ADDRESS OF DFHCOMMAREA TO ptr.

DFHEIB. The Execute Interface Block is a real 01 DFHEIB based block; its children keep the bare IBM names (EIBRESP, EIBAID, EIBCALEN, EIBTRMID, EIBCP05N, ...) so a program references them unqualified exactly as on z/OS. EXEC CICS ADDRESS EIB(ptr) hands back a handle addressing it.

Based records — BASED. A LINKAGE 01 declared BASED owns no storage of its own:

```

LINKAGE SECTION.
01 CUST-REC BASED.
   05 CUST-NO PIC X(8).
   05 CUST-NAME PIC X(30).

```

Until the program points it somewhere with SET ADDRESS OF, CUST-REC is NULL — any read or write through it is a clean NULL-pointer abend (ASRA-style), never a crash. After SET ADDRESS OF CUST-REC TO ptr (where ptr came from GETMAIN, ADDRESS, or ADDRESS OF another item), every reference to CUST-NO / CUST-NAME reads and writes the addressed bytes. See the [COBOL pointer model](#) in Chapter 22 for the full surface.

Cross-group name collisions — OF / IN qualification

A child name that appears under more than one group is fine; the parser accepts the declaration and tracks the collision in Program.AmbiguousNames. Unqualified access to such a name is a runtime error ("ambiguous reference") — disambiguate with OF (or its synonym IN):

```

01 SCR.
   05 CUSTNO PIC X(8).
01 DET.
   05 CUSTNO PIC X(8).
...
MOVE 'A1234567' TO CUSTNO OF SCR.
MOVE 'B7654321' TO CUSTNO OF DET.
DISPLAY CUSTNO OF SCR.

```

Single-step qualification (`X OF Y`) picks the matching child anywhere in `Y`'s subtree; multi-step (`X OF Y OF Z`) chains through nested groups. Bare names that are unique across the program continue to work without qualification — the pre-7 runtime/cobol/gust.cob convention of prefixed child names (`DCUSTNO`, `DNAME`, `DMSG`) still compiles and runs, it just isn't required any more.

A sibling-duplicate (two children of the SAME parent sharing a name) is still a parse-time error — no qualifier can disambiguate between siblings of one group.

Type coercion on MOVE

Source class	Target class	Behaviour
Alphanumeric → Alphanumeric	X / A → X / A	Copy, right-pad with spaces if target is wider; right-truncate if source is wider.
Numeric → Numeric	9 / S9 → 9 / S9	Load through the fixed-point decimal engine, rescale to the target's V position, store. Truncates high-order if the integer part doesn't fit (silent unless <code>ON SIZE ERROR</code> is set on the surrounding verb).
Numeric → Edited	9 / S9 → Z / \$ / * / ...	Read the source as a decimal (honouring V-scale), render the value through the edited mask.
Edited → Numeric	Z / \$ / * / ... → 9 / S9	De-edit on read by stripping insertion characters (., ., \$, *, spaces) and re-parsing the digits.
Alphanumeric → Numeric	X → 9	Best-effort <code>loadNumeric</code> — succeeds if the byte content parses as digits after trimming. A letter raises a runtime data exception.
Numeric → Alphanumeric	9 / S9 → X	Source stringified, then alphanumeric-padded.
Anything → Group	* → group item	Fans out to every elementary child (recursively).

A **numeric literal** source follows the Numeric → Numeric row: it is rescaled to the target's V position, so `MOVE 2 TO PIC S9(4)V99` stores 2.00 and `MOVE 1.5 TO PIC S9(4)` truncates to 1. `MOVE`, `COMPUTE` and a numeric `VALUE` clause all agree on the same literal. A **quoted** literal is alphanumeric and takes the byte path instead (`MOVE '150'` lays those digits into the field's bytes).

Figurative expansion

`SPACES` → ' ' repeated to fill the target; `ZEROS` → '0' for numeric targets, `\x00` for alphanumeric; `HIGH-VALUES` → `\xFF`; `LOW-VALUES` → `\x00`; `QUOTES` → '\ '. Group targets compute the fill length from the sum of the elementary children's storage.

Chapter 22. PROCEDURE DIVISION

The PROCEDURE DIVISION is a flat list of paragraphs. Each paragraph is a label (PARAGRAPH-NAME.) followed by one or more statements. Paragraphs are visible to PERFORM and GO TO.

Statement reference

Data movement

Verb	Syntax	Notes
MOVE	MOVE src TO tgt1, tgt2, ...	One or more targets; fan-out to group children supported. Type coercion table in Chapter 21 — Type coercion .
INITIALIZE	Not yet supported.	Initialise individual items with MOVE.

Arithmetic

Every arithmetic verb accepts the optional clauses ROUNDED, ON SIZE ERROR ... [END-VERB], and GIVING (where listed). See "ROUNDED and ON SIZE ERROR" below.

Verb	Syntax
COMPUTE	COMPUTE tgt [ROUNDED] = expr [ON SIZE ERROR ...] [END-COMPUTE]
ADD	ADD a TO b [GIVING c] [ROUNDED] [ON SIZE ERROR ...] [END-ADD]
SUBTRACT	SUBTRACT a FROM b [GIVING c] [ROUNDED] [ON SIZE ERROR ...] [END-SUBTRACT]
MULTIPLY	MULTIPLY a BY b [GIVING c] [ROUNDED] [ON SIZE ERROR ...] [END-MULTIPLY]
DIVIDE (INTO)	DIVIDE a INTO b [GIVING c] [ROUNDED] [ON SIZE ERROR ...] [END-DIVIDE]
DIVIDE (BY)	DIVIDE a BY b GIVING c [ROUNDED] [ON SIZE ERROR ...] [END-DIVIDE]

REMAINDER is not yet supported on DIVIDE. Arithmetic expressions support +, -, *, /, **, parentheses, and unary sign.

Signed numeric literals

A sign written **immediately** against the digits is part of the literal, exactly as on IBM, so a signed literal is legal anywhere a plain one is:

```
MOVE -1 TO WS-IND.  
MOVE +5 TO WS-N.  
ADD -1 TO WS-COUNT.  
DISPLAY -1.  
COMPUTE WS-X = -1.5.  
IF WS-N = -100 THEN ...  
01 WS-LIMIT PIC S9(4) VALUE -150.
```

The spacing is what distinguishes a literal from an operator: -1 is a signed literal, - 1 is the minus operator applied to 1. Inside an arithmetic expression both readings work out the same (COMPUTE X = A - 1 and COMPUTE X = -1 are each unambiguous), but in a list-taking statement they differ — DISPLAY A -1 displays two items, while DISPLAY A - 1 is a parse error, because DISPLAY takes a list of operands, not an expression.

A redundant + is not stored, so `DISPLAY +5` prints 5. Moving a negative literal into an unsigned PIC 9 target drops the sign, as IBM specifies.

Control flow

Verb	Syntax	Notes
IF	IF cond [THEN] stmts [ELSE stmts] [END-IF]	The period after the last unscoped statement closes the IF when END-IF is omitted.
EVALUATE	EVALUATE subject [ALSO subject ...] {WHEN object [ALSO object ...]} ... stmts [WHEN OTHER stmts] [END-EVALUATE]	Full COBOL-85 form: TRUE / FALSE / condition / value subjects, ANY, NOT, THRU ranges, stacked WHENs. See EVALUATE .
PERFORM	PERFORM para [THRU para2] PERFORM [para] [WITH TEST BEFORE AFTER] UNTIL cond PERFORM [para] N TIMES PERFORM [para] VARYING idx FROM x BY y UNTIL cond [AFTER ...] ...any of the above with the paragraph name omitted and an in-line body closed by END-PERFORM	Out-of-line (named paragraph) and in-line (END-PERFORM) forms both work. WITH TEST is accepted but inert on the N TIMES and no-clause shapes.
GO TO	GO TO para	Calculated GO TO ... DEPENDING ON not supported.
CONTINUE	CONTINUE	Explicit no-op (used in empty WHEN OTHER etc.).
STOP RUN	STOP RUN (or just STOP)	Terminate the task.
GOBACK	GOBACK	Synonym for EXIT PROGRAM.
EXIT	EXIT [PROGRAM] EXIT PERFORM [CYCLE] EXIT PARAGRAPH	EXIT alone is a no-op; EXIT PROGRAM returns to the caller. EXIT PERFORM leaves the innermost in-line PERFORM, EXIT PERFORM CYCLE ends only the current iteration. EXIT PARAGRAPH ends the current paragraph (see EXIT PARAGRAPH). EXIT SECTION is a parse error — bricks has no sections.

Strings

Verb	Syntax
STRING	STRING src ... DELIMITED [BY] (SIZE delim) ... INTO tgt [[WITH] POINTER ptr] [[ON] OVERFLOW stmts] [NOT [ON] OVERFLOW stmts] [END-STRING] — see STRING and UNSTRING
UNSTRING	UNSTRING src [DELIMITED [BY] [ALL] d1 [OR [ALL] d2] ...] INTO {t [DELIMITER [IN] d] [COUNT [IN] c]} ... [[WITH] POINTER ptr] [TALLYING [IN] n] [[ON] OVERFLOW stmts] [NOT [ON] OVERFLOW stmts] [END-UNSTRING]
INSPECT TALLYING	INSPECT subject TALLYING counter FOR (ALL LEADING CHARACTERS) [needle] [BEFORE/AFTER INITIAL delim] [, ...]
INSPECT REPLACING	INSPECT subject REPLACING (ALL LEADING FIRST CHARACTERS) [needle] BY replacement [BEFORE/AFTER INITIAL delim] [, ...]
INSPECT CONVERTING	INSPECT subject CONVERTING from TO to [BEFORE/AFTER INITIAL delim]

JSON

Verb	Syntax
JSON GENERATE	JSON GENERATE tgt FROM rec [COUNT [IN] n] [NAME [OF] item [IS] {'lit' OMITTED} ...] [SUPPRESS ...] [CONVERTING ...] [[ON] EXCEPTION stmts] [NOT [ON] EXCEPTION stmts] [END-JSON] — see JSON GENERATE and JSON PARSE
JSON PARSE	JSON PARSE src INTO rec [WITH DETAIL] [NAME ...] [SUPPRESS item ...] [CONVERTING ...] [[ON] EXCEPTION stmts] [NOT [ON] EXCEPTION stmts] [END-JSON]

I/O and external

Verb	Syntax	Notes
DISPLAY	DISPLAY item1 item2 ...	Concatenates items with no separator; newline after the list.
ACCEPT	Not yet supported.	Use EXEC CICS RECEIVE for terminal input.
EXEC CICS	EXEC CICS verb operand-list END-EXEC	Same verb set as REXX. See Parts 4-5 and 7.
EXEC SQL	EXEC SQL stmt END-EXEC	See Part 6.
CALL, SEARCH, ALTER	Not yet supported.	

Every target name above (the second operand of MOVE, the target of COMPUTE / ADD / SUBTRACT / MULTIPLY / DIVIDE / GIVING / STRING INTO / UNSTRING INTO / INSPECT, and the subject of EVALUATE / IF) accepts qualification via OF / IN (see the [Data Division](#) section on globally non-unique names) and subscripts via (idx) for OCCURS-resident items.

Reference modification (string slicing)

identifier(start:length) addresses a byte window inside a data item. Positions are **1-indexed**, length is a count of bytes, and both operands may be any integer arithmetic expression — literals, data-names, or WS-POS + 1. Omitting the length (identifier(start:)) means "from start to the end of the item", padding included.

Slices work on both sides of a statement:

```
01 WS-LINE  PIC X(80) VALUE 'HELLO WORLD'.
01 WS-PART  PIC X(5).
01 WS-POS   PIC 9(2)  VALUE 7.

MOVE WS-LINE(1:5)      TO WS-PART      *> read  -> 'HELLO'
MOVE WS-LINE(WS-POS:5) TO WS-PART      *> read  -> 'WORLD'
MOVE '*'              TO WS-LINE(6:1)  *> write -> 'HELLO*WORLD'
IF WS-LINE(6:1) = SPACE THEN ...      *> test one byte
```

A write touches **only** the named window; the bytes on either side keep their contents. That makes MOVE CH TO BUF(POS:1) the natural way to render a screen row one character at a time, without an OCCURS table or REDEFINES overlay.

Three rules worth knowing:

- **A slice is always category alphanumeric**, whatever the item's own PIC says. NUM-FLD(1:1) compares as bytes, so IF NUM-FLD(1:1) = '1' is the correct test and = 1 is not. On a write, the value

is left-justified and space-padded to the window width — never right-justified or zero-filled, even into a PIC 9 item.

- **A USAGE COMP / COMP-3 item is sliced through its decimal rendering**, not its raw bytes: PIC 9(4) COMP VALUE 1234, sliced (1:2), yields '12'.
- **Out-of-range windows are errors.** A read whose length overruns the item is space-padded, but a start below 1 or past the end of the item aborts the transaction, as does any write whose window would run past the item's last byte — those bytes belong to the next field in the same 01 and silently overwriting them corrupts it.

ROUNDED and ON SIZE ERROR

ROUNDED applied to COMPUTE and the four arithmetic verbs rounds the final value half-away-from-zero to the destination's PIC scale instead of truncating toward zero (the unROUNDED default).

ON SIZE ERROR runs its statement list when the calculation overflows the int64 fixed-point engine, the result's integer-part digits don't fit in the target PIC, or a DIVIDE divisor is zero; when present, the destination is left untouched on a size-error firing. Without an ON SIZE ERROR branch the value is silently truncated/wrapped as before so existing programs aren't disrupted.

PERFORM N TIMES and PERFORM VARYING

```
PERFORM FILL-ROW 15 TIMES.  
  
PERFORM FILL-ROW  
  VARYING I FROM 1 BY 1 UNTIL I > 15.
```

PERFORM ... N TIMES runs the named paragraph exactly N times; N can be a numeric literal or a numeric data item. PERFORM ... VARYING idx FROM x BY y UNTIL cond initialises idx to x, runs the body while cond is false, then increments idx by y between iterations.

In-line PERFORM

Leave the paragraph name out and the loop body sits between the PERFORM and a mandatory END-PERFORM:

```
PERFORM UNTIL WS-DONE = 'Y'  
  COMPUTE WS-N = WS-N + 1  
  IF WS-N > 9 THEN  
    MOVE 'Y' TO WS-DONE  
  END-IF  
END-PERFORM.  
  
PERFORM 10 TIMES  
  COMPUTE WS-TOT = WS-TOT + 1  
END-PERFORM.  
  
PERFORM VARYING I FROM 1 BY 1 UNTIL I > 15  
  MOVE SPACES TO ROW-TEXT(I)  
END-PERFORM.
```

All four loop shapes work in-line — no clause (the body runs once), UNTIL, N TIMES, and VARYING. In-line and out-of-line PERFORM nest freely in either direction.

END-PERFORM is **required** on the in-line form: unlike IF, a period does not close it. A period inside the body ends the sentence and is reported as in-line PERFORM needs END-PERFORM.

PERFORM THRU

```
PERFORM WORK-A THRU WORK-EXIT.
```

Runs every paragraph from WORK-A through WORK-EXIT in source order, then returns. THROUGH is accepted as a synonym. This enables the classic early-return idiom — a GO TO aimed at a paragraph **inside** the range transfers control there and the range still returns to the caller when its last paragraph ends:

```
PERFORM CHECK-IT THRU CHECK-EXIT.  
...  
CHECK-IT.  
    IF WS-AMT = 0 THEN  
        GO TO CHECK-EXIT  
    END-IF.  
    COMPUTE WS-FEE = WS-AMT * 2.  
CHECK-EXIT.  
EXIT.
```

A GO TO aimed **outside** the range still abandons the PERFORM and restarts at the target, which is bricks' general GO TO behaviour (see [Periods and GO TO](#)). A backwards range (PERFORM Z THRU A) is rejected at run time.

WITH TEST BEFORE / AFTER

```
PERFORM WITH TEST AFTER UNTIL WS-RESP = DFHRESP(NORMAL)  
    EXEC CICS READ FILE('CUSTOMER') INTO(WS-REC)  
        RIDFLD(WS-KEY) RESP(WS-RESP) END-EXEC  
END-PERFORM.
```

WITH TEST BEFORE is the default: the condition is tested at the top, so a condition that is already true runs the body zero times. WITH TEST AFTER tests at the bottom, so the body always runs at least once. WITH is optional (TEST AFTER alone is accepted). It applies to the UNTIL and VARYING shapes; on N TIMES and the no-clause form it is accepted but has no effect.

On a post-tested VARYING the flow is body → test → augment, so the index is left at its last in-range value rather than one step past the limit.

EXIT PERFORM

```
PERFORM VARYING I FROM 1 BY 1 UNTIL I > 100  
    IF TBL-KEY(I) = WS-WANTED THEN  
        MOVE I TO WS-HIT  
        EXIT PERFORM  
    END-IF  
    IF TBL-KEY(I) = SPACES THEN  
        EXIT PERFORM CYCLE  
    END-IF  
    COMPUTE WS-SCANNED = WS-SCANNED + 1  
END-PERFORM.
```

EXIT PERFORM breaks out of the innermost in-line PERFORM; EXIT PERFORM CYCLE ends only the current iteration and continues with the next one (still reaching a VARYING augment, so the loop advances rather than spinning).

Both are scoped to the **in-line** PERFORM that lexically contains them, matching IBM. An EXIT PERFORM written in a performed *paragraph* is a program error, not a way to break the caller's loop, and is reported as EXIT PERFORM outside an in-line PERFORM.

EXIT PARAGRAPH

EXIT PARAGRAPH (COBOL 2002, IBM Enterprise COBOL 6) ends the current paragraph at once. Control goes wherever the end of that paragraph would have gone:

- under PERFORM para (also TIMES / UNTIL / VARYING) — back to the performer, or on to the next iteration;
- under PERFORM a THRU b — on to the next paragraph of the range;
- when execution simply fell into the paragraph — on to the next paragraph.

It works at any nesting depth — inside IF, EVALUATE, an ON OVERFLOW phrase or an in-line PERFORM. Inside an in-line PERFORM it leaves the whole paragraph, not just the loop (use EXIT PERFORM for that).

```
VALIDATE-INPUT.  
  IF WS-NAME = SPACES  
    MOVE 'NAME REQUIRED' TO WS-MSG  
    EXIT PARAGRAPH  
  END-IF  
  PERFORM CHECK-DUPLICATE.
```

With WS-NAME blank, the PERFORM VALIDATE-INPUT returns with WS-MSG = NAME REQUIRED and CHECK-DUPLICATE is not run.

EXIT PARAGRAPH. written as a sentence of its own is fine. (Earlier bricks releases read it as EXIT. followed by a new paragraph named PARAGRAPH, so the program compiled and silently did the wrong thing.)

EXIT SECTION is rejected at parse time with EXIT SECTION needs SECTIONs, not supported in bricks COBOL: bricks has no PROCEDURE DIVISION sections, so there is nothing for it to leave.

PERFORM VARYING ... AFTER

```
PERFORM VARYING I FROM 1 BY 1 UNTIL I > 4  
  AFTER J FROM 1 BY 1 UNTIL J > 3  
  COMPUTE SLOT((I - 1) * 3 + J) = I * J  
END-PERFORM.
```

Each AFTER adds an inner dimension. Inner indices are re-initialised to their FROM value on every iteration of the dimension outside them, so the example above runs the body 12 times. AFTER may be repeated for further dimensions. A plain EXIT PERFORM unwinds the whole nest, not just the innermost dimension.

⚠ bricks rejects nested OCCURS, so a two-dimensional walk addresses a flat table with a computed subscript as shown above rather than SLOT(I, J).

OCCURS arrays

```
01 TABLE-AREA.  
  05 ROW OCCURS 15 TIMES.  
    10 K PIC X(8).  
    10 N PIC X(28).  
01 I PIC 9(4).  
...  
MOVE 'KEY-A' TO K(1).  
MOVE 'KEY-B' TO K(I).  
PERFORM SHOW-ROW VARYING I FROM 1 BY 1 UNTIL I > 15.
```

OCCURS *n* TIMES declares *n* copies of an item; references add a parenthesised 1-based subscript expression that picks one. The subscript can be any numeric expression — a literal, a data item, or arithmetic — and is bounds-checked at runtime. Subscripts work on group items (ROW(5)) and on elementaries inside the OCCURS group (K(5)) alike. The cap is 4096 occurrences per item; nested OCCURS (an OCCURS group inside another OCCURS group) is rejected at parse time.

Periods

Periods are scope-terminators, not per-statement: a statement inside an IF ... END-IF body has no period; only the unscoped statement at the paragraph level does.

Postfix NOT

IBM-style postfix NOT (X NOT = Y, X NOT > Y) is supported in IF / EVALUATE / PERFORM UNTIL conditions.

Relational operators

Both symbolic and English-spelled relational operators are accepted in IF, EVALUATE WHEN, and PERFORM UNTIL conditions. The two styles are interchangeable; legacy CICS applications typically use the English spelling.

The optional words IS, THAN, and TO are noise — present or omitted without changing meaning. NOT may prefix any operator and appears before or after IS (IF X IS NOT EQUAL TO Y).

Meaning	Symbolic	English-spelled forms (noise words bracketed)
equal	=	[IS] EQUAL [TO]
not equal	<>, NOT =	[IS] NOT EQUAL [TO]
greater	>	[IS] GREATER [THAN]
not greater ($\leq$)	NOT >	[IS] NOT GREATER [THAN]
less	<	[IS] LESS [THAN]
not less ($\geq$)	NOT <	[IS] NOT LESS [THAN]
greater or equal	>=	[IS] GREATER [THAN] OR EQUAL [TO]
less or equal	<=	[IS] LESS [THAN] OR EQUAL [TO]

The four lines below are equivalent — all four compile to the same comparison and run identically:

```

IF WK-MTH = 7          DISPLAY 'JULY'.
IF WK-MTH EQUAL 7      DISPLAY 'JULY'.
IF WK-MTH EQUAL TO 7   DISPLAY 'JULY'.
IF WK-MTH IS EQUAL TO 7 DISPLAY 'JULY'.

```

Combined conditions and parentheses

Conditions combine with NOT, AND and OR (precedence NOT > AND > OR) and can be grouped with parentheses anywhere a condition is allowed (IF, PERFORM UNTIL, EVALUATE objects):

```

IF A = 0 OR (A = 1 AND B = 2) ...
IF NOT (WS-CODE = 'X' OR WS-CODE = 'Y') ...
IF ((A + 1) * 2 = B) OR (C = 1) ...

```

NOT must be followed by a condition: IF NOT NOT A = 1 is rejected with NOT must be followed by a condition.

Division in a condition follows IBM's intermediate-result rule (Enterprise COBOL Programming Guide, Appendix A): a quotient keeps (d2 - d1) or dmax decimal places, whichever is greater. d2 and d1 are the decimal places of the dividend and the divisor; dmax is the largest number of decimal places of any operand in the relation, both sides counted, divisors and exponents excluded. The quotient is truncated, not rounded:

Condition	Values	Compares	Result
IF 5 / 2 > 2	—	2 > 2	false (as on z/OS)
IF A / 2 > 2	A PIC 9(3) = 5	2 > 2	false
IF B / 10 = 4.5	B PIC 9(3) = 45	4.5 = 4.5	true (the literal gives dmax 1)

An EVALUATE subject and each object are evaluated on their own operands (their own dmax): EVALUATE B / 10 WHEN 4.5 does **not** match for B = 45, because B / 10 alone has no decimal places and gives 4. COMPUTE, subscripts and reference modifiers are not affected — COMPUTE still divides at the receiving field's scale.

A parenthesis that opens an arithmetic operand still works: IF (A + 1) > B and IF (A + 1) * 2 = B OR (C = 1) read the first group as arithmetic because a relational or arithmetic operator follows it.

Abbreviated combined relation conditions

After a relation, later relations in the same condition may omit the subject, or the subject and the operator; the last stated subject and operator are filled in (IBM rules):

Written	Means
A = 1 OR 2 OR 3	A = 1 OR A = 2 OR A = 3
A > 1 AND < 9	A > 1 AND A < 9
A NOT = 1 AND 2	A NOT = 1 AND A NOT = 2
A > B AND NOT < C OR D	(A > B AND A NOT < C) OR A NOT < D
A > 5 AND NOT 9	A > 5 AND NOT (A > 9)
NOT A = B OR C	NOT (A = B) OR A = C

A NOT directly in front of a relational operator belongs to the operator and is carried along with it; any other NOT is the logical NOT. A level-88 condition-name after AND / OR is a condition, not an implied object (IF A =

1 OR WS-DONE). As in Enterprise COBOL, an omitted subject may not open a parenthesised group: A = B AND (< C OR D) is rejected with a relational operator cannot start a parenthesised condition. A group that starts with a full relation may use the short forms inside it: IF B = 0 AND (A = 4 OR 5 OR 6).

Changed behaviour. Earlier bricks releases accepted these forms but evaluated them wrongly: IF A = 1 OR 2 OR 3 was true whatever the value of A, and IF C NOT = 'L' AND 'H' dropped the NOT on the second part, so it was true for C = 'H'. Both now follow the IBM rules above (C NOT = 'L' AND C NOT = 'H' is false for C = 'H'). Check old programs that relied on the previous results.

Class condition tests

X IS [NOT] {NUMERIC | ALPHABETIC | ALPHABETIC-UPPER | ALPHABETIC-LOWER} tests the current **byte contents** of a data item against a character class — not its static PIC class. A PIC 9(4) field holding 0000 is NUMERIC; the same field after a partial move that injects letters is not.

IS is optional noise; NOT negates as in any other condition.

```
IF WS-INPUT IS NUMERIC
  PERFORM CALC
ELSE
  DISPLAY 'NON-NUMERIC INPUT'.
IF WS-NAME IS ALPHABETIC
  PERFORM ACCEPT-NAME.
```

ALPHABETIC accepts uppercase letters, lowercase letters, and space. ALPHABETIC-UPPER and ALPHABETIC-LOWER restrict to their case (plus space). For PIC S9 signed numerics the byte 0 sign position is allowed to hold ' ', '+', or '-'.

Sign condition tests

X IS [NOT] {POSITIVE | NEGATIVE | ZERO} tests the numeric sign of a numeric expression. The operand must be a numeric data item (PIC 9... or PIC S9...) or a numeric expression; a sign test on a non-numeric operand is a runtime error rather than a silent false.

```
IF WS-BALANCE IS NEGATIVE
  DISPLAY 'OVERDRAWN'.
IF AMOUNT IS ZERO
  PERFORM SKIP-ENTRY.
```

Signed zero is neither POSITIVE nor NEGATIVE; an unsigned PIC (PIC 9... without S) can never be NEGATIVE — that short-circuits to false without inspecting the value.

88-level condition names

A level-88 line in DATA DIVISION declares a boolean condition tied to the **preceding non-88 data item**. The condition is true when the parent's current value matches one of the listed values or falls within a THRU range. Reference as a bare name in any PROCEDURE DIVISION condition.

```

DATA DIVISION.
WORKING-STORAGE SECTION.
01 RESP-CODE PIC 9(3).
   88 RESP-OK          VALUE 0.
   88 RESP-MISSING     VALUE 13.
   88 RESP-RECOVER     VALUES 12, 26, 80.
   88 RESP-FATAL       VALUE 100 THRU 999.
PROCEDURE DIVISION.
MAIN.
   EXEC CICS READ FILE('CUSTOMER')
               INTO(REC) RIDFLD(KEY)
               RESP(RESP-CODE) END-EXEC.

   IF RESP-OK
       CONTINUE
   ELSE IF RESP-RECOVER
       PERFORM RETRY
   ELSE IF RESP-FATAL
       PERFORM ABEND.

```

VALUES (plural) and VALUE (singular) are interchangeable. Multiple values may be separated by commas or by spaces. THRU (synonym THROUGH) declares an inclusive range; ranges and single values can be mixed in one declaration: 88 OK VALUES 0, 1 THRU 5, 9.

Comparison style follows the parent's PIC class — numeric parents compare numerically, alphanumeric parents compare byte-by-byte (after rstripping trailing spaces). NOT and the AND/OR connectives compose normally: IF NOT RESP-OK OR RESP-FATAL.

A condition-name lives in its own namespace, separate from data items — 01 STATUS PIC X. and 88 STATUS-OK VALUE 'Y'. are two different lookups. Declaring an 88 whose name collides with an existing data item is rejected at parse time.

Intrinsic functions

The shipped intrinsic library is:

Function	Args	Result
FUNCTION UPPER-CASE(s)	1 alphanumeric	s uppercased.
FUNCTION LOWER-CASE(s)	1 alphanumeric	s lowercased.
FUNCTION TRIM(s)	1 alphanumeric	s with leading and trailing spaces removed (both ends; one-arg form).
FUNCTION REVERSE(s)	1 alphanumeric	s reversed (rune-aware).
FUNCTION LENGTH(s)	1 alphanumeric	numeric byte length of the operand's storage (PIC X(8) → 8).
FUNCTION NUMVAL(s)	1 alphanumeric	numeric value of s after stripping leading / trailing whitespace; errors on non-numeric content (does not silently coerce to zero).
FUNCTION POS(needle, haystack)	2 alphanumeric	1-based byte position of needle in haystack, or 0 when not found. Mirrors REXX's POS(). Empty needle returns 0.

Numeric-returning intrinsics (LENGTH, NUMVAL, POS) participate in IF / EVALUATE / arithmetic comparisons as numeric operands, so IF FUNCTION POS(NEEDLE, HAY) > 0 works without an intermediate COMPUTE. The remaining intrinsics return alphanumeric and compare by bytes (rstripped).

FUNCTION POS is the substring-search hook used by runtime/cobol/gusl.cob to filter the customers file when the operator types a search term at GUST's S action.

INSPECT

```
INSPECT subject TALLYING counter FOR (ALL | LEADING | CHARACTERS)
                                   [needle] [BEFORE/AFTER INITIAL delim]
                                   [, ...]
INSPECT subject REPLACING (ALL | LEADING | FIRST | CHARACTERS)
                        [needle] BY replacement
                        [BEFORE/AFTER INITIAL delim]
                        [, ...]
INSPECT subject CONVERTING from TO to
                        [BEFORE/AFTER INITIAL delim]
```

Quantifiers:

- **ALL** — every occurrence of `needle` (TALLYING) or every occurrence rewritten (REPLACING).
- **LEADING** — only successive occurrences at the start of the in-scope region.
- **FIRST** — only the first occurrence (REPLACING only).
- **CHARACTERS** — every character in the in-scope region; for TALLYING this just counts characters, for REPLACING it overwrites each one with the first byte of replacement.

Limiters:

- **BEFORE INITIAL delim** narrows the in-scope region to everything before the first occurrence of `delim`.
- **AFTER INITIAL delim** narrows to everything after the first occurrence of `delim`.

Both can be set on the same phrase; AFTER picks the start, BEFORE picks the end.

Multiple phrases can be combined per statement, optionally separated by commas:

```
INSPECT REC REPLACING ALL ',' BY '|' ALL ';' BY '/'.
INSPECT REC TALLYING C FOR LEADING '0', ALL 'X' AFTER INITIAL ' '.
```

Counter behaviour: TALLYING **accumulates** — successive INSPECT TALLYING statements add to the existing counter value rather than resetting, matching IBM semantics. REPLACING phrases run left-to-right in declared order, so later phrases see the output of earlier ones.

Converting: CONVERTING `from TO to` builds a positional per-character map and translates every character of the in-scope region through it — `from[k]` becomes `to[k]`, exactly like Unix `tr` or the old COBOL TRANSFORM. This is the verb for case folding:

```
INSPECT VMUSR CONVERTING 'abcdefghijklmnopqrstuvwxyz'
                        TO 'ABCDEFGHIJKLMNOPQRSTUVWXYZ'.
```

If `from` repeats a character, the **leftmost** occurrence defines its mapping and later duplicates are ignored. The sets need not be the same length: only the overlapping $\min(\text{len}(\text{from}), \text{len}(\text{to}))$ positions map and any surplus is silently ignored (no runtime error), matching the REPLACING leniency philosophy. IBM allows only **one** CONVERTING phrase per INSPECT (no comma list), and BEFORE/AFTER INITIAL `delim` narrows the region just like REPLACING.

Subject, counter, and operands all accept qualified data-names (INSPECT REC OF DET ...). Trailing PIC X padding is rstripped from both the subject and the needle before matching, so a 30-byte filter containing F00 plus trailing spaces matches against the meaningful content of the subject rather than against the padding.

IBM's CONVERTING form is now supported (see the Converting paragraph above).

EVALUATE

```
EVALUATE subject-1 [ALSO subject-2 ...]
  {WHEN object-1 [ALSO object-2 ...]} ...
    statements
  [WHEN OTHER statements]
[END-EVALUATE]
```

A subject is a data item, literal, arithmetic expression, a condition, TRUE or FALSE. Each WHEN must list exactly one object per subject, separated by ALSO; a mismatch is a parse error (WHEN has 1 object(s) but EVALUATE has 2 subject(s)). Objects pair with their subject like this:

Subject	Allowed objects
item / literal / arithmetic expression	[NOT] value, [NOT] value THRU value, ANY
TRUE, FALSE, or a condition	a condition (88-level names, NOT, AND / OR, parentheses), TRUE, FALSE, ANY

THRU (or THROUGH) is an inclusive range: WHEN 1 THRU 5 means subject >= 1 AND subject <= 5, compared with the same rules as IF. Several WHEN phrases in a row share one body. The first matching WHEN runs, with no fall-through; WHEN OTHER runs when none matched. Every subject is evaluated once. END-EVALUATE may be left out when a period ends the statement.

A WHEN takes one value. WHEN 'A' OR 'B' is rejected with WHEN takes one value – write alternatives as separate WHEN lines; stack the WHENS instead (WHEN 'A' / WHEN 'B' sharing one body). This does not apply to a condition object (under EVALUATE TRUE, FALSE or a condition subject), where OR is part of the condition: WHEN A = 1 OR B = 2 and the abbreviated WHEN A = 1 OR 3 both work.

```
EVALUATE TRUE
  WHEN WS-QTY = 0      MOVE 'EMPTY' TO WS-STATE
  WHEN WS-QTY < 10
  WHEN RUSH-ORDER      MOVE 'LOW'   TO WS-STATE
  WHEN OTHER           MOVE 'OK'    TO WS-STATE
END-EVALUATE

EVALUATE WS-ACTION ALSO WS-LEVEL
  WHEN 'A' ALSO 1 THRU 3  PERFORM ADD-BASIC
  WHEN 'A' ALSO ANY      PERFORM ADD-FULL
  WHEN NOT 'Q' ALSO ANY  PERFORM OTHER-ACTION
END-EVALUATE

EVALUATE WS-QTY > 0 ALSO RUSH-ORDER
  WHEN TRUE  ALSO TRUE    PERFORM SHIP-NOW
  WHEN TRUE  ALSO FALSE  PERFORM SHIP-LATER
  WHEN FALSE ALSO ANY     PERFORM BACKORDER
END-EVALUATE
```

STRING and UNSTRING

STRING joins the sending items into the receiver. **DELIMITED BY SIZE** moves the whole item — a data item at its full PIC width (trailing spaces included), a literal exactly as written; **DELIMITED BY x** stops at the first x. Several sending items may share one **DELIMITED BY**. bricks never trims a sending identifier: to drop its padding, use **DELIMITED BY SPACE**. The receiver must not be reference-modified (parse error **STRING INTO** receiver must not be reference-modified, as IBM requires).

- Without **POINTER** the receiver is space-filled first, then written (bricks behaviour; IBM leaves unwritten positions alone).
- With **[WITH] POINTER ptr** the IBM rules apply: writing starts at position ptr, ptr goes up by one per character, and the rest of the receiver keeps its contents — so successive **STRING**s can build a line piece by piece.
- Overflow: the data did not fit, or ptr was < 1 or past the end at the start (then nothing moves). **ON OVERFLOW** runs on overflow, **NOT ON OVERFLOW** otherwise.

```
01 WS-NAME  PIC X(10) VALUE 'JOE SMITH'.
01 WS-LINE  PIC X(12) VALUE '.....'.
01 WS-PTR   PIC 99.
01 WS-TRUNC PIC X VALUE 'N'.
...
MOVE 1 TO WS-PTR
STRING 'Name: ' DELIMITED BY SIZE
      WS-NAME  DELIMITED BY SPACE
      INTO WS-LINE WITH POINTER WS-PTR
      ON OVERFLOW MOVE 'Y' TO WS-TRUNC
END-STRING
DISPLAY WS-LINE '|' WS-PTR
```

Output: Name: JOE...|10 — the literal keeps its trailing space, the last three dots are untouched, and WS-PTR points at the next free position.

UNSTRING splits the sending item, scanning from **POINTER** (default 1) over its full length. The sending item must be a data item, as in IBM; a literal, figurative constant or function is rejected with **UNSTRING source** must be a data item, got "...".

- each receiver gets the characters up to the next delimiter; with several delimiters (**OR**) the first one found wins, and on a tie the first one listed; **ALL** treats a run of the same delimiter as one; two delimiters in a row (without **ALL**) give an empty field, i.e. spaces or zeros;
- **DELIMITER IN** receives the delimiter that ended the field (spaces at the end of the data), **COUNT IN** the field's length; both need **DELIMITED BY** (else **DELIMITER IN** / **COUNT IN** need **DELIMITED BY**);
- without **DELIMITED BY** each receiver takes as many characters as its own size;
- **TALLYING IN** is increased by the number of receivers filled; **POINTER** ends one past the last character examined;
- receivers that are not reached keep their contents. When the data is used up exactly on a delimiter ('a|b|' in a 4-byte field with three receivers), the third receiver is not reached and keeps its old value, as IBM specifies. Earlier bricks versions set it to spaces;
- overflow: **POINTER** < 1 or past the end at the start, or characters are left over when every receiver is filled. **END-UNSTRING** is optional.

```

01 WS-REC      PIC X(20) VALUE 'AB   CD,EF'.
01 WS-F1       PIC X(5).
01 WS-F2       PIC X(5).
01 WS-D1       PIC X.
01 WS-N1       PIC 99.
01 WS-PTR      PIC 99 VALUE 1.
01 WS-FIELDS   PIC 99 VALUE 0.
...
UNSTRING WS-REC DELIMITED BY ALL SPACES OR ','
  INTO WS-F1 DELIMITER IN WS-D1 COUNT IN WS-N1
  WS-F2
  WITH POINTER WS-PTR TALLYING IN WS-FIELDS
  ON OVERFLOW DISPLAY 'MORE DATA'
END-UNSTRING
DISPLAY WS-F1 '|' WS-D1 '|' WS-N1 '|' WS-F2 '|' WS-PTR '|' WS-FIELDS

```

Output: MORE DATA, then AB | |02|CD |09|02. The run of spaces after AB is one delimiter (ALL), WS-F2 stops at the comma, and EF plus the padding are left over, which is the overflow. POINTER 9 is where a second UNSTRING would carry on.

Changed behaviour. Earlier bricks releases read INTO B COUNT IN P as a second receiver (the qualified name COUNT IN P); the program compiled and then failed at run time. DELIMITER IN, COUNT IN, WITH POINTER and TALLYING IN are now recognised wherever IBM allows them.

JSON GENERATE and JSON PARSE

The IBM Enterprise COBOL 6.4 statements turn a record into JSON text and back, so a web API program no longer has to build or pick apart JSON by hand. Both work on plain alphanumeric fields, which is exactly what EXEC CICS WEB RECEIVE INTO and WEB SEND FROM use.

```

01 LST-REC.
  05 LST-KEY.
    10 LST-USER PIC X(16).
    10 LST-ID   PIC 9(6).
  05 LST-NEXTITEM PIC 9(8).
  05 LST-COUNT   PIC 9(5).
  05 LST-NAME    PIC X(80).
01 JSON-BUF      PIC X(4000).
01 JS-LEN        PIC 9(5).

  JSON GENERATE JSON-BUF FROM LST-REC COUNT IN JS-LEN
  NAME OF LST-REC IS OMITTED LST-ID IS 'id' LST-NAME IS 'name'
  SUPPRESS LST-USER LST-NEXTITEM
  ON EXCEPTION PERFORM FAIL-JSON
END-JSON
EXEC CICS WEB SEND FROM(JSON-BUF) LENGTH(JS-LEN) END-EXEC

```

gives {"LST-KEY":{"id":12},"LST-COUNT":3,"name":"Buy milk"}.

Special registers

Register	Shape	Meaning
JSON-CODE	PIC S9(9) BINARY	0 after a successful statement, else the exception code (tables below).
JSON-STATUS	PIC 9(9) BINARY	Set by JSON PARSE: the sum of the non-exception reason codes.

Both are declared automatically, like EIBRESP and SQLCODE. JSON-STATUS is additive (1 + 2 = 3 means "an item had no member and a member had no item"); IF JSON-STATUS = 0 tests for a perfect match.

JSON GENERATE

```
JSON GENERATE identifier-1 FROM identifier-2
[COUNT [IN] identifier-3]
[NAME [OF] identifier-4 [IS] {literal | OMITTED} ...]
[SUPPRESS {identifier-5 [WHEN fig [OR fig] ...]
           | EVERY [NUMERIC | NONNUMERIC] WHEN fig [OR fig] ...} ...]
[CONVERTING {identifier-6 TO JSON BOOLEAN USING {condition-name | literal}
            | identifier-6 TO JSON NULL USING fig} ...]
[[ON] EXCEPTION imperative] [NOT [ON] EXCEPTION imperative]
[END-JSON]
```

fig is ZERO, SPACE, LOW-VALUE or HIGH-VALUE (singular or plural). The phrases must appear in the order shown (a phrase out of order is a parse error, as in IBM's syntax diagram); entries of SUPPRESS and CONVERTING may be separated by ALSO.

- **Shape.** {"REC":{...}} names the FROM item. NAME OF rec IS OMITTED drops that wrapper, which is what web clients want. The FROM item may be a group, an elementary item, a subscripted occurrence, or a whole OCCURS table — with OMITTED the table becomes an anonymous array [{...}, {...}], and an elementary item becomes just its value (1234, "B0B"; IBM APAR PH58384).
- **Names.** A member's name is its data name; NAME OF item IS 'lit' replaces it. bricks upper-cases data names when it reads the source, so 05 Cust-Name comes out as "CUST-NAME" (IBM keeps the source spelling). Use NAME for lower-case JSON names.
- **Groups** become objects, OCCURS items arrays. FILLER, REDEFINES items (and everything below them) and USAGE POINTER items are left out. A group whose members are all left out or suppressed disappears; the outermost item is always written, even as {"REC":{}}.
- **Numbers** are written as JSON numbers without leading zeros, with a - sign and a decimal point from the PIC's V: PIC 9(6) 12 gives 12, PIC S9(3)V999 gives -63.435, PIC 9V99 1.5 gives 1.50 (fraction zeros stay), zero gives 0. COMP and COMP-3 work the same. Numeric-edited items are strings, trimmed of leading and trailing spaces.
- **Text** loses its trailing spaces; leading spaces stay. A field of only spaces is written as " " (IBM). " and \ are escaped, as are control characters (\t \n \r \b \f, others as \u00hh). UTF-8 passes through; a byte that is not valid UTF-8 (HIGH-VALUES) is written as \u00hh.
- **SUPPRESS** item leaves that item (or group) out. item WHEN fig and EVERY [NUMERIC | NONNUMERIC] WHEN fig leave out an elementary item whose value matches: ZERO matches numbers equal to zero and text of only 0 characters; SPACE, LOW-VALUE, HIGH-VALUE match DISPLAY items only. Elements of an elementary OCCURS table are not suppressed one by one (that would shift the positions of the rest).
- **CONVERTING** item TO JSON BOOLEAN USING cond-name | literal writes true when the item satisfies the condition-name (which must be an 88 of that item) or equals the literal, otherwise false. item TO JSON NULL USING fig writes null when the item is all fig.
- **COUNT IN** receives the number of characters generated. The rest of the receiver is filled with spaces. The receiver may be subscripted or reference-modified (JSON GENERATE BUF(8:30) ...); then only that window is written and space-filled, and the too-small test and COUNT IN are relative to it.
- **Exception.** If the receiver is too small, JSON-CODE is 1, the receiver holds what fit, COUNT IN holds the receiver's length and ON EXCEPTION runs. Without ON EXCEPTION the program just continues. Otherwise JSON-CODE is 0 and NOT ON EXCEPTION runs.

JSON PARSE

```
JSON PARSE identifier-1 INTO identifier-2
  [WITH DETAIL]
  [NAME [OF] identifier-3 [IS] {literal | OMITTED} ...]
  [SUPPRESS identifier-4 ...]
  [CONVERTING {identifier-5 FROM JSON BOOLEAN USING
               {cond-1 [AND cond-2] | literal-1 [AND literal-2]}} ...]
  [[ON] EXCEPTION imperative] [NOT [ON] EXCEPTION imperative]
[END-JSON]
```

```
01 IN-REC.
   05 IN-TITLE PIC X(80).
   05 IN-PRIO  PIC X.
   05 IN-DONE  PIC X.

*> body: {"title":"Buy milk","priority":"H","done":true}
JSON PARSE WS-BODY INTO IN-REC
  NAME OF IN-REC IS OMITTED IN-TITLE IS 'title'
  IN-PRIO IS 'priority' IN-DONE IS 'done'
  CONVERTING IN-DONE FROM JSON BOOLEAN USING 'Y' AND 'N'
  ON EXCEPTION PERFORM BAD-REQUEST
END-JSON
```

- identifier-1 holds the text (it may be reference-modified); trailing spaces and LOW-VALUES are ignored.
- **Matching.** Without OMITTED the text must be {"IN-REC":{...}}. With NAME OF IN-REC IS OMITTED the outermost object's members map straight onto IN-REC's subordinates (or, for a whole OCCURS table, the outermost array onto its occurrences). Data names match case-insensitively; a NAME literal replaces the data name and matches case-sensitively. FILLER and REDEFINES items never match; SUPPRESS items are skipped.
- **Values** are assigned with MOVE rules: a number (exponent allowed) or a string holding a number goes to numeric and numeric-edited items, aligned on the decimal point (extra fraction digits are truncated); strings and numbers go to alphanumeric items (padded / truncated on the right). Objects go to groups, arrays to OCCURS items. true / false need a CONVERTING entry: the item gets literal-1 / cond-1's value for true and literal-2 / cond-2's value for false. null leaves the item unchanged. \uxxxx escapes are stored as UTF-8.
- Members with no matching item are ignored and items with no member keep their contents; JSON-STATUS reports it:

JSON-STATUS	Meaning (IBM)
1	An item had no matching name, so it was not changed.
2	A JSON name matched no item.
4	An item had duplicate names with the same value.
8	A table had more elements than the JSON array.
16	A JSON array had more values than the table (extras ignored).
32	An item was not changed because the value was null.
64	A table element was not changed because the value was null.
128	SIZE ERROR in a numeric assignment (stored anyway, high digits lost).
256	Text was truncated in an alphanumeric assignment (stored anyway).

JSON-CODE	Exception (IBM)
100	The JSON text is invalid (also: nesting deeper than 64 levels, top level not an object).
101	The JSON text is empty or only whitespace.
102	Extra characters after the outermost value.
103	Duplicate names with different values; the leftmost was kept.
104	A value does not fit its item (object into an elementary item, text into a numeric item, ...).
106	No JSON name matched any item. In the {"REC": {...}} form the outer name matching the INTO item already counts as a match, so {"REC": {"foo": 1}} is code 0 with JSON-STATUS 3.
107	true / false for an item without CONVERTING.
108	A top-level array without NAME ... OMITTED on an OCCURS item.

An exception stops the statement and runs ON EXCEPTION. The whole text is checked before anything is stored, so invalid JSON (100-102) leaves the record untouched; a later exception (103, 104, 107) leaves the members before it assigned. WITH DETAIL writes one log line per statement with JSON-CODE, JSON-STATUS and the reason.

Differences from IBM: names come out upper-case (see above); the single condition-name form USING cond-1 cannot set "false" because bricks has no 88 ... WHEN SET TO FALSE — JSON false gives code 104, so use USING cond-1 AND cond-2. Invalid JSON text leaves the whole record untouched, where IBM may already have updated part of it. ENCODING, IGNORING, INDICATING, CONVERTING ... FROM JSON NULL and OCCURS DEPENDING ON are not supported.

COBOL pointer model: USAGE POINTER / SET ADDRESS OF

Bricks implements the IBM COBOL/CICS pointer surface so a program can address dynamically allocated or CICS-owned storage. **Pointers are opaque handles, not machine addresses** — there is no pointer arithmetic and they are not displayable numbers.

USAGE POINTER **items**

Declared with no PIC clause:

```
01 WS-PTR  USAGE POINTER.          *> 'USAGE IS POINTER' also accepted
```

A pointer occupies a 4-byte fullword slot and starts **NULL**. Its value comes from EXEC CICS ADDRESS / GETMAIN, from ADDRESS OF identifier, or from another pointer; it can be compared against NULL and against ADDRESS OF X, and moved like any other item. (Internally a pointer's value is a small integer handle; the program never sees or depends on the integer.)

ADDRESS OF identifier

ADDRESS OF X is the address of data item X — usable on the right of SET, and in an equality test:

```
SET WS-PTR TO ADDRESS OF WS-REC.
IF WS-PTR = ADDRESS OF WS-REC THEN DISPLAY 'same' END-IF.
```

ADDRESS OF of the same item always compares equal (handles are interned), so pointer equality is meaningful.

NULL / NULLS

The figurative NULL (synonym NULLS) is the null pointer:

```
SET WS-PTR TO NULL.  
IF WS-PTR = NULL THEN DISPLAY 'unset' END-IF.
```

A freshly declared USAGE POINTER item, and any based 01 that has not been pointed anywhere, are NULL.

SET P TO {ADDRESS OF X | NULL | ptr}

Assigns a pointer variable: to the address of an item, to NULL, or to the value of another pointer.

SET ADDRESS OF L TO {ptr | ADDRESS OF X | NULL}

Redirects a **based** LINKAGE 01 onto bytes. After the SET, every reference to L's children reads / writes the addressed storage:

```
LINKAGE SECTION.  
01 GM-AREA BASED.  
    05 GM-TEXT PIC X(16).  
PROCEDURE DIVISION.  
    EXEC CICS GETMAIN SET(WS-PTR) FLENGTH(16) END-EXEC.  
    SET ADDRESS OF GM-AREA TO WS-PTR.      *> map the record onto the buffer  
    MOVE 'POINTER-OK' TO GM-TEXT.          *> writes into the GETMAIN bytes
```

SET ADDRESS OF L TO NULL un-addresses L; a subsequent read or write through L raises a **clean NULL-pointer abend** — the task ends with a diagnostic, the server never crashes. The same applies to a based 01 that was never SET, and to dereferencing through a pointer whose GETMAIN buffer was already FREEMAIN-ed.

Based LINKAGE 01s, DFHEIB, DFHCOMMAREA

DFHCOMMAREA and DFHEIB are themselves based 01s. DFHCOMMAREA ships with a default buffer (so a plain MOVE ... TO DFHCOMMAREA works without any SET) yet can be redirected: SET ADDRESS OF DFHCOMMAREA TO ADDRESS OF WS-NEW. DFHEIB's children keep the bare IBM names so EIBRESP, EIBCALEN, etc. are referenced unqualified. See [LINKAGE SECTION](#) in Chapter 21.

SET ptr UP BY n / SET ptr DOWN BY n — **pointer arithmetic**

SET ptr UP BY n advances a USAGE POINTER item by n bytes; SET ptr DOWN BY n moves it back. n is a **non-negative integer** displacement — a literal, a data item, or an arithmetic expression that evaluates to a whole number (the same evaluator COMPUTE uses). A non-integer (e.g. 1.5) or a negative amount is rejected at SET time; UP/DOWN carry the direction, so use DOWN BY n rather than UP BY -n:

```

01 WS-REC PIC X(10) VALUE 'AAAAAAAAAA'.
01 P      USAGE POINTER.
01 L BASED.
   05 L-FLD PIC X(6).
...
   SET P TO ADDRESS OF WS-REC.
   SET P UP BY 4.           *> P now addresses byte 4 of WS-REC
   SET ADDRESS OF L TO P.
   MOVE 'XXXXXX' TO L-FLD.  *> writes WS-REC bytes 4..9
   SET P DOWN BY 4.        *> back to the start of WS-REC

```

Nuance — arithmetic stays within the pointer's current target buffer. Bricks pointers are interned (buffer, offset) handles, not flat region addresses, so UP/DOWN BY adjusts the offset *inside the same buffer* the pointer already names; it does not let you walk across unrelated allocations. Moving a pointer past the end of its buffer is allowed at SET time (the address is interned, never validated), but a **dereference** through it raises a clean bounds abend — exactly as with any out-of-range based access. Moving below offset 0, or arithmetic on a NULL/never-set pointer, raises a clean diagnostic at SET time. To walk an array inside a GETMAIN buffer you may still prefer a based 01 with an OCCURS table and subscript it, which keeps the bounds check per element.

See [EXEC CICS ADDRESS](#) and [GETMAIN / FREEMAIN](#) in Chapter 6 for the verbs that produce pointer handles, and runtime/cobol/gmpt.cob (TRANSID GMPT) for a worked example.

Chapter 23. The EIB block in COBOL

EIBRESP, EIBRESP2, EIBAID, EIBCPOSN, EIBCALEN, EIBTRMID, RC, and DFHCOMMAREA are auto-injected if not declared, so most programs need no boilerplate. The same goes for the SQLCA fields (SQLCODE, SQLSTATE, SQLERRMC, ... — see Chapter 26) and the JSON-CODE / JSON-STATUS special registers set by JSON GENERATE and JSON PARSE (see [JSON GENERATE](#) and [JSON PARSE](#)). After every EXEC CICS, EIBRESP and EIBRESP2 are populated by the handler the same way they are for REXX. The legacy "test EIBRESP after every verb" pattern is the most common idiom; for non-local error handling without a REXX-style SIGNAL ON ERROR, use EXEC CICS HANDLE CONDITION / EXEC CICS HANDLE ABEND ([Chapter 10](#)).

See [Chapter 11](#) for the field meanings.

Chapter 24. EXEC CICS in COBOL

The COBOL parser collects every token between EXEC CICS and END-EXEC and reconstructs the body verbatim for the same cics.ParseCommand REXX uses. Three consequences:

- **Map-field names must match group-child names exactly.** When EXEC CICS SEND MAP('CUST1') FROM(SCR) fires, the CICS handler asks the frame for SCR.INFOLINE, SCR.MSG, etc. — so SCR must have children spelled exactly as the map declares fields. Real COBOL solves this with BMS-generated copybooks; bricks doesn't ship those, so the operator declares fields by hand.
- **DFHCOMMAREA is the COMMAREA marshalling slot.** Caller passes bytes in via the frame; sub-program reads with MOVE DFHCOMMAREA TO ..., mutates a working copy, and MOVE ... TO DFHCOMMAREA. Trailing space is stripped by the dispatcher on the way back out.
- **Command-line arguments arrive via EXEC CICS RECEIVE INTO(...).** When the operator types EXAM 1 2 3 at the blank prompt, the dispatcher hands the unedited line (TRANSID prefix included) to the

first task in the chain. The program reads it with `EXEC CICS RECEIVE INTO(W$-INPUT) LENGTH(W$-LEN) END-EXEC` and parses with `UNSTRING W$-INPUT DELIMITED BY ' ' INTO TRANSID A B C END-UNSTRING`. Single-shot per task, then `E0C`. See `runtime/cobol/exam.cob`.

`EXEC CICS SEND TEXT FROM(area) [ERASE]` works the same way. Reconstruction routes through the shared `cics.Handler`, so any verb REXX can issue, COBOL can issue too. The body is laid out as a flat row-major buffer (every `cols` bytes = one row, default 80) — 3270 has no LF, so build area as a group of `PIC X(80)` children and fill them row by row, or compose a single `PIC X(n*80)` with a matching `MOVE` per row.

The full `EXEC CICS` verb set is documented in [Part 2](#); behaviour is identical between REXX and COBOL.

Chapter 25. Copybooks

A copybook is a fragment of COBOL source — typically declarations for constants and groups — kept in a separate file and pulled into a program at parse time with a single line:

```
WORKING-STORAGE SECTION.  
COPY DFHAID.  
COPY DFHRESP.  
01 OWN-COUNTER PIC 9(4).
```

`bricks` expands every `COPY name.` directive at the source-text level before the lexer runs, so the rest of the parser sees one unbroken stream of COBOL.

The `COPY` directive

Syntax:

```
COPY <name>.
```

- The directive must be on its own line. Anything before or after the `COPY` keyword on the same line is a parse error.
- `<name>` is the basename of a copybook file; leave the extension off. The lookup tries `<name>.cpy`, then `<name>.cbl`, then the exact `<name>`, all case-insensitive.
- The trailing period is required.
- Leading whitespace is tolerated, and an `*>` inline comment on the same line is fine: `COPY DFHAID. *> 3270 AID bytes.`
- The directive is case-insensitive (`COPY`, `Copy`, `copy` all work), and so is the name (`COPY DFHAID.` and `copy dfhaid.` resolve to the same file).

Where `bricks` searches

The search directory is configured by `copybook_dir=` in `bricks.cnf`. The default is `runtime/cobolcopy/` (auto-derived from `runtime_dir` when that line is set). The standard CICS copybooks ship under this directory, so a fresh install already has `DFHAID.cpy` and `DFHRESP.cpy` available.

Copybook names are restricted to a flat alphanumeric namespace: `A-Z`, `a-z`, `0-9`, `-`, `_`, and `..`. Names containing a path separator, a `..` sequence, or a leading dot are rejected before any filesystem access. A copybook itself may contain `COPY` directives; nesting is capped at depth 5 and cycles are detected explicitly.

Delivered copybook: DFHAID

COPY DFHAID. brings in the standard 3270 attention-identifier (AID) byte constants. Every AID is declared **twice**, once under the friendly bricks mnemonic and once under the IBM-traditional DFH alias — both names resolve to the same byte, and a program may use either style or mix them.

bricks mnemonic	IBM alias	AID byte	Key
ENTER	DFHENTER	X'7D'	Enter
CLEAR	DFHCLEAR	X'6D'	Clear
PA1	DFHPA1	X'6C'	PA1
PA2	DFHPA2	X'6E'	PA2
PA3	DFHPA3	X'6B'	PA3
PF01	DFHPF1	X'F1'	PF1
PF02	DFHPF2	X'F2'	PF2
PF03	DFHPF3	X'F3'	PF3
PF04	DFHPF4	X'F4'	PF4
PF05	DFHPF5	X'F5'	PF5
PF06	DFHPF6	X'F6'	PF6
PF07	DFHPF7	X'F7'	PF7
PF08	DFHPF8	X'F8'	PF8
PF09	DFHPF9	X'F9'	PF9
PF10	DFHPF10	X'7A'	PF10
PF11	DFHPF11	X'7B'	PF11
PF12	DFHPF12	X'7C'	PF12
PF13	DFHPF13	X'C1'	PF13
PF14	DFHPF14	X'C2'	PF14
PF15	DFHPF15	X'C3'	PF15
PF16	DFHPF16	X'C4'	PF16
PF17	DFHPF17	X'C5'	PF17
PF18	DFHPF18	X'C6'	PF18
PF19	DFHPF19	X'C7'	PF19
PF20	DFHPF20	X'C8'	PF20
PF21	DFHPF21	X'C9'	PF21
PF22	DFHPF22	X'4A'	PF22
PF23	DFHPF23	X'4B'	PF23
PF24	DFHPF24	X'4C'	PF24

The bricks mnemonic uses uniform-width PF01..PF24; the IBM alias keeps the no-leading-zero form (DFHPF1..DFHPF24) everyone recognises from real CICS code.

Delivered copybook: DFHRESP

COPY DFHRESP. brings in the EXEC CICS condition-code constants for EIBRESP. Each code is declared twice as well — RESP-X mnemonic and DFHRESP-X traditional alias. Numeric values match the runtime's emitter

table (`cics/resp.go`), so any code bricks returns from a verb has a named constant here.

The most useful entries:

bricks mnemonic	IBM alias	EIBRESP	Condition
RESP-NORMAL	DFHRESP-NORMAL	0	success
RESP-ERROR	DFHRESP-ERROR	1	generic error
RESP-EOC	DFHRESP-EOC	6	end-of-chain / nothing to RECEIVE
RESP-NOTFND	DFHRESP-NOTFND	13	record / item / file not found
RESP-DUPREC	DFHRESP-DUPREC	14	duplicate record on WRITE
RESP-DUPKEY	DFHRESP-DUPKEY	15	duplicate key in browse
RESP-INVREQ	DFHRESP-INVREQ	16	invalid request (bad args, sandbox violation)
RESP-IOERR	DFHRESP-IOERR	17	underlying I/O failure
RESP-NOSPACE	DFHRESP-NOSPACE	18	TS / TD queue full
RESP-NOTOPEN	DFHRESP-NOTOPEN	19	file not open / dataset not enabled
RESP-ENDFILE	DFHRESP-ENDFILE	20	end of browse
RESP-LENGERR	DFHRESP-LENGERR	22	length mismatch
RESP-QZERO	DFHRESP-QZERO	23	TS / TD queue empty (READQ TD returns this)
RESP-ITEMERR	DFHRESP-ITEMERR	26	TS item number out of range
RESP-PGMIDERR	DFHRESP-PGMIDERR	27	LINK / XCTL program not found
RESP-MAPFAIL	DFHRESP-MAPFAIL	36	RECEIVE MAP with no input
RESP-INVMP SZ	DFHRESP-INVMP SZ	38	map size mismatch
RESP-QIDERR	DFHRESP-QIDERR	44	TS queue id unknown
RESP-ROLLEDBACK	DFHRESP-ROLLEDBACK	82	unit-of-work rolled back

The full list (every code bricks emits) is in `runtime/cobolcopy/DFHRESP.cpy`. Open the file to see the codes not summarised above (RESP-ILLOGIC, RESP-SIGNAL, RESP-QBUSY, the various RESP-INV... codes, RESP-SYSIDERR, etc.).

Delivered copybook: SQLCA

COPY SQLCA. brings in named constants for SQLCODE and the most common SQLSTATE values, so embedded-SQL programs can write:

```
COPY SQLCA.
...
EXEC SQL SELECT name INTO :NM
        FROM customers WHERE id = :CUSTID END-EXEC.
EVALUATE SQLCODE
    WHEN SQL-OK                PERFORM SHOW-ROW
    WHEN SQL-NODATA            PERFORM SHOW-NOT-FOUND
    WHEN SQL-MULTIPLEROWS     PERFORM SHOW-AMBIGUOUS
    WHEN OTHER                 PERFORM SHOW-SQL-ERROR
END-EVALUATE.
```

The SQLCA fields themselves (SQLCODE, SQLSTATE, SQLERRMC) are auto-injected by the bricks COBOL parser — no COPY needed for the data items. SQLCA only adds the constants programs compare against.

bricks mnemonic	IBM alias	SQLCODE	Meaning
SQL-OK	DB2-SUCCESS	0	success
SQL-NODATA	DB2-NOTFOUND	+100	no row / end-of-cursor
SQL-NOTFND	DB2-NOTFOUND	+100	alias
SQL-NOCONFIG	—	-1	SQL not configured in bricks.cnf
SQL-DDLREJECTED	—	-2	DDL rejected (CEDA-only)
SQL-GENERIC	—	-100	generic PG error (see SQLERRMC + log)
SQL-MULTIPLEROWS	DB2-DUPLICATE	-811	SELECT INTO returned >1 row

SQLSTATE constants are PIC X(5) and compare directly:

```
IF SQLSTATE = SQLSTATE-DUP-KEY THEN
  MOVE 'Already exists.' TO MSG
END-IF.
```

Constant	SQLSTATE	Class
SQLSTATE-OK	00000	success
SQLSTATE-NODATA	02000	no-data
SQLSTATE-WARNING	01000	warning
SQLSTATE-NOT-NULL	23502	NOT NULL constraint violated
SQLSTATE-FK-VIOL	23503	foreign-key constraint violated
SQLSTATE-UQ-VIOL	23505	UNIQUE constraint violated
SQLSTATE-DUP-KEY	23505	alias
SQLSTATE-CHK-VIOL	23514	CHECK constraint violated
SQLSTATE-CONN-FAIL	08001	unable to connect to PG
SQLSTATE-CONN-LOST	08006	connection dropped mid-flight
SQLSTATE-INV-AUTH	28000	PG authentication failed
SQLSTATE-INV-PWD	28P01	wrong password
SQLSTATE-SYNTAX	42601	SQL syntax error
SQLSTATE-UNDEF-TBL	42P01	table doesn't exist
SQLSTATE-UNDEF-COL	42703	column doesn't exist
SQLSTATE-UNDEF-FN	42883	function doesn't exist
SQLSTATE-AMBIG-COL	42702	ambiguous column reference
SQLSTATE-INSUF-PRIV	42501	permission denied

The bricks SQL executor also logs every non-zero SQLCODE's full PG error text via `brickslog.Sys` (one line on the operator console + the per-run log file). Programs render only the short mnemonic-driven message on the 3270 screen; the diagnostic detail lives in the log so PG hostnames / port numbers / wrapper text don't leak to terminal users.

Idiom: replacing raw literals

Before — opaque hex and bare integers, easy to mis-key:

```

IF EIBAID = X'F3' THEN
    MOVE 'Y' TO EXIT-FLAG
END-IF.
IF EIBRESP = 13 THEN
    MOVE 'Customer not found.' TO MSG
END-IF.
IF EIBRESP = 23 THEN          *> TD queue empty
    MOVE 'Y' TO DONE-FLAG
END-IF.

```

After — names a reader can scan:

```

WORKING-STORAGE SECTION.
COPY DFHAID.
COPY DFHRESP.
...
IF EIBAID = PF03 THEN          *> or DFHPF3 if you prefer
    MOVE 'Y' TO EXIT-FLAG
END-IF.
IF EIBRESP = RESP-NOTFND THEN
    MOVE 'Customer not found.' TO MSG
END-IF.
IF EIBRESP = RESP-QZERO THEN
    MOVE 'Y' TO DONE-FLAG
END-IF.

```

Live examples: `qagc.cob` (PF03 cancel branch), `gusl.cob` (EVALUATE EIBAID with PF03 / PF07 / PF08 / PF12), `ordr.cob` (RESP-QZERO and RESP-DUPREC against READQ TD / WRITE FILE), `exam.cob` (RESP-E0C after RECEIVE). All under `runtime/cobol/`.

Restrictions

The bricks COPY preprocessor is deliberately minimal. The following real-CICS extensions are **not** supported:

- `COPY ... REPLACING ==X== BY ==Y==.` — token substitution. Rare in modern code; deferred until a concrete use case appears.
- `COPY xyz OF DFHCOB. / COPY xyz IN libname.` — library qualifier. bricks has one search root per process.
- `COPY xyz SUPPRESS.` — suppress copybook listing. bricks produces no listing to suppress.
- Embedded directives: multiple statements on one line including a `COPY` are rejected. The line must be `COPY <name>.` and nothing else.

If a copybook is edited, programs that include it will pick up the new contents only after the program's own source mtime changes (the program cache keys by the program path + mtime, not the copybook's). A trivial way to force a reload is to touch the COBOL file:

```
touch runtime/cobol/qagc.cob
```

Or use `CEMT P C` from a TSO/3270 session to flush the entire program cache.

Part 6. EXEC SQL command reference

Part 6 begins here. The chapter that follows is the complete EXEC SQL surface — identical for COBOL and REXX. It covers every embedded-SQL verb Bricks supports, the SQLCA layout and SQLCODE catalogue, the WHENEVER declarative error handler, cursor lifecycle, per-task CONNECT TO, null indicators, and SYNCPOINT integration with the bricks unit of work.

Chapter 26. Embedded SQL (COBOL and REXX)

Programming Note. Although this chapter was originally written against COBOL, every EXEC SQL shape and behaviour described here applies identically to REXX. The REXX-specific notes appear later in the chapter under *REXX equivalent*. Programs in either language read and write the SQLCA fields (SQLCODE, SQLSTATE, SQLERRMC) the same way the COBOL samples do.

bricks accepts the classic IBM-style EXEC SQL ... END-EXEC embedded-SQL surface and runs it against the Postgres connection configured via db_* in bricks.cnf (see [SQL support — Phase 1](#)). Programs read and write rows through prepared statements; the SQLCA fields SQLCODE, SQLSTATE, SQLERRMC are auto-injected into every program's DATA DIVISION so the program can test them without declaring them.

Scope: single-row CRUD (SELECT INTO, INSERT, UPDATE, DELETE, COMMIT, ROLLBACK), the four-verb cursor lifecycle (DECLARE / OPEN / FETCH / CLOSE), per-task CONNECT TO 'name' database switching, null indicators (:hostvar :indicator), and the WHENEVER declarative error-handling directive. The same surface works identically in COBOL and REXX.

Format

```
EXEC SQL <statement> END-EXEC.
```

Where <statement> is any of the supported verbs (below). Host variables are written with a leading colon: :CUSTID, :NM, :BALANCE. The variable name must already exist in DATA DIVISION — the SQL executor reads its current value (for INPUT bindings) and writes back (for SELECT INTO targets).

Supported verbs

Verb	Notes
SELECT cols INTO :v, :v FROM ...	Single-row read. Returns +100 (no-data) when 0 rows, -811 when >1 row.
INSERT INTO t (...) VALUES (...)	Host-variable references bound as \$N placeholders.
UPDATE t SET ... WHERE ...	Returns RESP-NORMAL even when 0 rows match (matches DB2).
DELETE FROM t WHERE ...	Same row-count semantics as UPDATE.
COMMIT WORK	Commits the per-task PG transaction. Also fired by EXEC CICS SYNCPOINT.
ROLLBACK WORK	Rolls back. Also fired by EXEC CICS SYNCPOINT ROLLBACK.

DDL (CREATE DATABASE, DROP DATABASE, CREATE USER, DROP USER, ALTER ROLE, CREATE TABLESPACE, ...) is **rejected** with SQLCODE = -2 and SQLSTATE = 42501. CEDA DATABASE owns those operations and audits each one through brickslog.Audit. CREATE TABLE / ALTER TABLE / similar non-privileged DDL is not blocked but is also unsupported in the executor's statement classification — porters should run schema migrations with psql.

SQLCODE catalog

bricks adopts the IBM DB2 convention adapted to Postgres. The catalog below covers what programs typically test:

SQLCODE	SQLSTATE	Meaning
0	00000	Success.
+100	02000	No data: SELECT INTO returned 0 rows, or an update/delete affected 0 rows.
-1	empty	No database configured (no db_* lines in bricks.cnf).
-2	42501	DDL rejected (CEDA owns CREATE/DROP DATABASE
-100	varies	Generic SQL failure — PG returned an error not specifically mapped.
-811	21000	SELECT INTO returned more than one row.

SQLERRMC always carries the human-readable message from Postgres (truncated to 70 chars to match the PIC X(70) auto-injection). When the error path includes a specific SQLSTATE, that 5-char code lands in SQLSTATE; otherwise SQLSTATE is the relevant standard code (00000 on success, 02000 on no-data, HV000 when the failure didn't carry a state).

SYNCPOINT integration

Each task lazily begins a Postgres transaction on its first EXEC SQL that touches data. The transaction commits when:

- EXEC CICS SYNCPOINT runs.
- EXEC SQL COMMIT WORK runs.
- The task ends cleanly (RETURN, GOBACK, STOP RUN).

And rolls back when:

- EXEC CICS SYNCPOINT ROLLBACK runs.
- EXEC SQL ROLLBACK WORK runs.
- The task abends via EXEC CICS ABEND.

Postgres-side commit fires *before* the bricks-side bbolt journal commit. That ordering avoids a hung PG commit leaving the bricks journal already finalised; failures during the SQL commit surface as ROLLEDBACK from SYNCPOINT, and the bricks journal then rolls back too.

Idiomatic test

```
EXEC SQL
  SELECT name INTO :NM
  FROM customers_sql
  WHERE id = :CUSTID
END-EXEC.

EVALUATE SQLCODE
  WHEN RESP-NORMAL    PERFORM SHOW-NAME      *> +0
  WHEN RESP-NOTFND    PERFORM SHOW-NOTFOUND   *> wait -- that's an EIBRESP code; SQLCODE
uses RESP-NORMAL/+100/etc.
  WHEN OTHER          PERFORM SHOW-SQL-ERR
END-EVALUATE.
```

The RESP-NORMAL constant (from COPY DFHRESP) is value 0, the same as the SQLCODE-OK value, so it conveniently doubles as the success test for both EIBRESP (EXEC CICS) and SQLCODE (EXEC SQL). For the no-data case, test SQLCODE = +100 directly — there's no DFHSQLCODE copybook yet.

Worked example: SQLD

runtime/cobol/sql.d.cob (TRANSID SQLD) is the bundled demonstration. It reads a single customers_sql row by id and renders the name + SQLCA verdict on screen SQLD1. The expected schema is documented in the program header. The operator runs psql once to seed the table:

```
CREATE TABLE customers_sql (  
  id    text PRIMARY KEY,  
  name  text NOT NULL  
);  
INSERT INTO customers_sql VALUES  
  ('K001', 'Alice'),  
  ('K002', 'Bob'),  
  ('K003', 'Carol');
```

Then bricks → sign on → SQLD → type K001 → ENTER. The operator sees Alice in the NM cell with SQLCODE = 0 and SQLSTATE = 00000.

REXX equivalent

REXX programs use the same EXEC SQL shape as COBOL, with :name binding REXX variables. There's no DATA DIVISION analog -- REXX is dynamically typed -- so SQLCODE, SQLSTATE, SQLERRMC arrive as ordinary REXX variables the program reads after each statement:

```
ADDRESS CICS  
  
CUSTID = 'K001'  
EXEC SQL  
  SELECT name INTO :NM  
  FROM customers_sql  
  WHERE id = :CUSTID  
END-EXEC  
  
IF SQLCODE = 0 THEN  
  SAY 'Customer:' NM  
ELSE IF SQLCODE = 100 THEN  
  SAY 'Not found'  
ELSE  
  SAY 'SQL error' SQLCODE ':' SQLERRMC
```

Behind the scenes the REXX preprocessor (rexex.PreprocessExecCICS) rewrites the EXEC SQL block into a sentinel bare-string that runCommand recognises and routes to the same per-task SQL executor COBOL uses. The two language surfaces are parity-clean: the executor doesn't know or care which language drove it. Live sample: runtime/rexx/sqlr.rexx, TRANSID SQLR.

REXX stem-tail gotcha. Classic REXX resolves STEM.NAME by looking up the simple symbol NAME and using its value as the tail. If a program does EXEC SQL ... INTO :NM then writes SCR.NM = NM, the LHS resolves to SCR.<value-of-NM> = SCR.Alice — the value lands in the wrong slot. Workaround: pick a host-var name that is NOT also a stem tail. The bundled SQLR uses CUSTNM instead of NM for exactly this reason.

WHENEVER (declarative error handling)

Instead of testing `SQLCODE` after every statement, a program can declare a standing rule with `EXEC SQL WHENEVER`:

```
EXEC SQL WHENEVER SQLERROR  GO TO SQL-ERROR-EXIT END-EXEC.
EXEC SQL WHENEVER NOT FOUND GO TO NO-MORE-ROWS  END-EXEC.

EXEC SQL SELECT name INTO :NM FROM customers_sql
        WHERE id = :CUSTID END-EXEC.
*> no SQLCODE test here -- WHENEVER handles it
```

Three conditions, checked after every subsequent `EXEC SQL`:

Condition	Fires when
SQLERROR	SQLCODE is negative
NOT FOUND	SQLCODE = +100
SQLWARNING	SQLWARN0 = 'W' (bricks rarely sets this — PG errors rather than warns)

Each condition takes one of two actions:

- `CONTINUE` — ignore the condition, fall through (the default for any condition never declared).
- `GO TO label / GOTO label` — branch to the paragraph (COBOL) or label (REXX). The branch unwinds any enclosing `PERFORM` exactly as an explicit `GO TO` would.

A later `WHENEVER` for the same condition replaces the earlier one (e.g. declare `GO TO` for a block of statements, then `CONTINUE` to resume manual handling).

REXX uses the identical syntax; the branch is a `SIGNAL` to the named label:

```
EXEC SQL WHENEVER SQLERROR GOTO SQLERR END-EXEC
EXEC SQL SELECT name INTO :CUSTNM FROM customers_sql
        WHERE id = :CUSTID END-EXEC
SAY 'customer:' CUSTNM
EXIT
SQLERR:
SAY 'SQL failed, SQLCODE' SQLCODE
EXIT
```

Scoping caveat. bricks's `WHENEVER` is *execution-order* scoped: the directive set by the most recently **executed** `EXEC SQL WHENEVER` governs later statements. Real DB2 scopes *lexically* (by source position) because its precompiler inlines a `SQLCODE` test after every statement. For the dominant pattern — one `WHENEVER` near the top of the program — the two are identical. Programs that re-declare `WHENEVER` inside conditional branches see execution-order semantics; keep `WHENEVER` declarations at the top of a paragraph to avoid surprises.

Cursors (DECLARE / OPEN / FETCH / CLOSE)

Multi-row result sets use the four-verb cursor lifecycle. The executor parks one `*sql.Rows` per cursor per task; cursors are closed automatically on `COMMIT`, `ROLLBACK`, or task end so a program that forgets `CLOSE` doesn't pin Postgres MVCC.

```
EXEC SQL DECLARE C1 CURSOR FOR
    SELECT id, name FROM customers_sql ORDER BY id
END-EXEC.

EXEC SQL OPEN C1 END-EXEC.

PERFORM UNTIL SQLCODE = 100
    EXEC SQL FETCH C1 INTO :ID, :NM END-EXEC
    IF SQLCODE = 0 THEN
        DISPLAY 'row: ' ID ' ' NM
    END-IF
END-PERFORM.

EXEC SQL CLOSE C1 END-EXEC.
```

REXX uses the same four verbs (the preprocessor routes them all through the same executor):

```
EXEC SQL DECLARE C1 CURSOR FOR SELECT id, name FROM customers_sql END-EXEC
EXEC SQL OPEN C1 END-EXEC
DO FOREVER
    EXEC SQL FETCH C1 INTO :ID, :CUSTNM END-EXEC
    IF SQLCODE = 100 THEN LEAVE
    SAY 'row:' ID CUSTNM
END
EXEC SQL CLOSE C1 END-EXEC
```

End-of-data is `SQLCODE = +100`. A closed cursor stays in the registry, so `OPEN c1` can rewind it without re-`DECLARE`.

Database binding (transactions.conf 5th field)

Every transaction starts on exactly one database. Which one is fixed by the 5th colon-separated field of its row in `runtime/transactions.conf`:

```
BANK:cobol:bank.cob:public,users,admin:bank
```

The grammar of the file is therefore:

```
transid:lang:program[:groups[:database]]
```

- **database** -- name of a row in `runtime/databases.conf`. The match is **case-insensitive**: both sides are folded to lower case, the way Postgres folds unquoted identifiers, so `:BANK` and `:bank` reach the same database. (They used to be two different records opening two pools onto one physical database, and a `:BANK` binding against a `bank` row silently failed at run time with `SQLCODE = -1`.) Whitespace around the name is trimmed, and the name may not contain `:`, a newline or a whitespace-preceded `#` -- `CEDA TRANSACTION` refuses those, because they produce a line the parser reads back as a different record.
- **Absent 5th field** -- the transaction binds to the **first** row of `databases.conf`. The first row is the deployment-wide default; `CEDA DATABASE` labels it `(def)` and refuses `delete` on it.
- **Unknown name** -- the dispatcher logs the misbinding at task start and the first `EXEC SQL` returns `SQLCODE = -1` with an empty `SQLSTATE`.

A transaction whose program never runs `EXEC SQL` may leave the field empty -- the implicit default never gets touched. A transaction whose program **does** run `EXEC SQL` against a non-default database **must** list

the binding explicitly: omitting it sends every statement to the default pool, which typically does not host the requested table. The operator-visible symptom is a flood of paired errors:

```
SQL error SQLCODE=-204 SQLSTATE=42P01: relation "accounts" does not exist
SQL error SQLCODE=-100 SQLSTATE=25P02: current transaction is aborted, commands ignored until end of transaction block
```

The -204 is the first statement landing on the wrong pool; the -100 / 25P02 cascade is Postgres marking the per-task PG transaction aborted, so every follow-up statement under a `WHENEVER SQLERROR CONTINUE` declaration logs the same pair until `SYNCPOINT` runs at task end. Fix: add the correct 5th field and restart the task.

CEDA TRANSACTION (A / U actions) and operator-vi edits of `transactions.conf` both honour the field; the file is hot-reloaded on mtime change.

CONNECT TO 'name' (per-task database switch)

A program that needs to touch a second database during a task issues `EXEC SQL CONNECT TO 'name'`. The named database must exist in `databases.conf` (the CEDA DATABASE catalogue); an unknown name returns `SQLCODE = -1`. The name is matched case-insensitively — it is folded to lower case, like every other database-name path in bricks — so `CONNECT TO 'ORDERS'` finds the `orders` row. An in-flight PG transaction on the previous database is committed implicitly (matches DB2's behaviour) -- programs that want to discard the previous work should `EXEC SQL ROLLBACK` before connecting.

```
EXEC SQL CONNECT TO 'orders' END-EXEC.
EXEC SQL INSERT INTO order_lines VALUES (:ID, :QTY) END-EXEC.
EXEC SQL COMMIT END-EXEC.
EXEC SQL CONNECT TO 'customers' END-EXEC.
EXEC SQL SELECT email INTO :EM FROM customers WHERE id = :CID END-EXEC.
```

Implicit binding still applies to the first statement of every task -- the transaction's Database column in `transactions.conf` picks the starting pool. `CONNECT TO` is only needed when the program needs to move between databases inside a single task.

Null indicators (:hostvar :indicator)

DB2 programs use a second host var beside the value to signal NULL. Bricks supports the same pair syntax in both COBOL and REXX:

```
EXEC SQL
  SELECT name INTO :NM :NIND
  FROM customers_sql WHERE id = :CUSTID
END-EXEC.

EVALUATE NIND
  WHEN 0   DISPLAY 'Name: ' NM
  WHEN -1  DISPLAY 'Name is NULL'
END-EVALUATE.
```

Output side (SELECT INTO / FETCH INTO):

- Column was non-NULL → indicator = 0, value holds the result.
- Column was NULL → indicator = -1, value is cleared to the empty string so a program that ignored its indicator can't read a stale value from a previous statement.

Input side (WHERE / VALUES / SET clauses):

```
MOVE -1 TO NIND.  
EXEC SQL  
    INSERT INTO customers_sql (id, name) VALUES (:K, :NM :NIND)  
END-EXEC.
```

- Indicator = -1 at bind time → bricks passes NULL to PG, regardless of what NM currently holds.
- Indicator = 0 (or unset) → the value var's frame contents bind normally.

Pair detection: a value/indicator pair is two :name tokens separated only by whitespace. :NM, :NIND (with a comma) means two independent host vars, not a pair. The INDICATOR keyword form (:NM INDICATOR :NIND) is not currently supported -- use the juxtaposed form.

Column-value coercion

PG's wire format doesn't always match what COBOL or REXX want to read. Bricks normalises three cases before the value reaches the program's host variable:

- **bool columns** -- PG returns "t" / "f" (or "true" / "false"). Bricks converts to "1" / "0" so COBOL PIC X(1) flags and REXX IF VAR = 1 idioms work the way they do in DB2.
- **NUMERIC(p, s) columns** -- PG drops trailing fractional zeros for computed expressions (SELECT 1.5 returns "1.5" even though the result type has scale 2). Bricks pads the value to the column's declared scale, so a COBOL PIC 9(3)V99 target receives "1.50" → digits "150" → stored as 001.50, not 000.15.
- **CHAR(N) columns** -- PG returns space-padded values. Bricks trims trailing spaces so a REXX literal compare (IF VAR = 'X') doesn't fail because the CHAR(5) value was "X ". VARCHAR is left alone (PG never pads it; its trailing spaces are real).

Other types (INT, BIGINT, TEXT, DATE, TIMESTAMP, JSON, ...) flow through verbatim. COBOL's MOVE semantics handle integer sign and zero-pad automatically; REXX is dynamically typed so the string form works for both display and arithmetic.

Restrictions

Out of scope	Why
Stored procedure calls (EXEC SQL CALL)	Out of scope — porters use plain SELECT-from-function.
Multi-row INSERT VALUES (...), (...)	Out of scope; rewrite as one statement per row.
CREATE / `DROP DATABASE	USER`
Long-running cursors across EXEC CICS RETURN chains	Out of scope; each task gets a fresh PG connection.

Chapter 27. Restrictions and deferred features

This chapter is the canonical list of language features bricks does **not** accept. Anything not listed here can be assumed implemented; refer back to the relevant Chapter 14–26 section for the supported syntax.

COBOL — disallowed at parse time

- **Calculated GOTOs.** GO TO ... DEPENDING ON expr is rejected. Every task has its own heap and stack with no static control-block aliasing; supporting calculated GOTOs would require a per-program

label-table the parser deliberately doesn't build.

- **Nested OCCURS.** An array of arrays is rejected at parse time. One level of OCCURS is fine.
- **EXIT SECTION.** bricks has no PROCEDURE DIVISION sections.
- **Forms IBM itself forbids.** These are rejected with a short message rather than run with a guessed meaning: a literal, figurative constant or function as the UNSTRING source; a reference-modified STRING INTO receiver; DELIMITER IN / COUNT IN without DELIMITED BY; WHEN 'A' OR 'B' (stack two WHENS); an EVALUATE WHEN whose ALSO count differs from the subjects; NOT NOT ...; and an abbreviated relation that opens a parenthesised group (A = B AND (< C OR D)). See [Chapter 22](#) for each.

COBOL — deferred (would be implementable, not yet wired)

Area	Specifics
DATA DIVISION	USAGE COMP-1 (single float), USAGE COMP-2 (double float), SYNC, INDEXED BY, OCCURS DEPENDING ON, 66 / 78 levels, BLANK WHEN ZERO, and IBM's extension allowing a level-01 redefiner <i>larger</i> than its target (bricks requires the redefining item to fit at every level). USAGE COMP / COMP-4 / COMPUTATIONAL / BINARY, USAGE COMP-3 / COMPUTATIONAL-3 / PACKED-DECIMAL, and REDEFINES ARE supported — see Chapter 21.
Edited PIC	Edit characters +, -, CR, DB, B, 0, /. (Z, ., ,, \$, * are supported.)
PROCEDURE DIVISION	CALL (external program), SEARCH / SEARCH ALL, ALTER, INITIALIZE, ACCEPT, GO TO ... DEPENDING ON, DIVIDE ... REMAINDER. In-line PERFORM ... END-PERFORM, PERFORM ... THRU, WITH TEST BEFORE/AFTER, VARYING ... AFTER, EXIT PERFORM [CYCLE], EXIT PARAGRAPH, full EVALUATE, parenthesised and abbreviated combined conditions, the full IBM STRING / UNSTRING (POINTER, OVERFLOW, DELIMITER IN, COUNT IN, TALLYING IN) and JSON GENERATE / JSON PARSE ARE supported — see Chapter 22 .
Reference modification	Two forms are rejected at parse time: a slice on an OF / IN qualifier (K(1:2) OF ROW), and a subscript combined with a slice in one reference (TBL(idx)(start:length)). Slice the whole reference, or MOVE the OCCURS slot to a work field first. Plain identifier(start:length) on either side of a statement IS supported — see Chapter 22 .
Intrinsic functions	Only UPPER-CASE, LOWER-CASE, LENGTH, NUMVAL, TRIM, REVERSE, POS are implemented. The statistical / date / time families are not. FUNCTION TRIM does NOT accept the LEADING / TRAILING modifier.
JSON	JSON GENERATE / JSON PARSE phrases ENCODING, IGNORING, INDICATING and CONVERTING ... FROM JSON NULL; the single condition-name form of JSON PARSE ... CONVERTING cannot set "false" (no 88 ... WHEN SET TO FALSE). The statements themselves ARE supported — see JSON GENERATE and JSON PARSE .
Copybooks	COPY ... REPLACING ==X== BY ==Y==., library qualifiers (COPY name OF lib), SUPPRESS.
EXEC SQL	Multi-row INSERT VALUES (...), (...). Stored-procedure calls (EXEC SQL CALL). DDL inside the executor — use psql or CEDA DATABASE.
EXEC CICS	SCREENHT-based map family suffix (e.g. CUST1L on a mod-4 screen) is REXX-only; COBOL twins always render the unsuffixed mod-2 map.

REXX — deferred / accepted but inert

Area	Specifics
Tracing	TRACE is parsed but emits no trace output.
Settings	OPTIONS ... parsed and accepted; recognised options are none. NUMERIC FORM SCIENTIFIC / ENGINEERING accepted syntactically, rendering is always plain decimal.
External routines	No dynamic linking — every callable is either a PROCEDURE in the program or a built-in.

Area	Specifics
Address environments	Only ADDRESS CICS is wired for host-command dispatch. ADDRESS COMMAND / ADDRESS SYSTEM / ADDRESS TSO do parse, and ADDRESS () reports whichever name is current, but the first host command issued under one fails with ADDRESS handler not registered.
Terminal queue	PUSH, QUEUE, PULL, PARSE PULL are parsed; the queue is always empty. Use EXEC CICS RECEIVE for terminal input.
PARSE	PARSE SOURCE, PARSE VERSION, PARSE NUMERIC are not wired. PARSE LOWER is not parsed (only UPPER).
Conditions	FAILURE, NOTREADY, LOSTDIGITS are not implemented. HALT parses but never fires (bricks has no external-halt fan-in).
Numeric	Arithmetic is float64 internally, then rounded to NUMERIC DIGITS significant figures — strict IEEE behaviour is not preserved.
Built-ins	ERRORTEXT (code) echoes the code rather than returning the standard REXX error-text string.
Variables / DROP	Two divergences, both visible through SYMBOL (). First, DROP A.J does not derive the tail: the DROP parser consumes a single identifier and . is an ordinary symbol character, so the literal tail J is dropped where TRL2 drops A.3 — clear the slot with A.J = '' if you need the derived name. Second, after S. = 'd' a DROP S.1 removes the tail but leaves the stem default in force, so S.1 reads back as d and SYMBOL ('S.1') answers VAR where TRL2 answers LIT.

Part 8. Sample programs

Part 8 begins here. Bricks ships a curated set of REXX and COBOL programs that exercise every command family in the earlier parts. Each transaction has been built to be operator-reachable from the blank prompt and self-explaining when run — reading the source is the fastest way to internalise the bricks programming model.

Chapter 28. Pre-installed sample transactions

COBOL

Transid	File	Notes
HELC	runtime/cobol/hello.cob	Hello-world; SEND MAP('HEL01') FROM(SCR), RETURN. Smallest end-to-end demo.
QAGC	runtime/cobol/qagc.cob	COBOL twin of QAGR (REXX). Validates the QAGE1 birthdate, computes age in years and approximate days, sends QAGR1. Pseudo-conversational redisplay of QAGE1 on validation errors.
GUST	runtime/cobol/gust.cob	COBOL CUST. A=Add, Q=Query, U=Update, D=Delete, L=List, S=Search; the S action LINKs to GUSL with the search term in COMMAREA and renders the match count returned.
GUSV	runtime/cobol/gusv.cob	COBOL twin of CUSV. Validates the customer-number COMMAREA (LINK target).
GUSL	runtime/cobol/gusl.cob	COBOL twin of CUSL. Renders the customers file via STARTBR / READNEXT / ENDBR. Blank inbound COMMAREA = paginated all-records list (PF7/PF8). Non-blank inbound COMMAREA = filtered single-page browse: every record is FUNCTION POS-tested against the upper-cased filter, the first 15 matches populate ROW1..ROW15, and the total match count flows back through DFHCOMMAREA so the caller (GUST) can render a summary.
EXAM	runtime/cobol/exam.cob	Worked example of reading the operator's command-line arguments. Type EXAM 1 2 3 at the blank prompt.
ORDR	runtime/cobol/ordr.cob	Conversational import: reads runtime/tmp/orders.sample.txt via READQ TD, parses pipe-delimited rows, and WRITE FILE('ORDERS') keyed on customer-id. Tolerates duplicates (EIBRESP = RESP-DUPREC). Summary screen shows counts. See worked example E .
TIMC	runtime/cobol/timc.cob	Timed-reminder demo: exercises CONVERSE, START (with INTERVAL + FROM), and RETRIEVE end-to-end. Shares tim1.map + tim2.map with TIMR.
SQLD	runtime/cobol/sqld.cob	Embedded-SQL demo: EXEC SQL SELECT name INTO :NM FROM customers_sql WHERE id = :CUSTID. Renders the row + SQLCODE / SQLSTATE / SQLERRMC. Requires Postgres seeded with the schema shown in Chapter 26 .
SQLR	runtime/rexx/sqlr.rexx	REXX twin of SQLD. Shares the customers_sql table; demonstrates how REXX accesses the same EXEC SQL surface (SQLCODE / SQLSTATE / SQLERRMC arrive as plain REXX variables).
GMPT	runtime/cobol/gmpt.cob	COBOL pointer / storage demo: EXEC CICS GETMAIN SET(WS-PTR) FLENGTH(16), SET ADDRESS OF a based LINKAGE 01 onto the buffer, write-and-read through it, then EXEC CICS ADDRESS COMMAREA / ADDRESS EIB. Result painted with SEND TEXT (no map). See GETMAIN/FREEMAIN and the COBOL pointer model .

Transid	File	Notes
TOD0	runtime/cobol/todo.cob	First server-side COBOL WAPI sample: a multi-user to-do REST API under /api/todo/ over the VSAM files TOD0USR / TOD0LST / TOD0ITM. Uses JSON GENERATE / JSON PARSE, EVALUATE ... ALSO, EXIT PARAGRAPH and SYNCPOINT ROLLBACK units of work. Browser front end in runtime/web/todo/, test script scripts/todo_api_test.bash. See COBOL server-side sample .

All five non-trivial COBOL samples (QAGC, GUST, GUSL, ORDR, EXAM) COPY DFHAID and/or COPY DFHRESP instead of hard-coding hex AID bytes or numeric EIBRESP codes. Skim any of them as living examples of the named-constant idiom from [Chapter 25](#).

REXX (selected)

Transid	File	Notes
HEL0	runtime/rexx/hello.rexx	Hello-world with system-info pane on mod 4.
CUST	runtime/rexx/cust.rexx	Full customer maintenance — Add / Query / Update / Delete / List / Search; LINKs to CUSV for validation.
CUSV	runtime/rexx/cusv.rexx	Validation sub-program; called by CUST via EXEC CICS LINK.
CUSL	runtime/rexx/cusl.rexx	Paginated customer list using STARTBR / READNEXT / ENDBR; adapts paging to mod 2 vs mod 4.
QAGE / QAGR	runtime/rexx/qage.rexx / qagr.rexx	Pseudo-conversational chain; QAGE prompts for a birthdate and chains to QAGR to render the result.
PROD / CONS	runtime/rexx/prod.rexx / cons.rexx	TS queue producer / consumer pair. Conversational; PF3 to exit.
GETC	runtime/rexx/getc.rexx	RECEIVE of command-line + READ FILE + SEND TEXT (no map).
INQR	runtime/rexx/inqr.rexx	EXEC CICS INQUIRE demo: direct INQUIRE TASK on this task plus a STARTBROWSE FILE / INQUIRE FILE NEXT / ENDBROWSE loop listing every file's OPEN/ENABLE status and record/key sizes. Result painted with SEND TEXT (no map). See INQUIRE / SET resources .
TIMR	runtime/rexx/timr.rexx	REXX twin of TIMC: START schedules a reminder; RETRIEVE discriminates cold vs scheduled entry. Shares tim1.map + tim2.map with TIMC.
SCHD	runtime/rexx/schd.rexx	Demonstrates three features in one transaction: START ... AFTER SECONDS(n), ADDRESS() and SYMBOL(). Paints one PASS/FAIL row per check, then schedules itself with AFTER SECONDS(10); the timer-fire path reports the payload and issues no further START, so the chain ends on its own. Contrast the plain AFTER SECONDS(...) here with the manual HH/MM/SS packing TIMR needs for INTERVAL(hhmmss).
CHAT	runtime/rexx/chat.rexx	Real-time multi-user chat. Self-refreshes every 2 seconds via EXEC CICS START TRANSID('CHAT') INTERVAL(000002); the tick handler does SEND MAP ... DATAONLY so the operator's in-progress typing at the bottom of the screen is not clobbered by the refresh. Messages persist in the auto-created KSDS file CHATLOG (one record per message, key shape YYYYMMDDHHMMSS-NNNN-TTTT for lexicographic / chronological sort). Adapts between Model 2 (16 history rows, chatm2.map) and Model 4 (35 history rows, chatm4.map) via EXEC CICS ASSIGN SCREENHT. F3 exits cleanly — no further tick is scheduled, the self-refresh chain dies on its own. Colours mirror the original tsu/chat.go palette: BLUE/BRIGHT title, TURQUOISE clock + footer, YELLOW topic, RED status line, GREEN history rows + prompt, WHITE underscored input.

Run any TRANSID by typing it at the blank prompt after CSSN sign-on. Refer to runtime/transactions.conf for the full list and ACL configuration.

Bank-domain bindings. BANK, BALC, BALR, and BRDS carry an explicit :bank 5th field in `transactions.conf` so their EXEC SQL statements land on the retail-banking pool rather than the default bricks database. BALC / BALR / BRDS do not currently issue SQL, but they share the bank-domain map set and are pinned to the same pool for forward compatibility — adding an EXEC SQL later will Just Work without re-binding. See [Chapter 26 — Database binding](#).

Built-in transactions (no entry in `transactions.conf`)

The bricks core dispatches a handful of TRANSIDs directly, without consulting `transactions.conf`. They take precedence over a same-name entry in the table.

Transid	Purpose	Notes
CSSN	Sign-on	Default authentication flow. Configurable via <code>secure_login_transacton</code> in <code>bricks.cnf</code> .
CSSF	Sign-off	CSSF LOGOFF clears the session's identity; bare CSSF is a no-op.
CEMT	Master-operator	INQUIRE / MONITOR / PERFORM trees; CONTROLBLOCKS sub-tree and PERFORM gated on the admin group.
CEDA	Resource definitions	TRANSACTION / PROGRAM / USER screens; admin-only.
CECI	Command-level interpreter	Type one EXEC CICS / EXEC SQL / EXEC WEB command (up to 5 input rows), press PF5, see RESPONSE / RESP2 / LEN / ELAPSED plus any variables the handler set or INTO buffers it filled. The TxCB and <code>cics.Handler</code> are session-scoped — built once on CECI entry, torn down on PF3 / CLEAR / PA1 / disconnect. Each PF5 press is a per-PF5 commit boundary via <code>h.CommitImplicit()</code> on success / <code>h.RollbackImplicit()</code> on failure; SYNCPOINT in bricks does not touch cursors, so the per-task TS read cursor, all open browses, TD handles, WEB sessions, and DOCUMENT tokens survive across PF5 presses (a STARTBR on PF5 #1 drives READNEXT on PF5 #2). Variable frame is also session-scoped. READ UPDATE + REWRITE still cannot span PF5 — the per-PF5 implicit SYNCPOINT releases non-HOLD record locks; see the REWRITE deviation note for the recommended workaround. Verbs that would unwind the CECI task itself (RETURN, XCTL, LINK, ABEND), repaint the screen (SEND MAP, SEND TEXT, RECEIVE MAP, CONVERSE), schedule deferred work (START, RETRIEVE, CANCEL, DELAY), or set per-task trap tables (HANDLE, IGNORE, WHENEVER) are refused with a short red-status message. Every command is wrapped in a 7-second wall-clock cap; a cap-hit <i>poisons</i> the session (every subsequent PF5 short-circuits with TIMEOUT, and the session-end defer skips the handler-owned closers to avoid racing the abandoned goroutine — PF3 is the only exit). Gated on the dev group — and the same gate applies to EXEC CICS LINK PROGRAM('CECI') / XCTL PROGRAM('CECI'), so a non-DEV caller sees PGMIDERR (LINK) or a 403 task-error screen (XCTL) without the built-in's own Run ever being reached. Operator manual: the CECI section in README.md covers layout, key bindings, verb policy, the response-line state mapping, and the EBCDIC-037 input rules.
ISPF	Source editor	Browse and edit the REXX, COBOL, and BMS-map source trees. Gated on the dev group. Operator manual: ISPF_editor.md — covers every PF key, every command-line word, every line-prefix command (D / I / C / M / R / U / L /) / (/ X / O / A / B plus the doubled block forms), the file browser, the warn-then-save flow, multi-file editing, and edit locks.

Every built-in TRANSID above can be reached via EXEC CICS LINK PROGRAM(name) and EXEC CICS XCTL PROGRAM(name) as well as by typing it at the blank prompt. The two paths share one dispatch helper, so the per-built-in auth gate (dev for ISPF/CECI, admin for CEDA/IDCA) fires uniformly regardless of how the operator reached it. CEMT carries a mixed model — the outer menu is open to any signed-on user (INQUIRE TRANSACTION, MONITOR), and the CONTROLBLOCKS / PERFORM sub-trees gate on admin inside `walk()`. The ACL check at the LINK / XCTL boundary refuses a denied caller before the built-in's Run is reached: the LINK returns the access-denied error to `programControl` which surfaces PGMIDERR via EIBRESP, and the XCTL renders the familiar 403 task-error screen — no internal refusal paint hijacks the caller's session. Built-ins do not have a COMMAREA-style interface; LINK PROGRAM('CEMT') COMMAREA(VAR) returns

an empty COMMAREA (VAR is set to "" on return) and the caller's sess.Commarea is restored unchanged. CSSN is the one exception: it refuses LINK/XCTL with a short message because the auth-prompt loop cannot be wedged into a caller's task stack sanely; operators sign on at a blank prompt.

Chapter 29. Worked examples

A. Producer / consumer over a TS queue

PROD writes one item per ENTER from an interactive map; CONS reads with the implicit per-task cursor. PF4 in CONS deletes the queue; PF5 rewinds the cursor by chaining back to itself (RETURN TRANSID('CONS')), so the dispatcher's task-end hook clears the cursor.

```
/* PROD: write one item per ENTER */
ADDRESS CICS
DO FOREVER
  EXEC CICS SEND      MAP('PROD1') FROM(SCR.) ERASE END-EXEC
  EXEC CICS RECEIVE MAP('PROD1')              END-EXEC
  IF C2X(EIBAID) = 'F3' THEN EXEC CICS RETURN END-EXEC
  EXEC CICS WRITEQ TS QUEUE(MAP.QNAME) FROM(MAP.PAYLOAD) END-EXEC
END
```

```
/* CONS: cursor-advancing read */
ADDRESS CICS
DO FOREVER
  EXEC CICS SEND      MAP('CONS1') FROM(SCR.) ERASE END-EXEC
  EXEC CICS RECEIVE MAP('CONS1')              END-EXEC
  IF C2X(EIBAID) = 'F3' THEN EXEC CICS RETURN END-EXEC
  EXEC CICS READQ TS QUEUE(QNM) INTO(REC) ITEM(GOTI) END-EXEC
  IF EIBRESP = 26 THEN /* ITEMERR – end of queue */ NOP
END
```

B. SEND TEXT + RECEIVE – no BMS map

GETC takes a customer number from the operator's command line (GETC 100), reads the customers KSDS, and paints the record as free-form text. Because **3270 has no LF**, the body is built as a flat row-major buffer — each logical line padded to the 80-column screen width with LEFT(s,80).

```
/* GETC: lookup-and-display, no map */
ADDRESS CICS
EXEC CICS RECEIVE INTO(BUF) END-EXEC          /* "GETC 100" */
PARSE VAR BUF TID CKEY .
EXEC CICS READ FILE('customers') INTO(REC) RIDFLD(CKEY) END-EXEC
PARSE VAR REC NM '|' AD '|' CY '|' PH
TXT = LEFT('Customer #' || CKEY, 80)          /* row 0 */
TXT = TXT || LEFT('', 80)                     /* row 1 (blank) */
TXT = TXT || LEFT('Name: ' || NM, 80)        /* row 2 */
TXT = TXT || LEFT('Address: ' || AD, 80)      /* row 3 */
EXEC CICS SEND TEXT FROM(TXT) ERASE END-EXEC
EXEC CICS RETURN END-EXEC
```

The COBOL companion runtime/cobol/exam.cob shows the same RECEIVE INTO pattern with UNSTRING ... DELIMITED BY ' ' doing the tokenisation:

```
EXEC CICS RECEIVE INTO(WS-INPUT) LENGTH(WS-LEN) END-EXEC
UNSTRING WS-INPUT DELIMITED BY ' '
    INTO WS-TID WS-A WS-B WS-C
END-UNSTRING.
```

C. Browse with GENERIC prefix

A paginated customer list filtered by a 3-character key prefix:

```
EXEC CICS STARTBR FILE('CUSTOMERS')
    RIDFLD('NY-')
    GENERIC KEYLENGTH(3)
END-EXEC
DO FOREVER
    EXEC CICS READNEXT FILE('CUSTOMERS') INTO(REC) RIDFLD(K) END-EXEC
    IF EIBRESP = 20 THEN LEAVE /* ENDFILE */
    CALL ADD_ROW K, REC
END
EXEC CICS ENDBR FILE('CUSTOMERS') END-EXEC
```

D. Synchronous LINK with COMMAREA

CUST LINKs to CUSV to validate a customer number:

```
SAV = CKEY /* in: number to validate */
EXEC CICS LINK PROGRAM('CUSV') COMMAREA(SAV) END-EXEC
PARSE VAR SAV STATUS '|' MSG /* out: status, message */
IF STATUS <> 'OK' THEN ...
```

CUSV reads DFHCOMMAREA, runs its check, and writes the result back into DFHCOMMAREA before RETURN. The dispatcher hands the final value back to the caller's SAV.

E. Sequential import via READQ TD + WRITE FILE

The ORDR transaction (runtime/cobol/ordr.cob) demonstrates the "text file → VSAM" pattern that the tmp_dir sandbox is built for. The sample data ships in runtime/tmp/orders.sample.txt:

```
C0000001|WIDGET-A|10|19.95
C0000002|WIDGET-B|2|249.00
...
```

The loop body is just two EXEC CICS calls and an UNSTRING:

```

IMPORT-ONE.
  MOVE SPACES TO REC.
  EXEC CICS READQ TD QUEUE('orders.sample.txt') INTO(REC) END-EXEC.
  IF EIBRESP = 12 THEN
    MOVE 'Y' TO DONE-FLAG      *> QZERO -- end of file
  END-IF.
  IF EIBRESP = 0 THEN
    COMPUTE N-READ = N-READ + 1
    MOVE SPACES TO CUST-ID PRODUCT QTY PRICE
    UNSTRING REC DELIMITED BY '|'
      INTO CUST-ID PRODUCT QTY PRICE
    END-UNSTRING
    PERFORM WRITE-ORDER
  END-IF.

WRITE-ORDER.
  MOVE SPACES TO OREC.
  STRING PRODUCT DELIMITED BY SIZE
    '|' DELIMITED BY SIZE
    QTY DELIMITED BY SIZE
    '|' DELIMITED BY SIZE
    PRICE DELIMITED BY SIZE
  INTO OREC
  END-STRING.
  EXEC CICS WRITE FILE('ORDERS') FROM(OREC) RIDFLD(CUST-ID) END-EXEC.
  IF EIBRESP = 0 THEN COMPUTE N-WRITE = N-WRITE + 1 END-IF.
  IF EIBRESP = 14 THEN COMPUTE N-DUP = N-DUP + 1 END-IF.

```

Notes:

- The READQ loop tracks EOF with EIBRESP = 12 (QZERO) — the CICS-canonical "queue empty / no more records" signal. The handle is closed automatically at task end.
- EIBRESP = 14 (DUPREC) is treated as a counted no-op, not an error: WRITE FILE against an existing key fails atomically without overwriting, so a re-run of ORDR on the same file is idempotent on the customer-id set.
- ORDR ships public in runtime/transactions.conf so the sample runs out of the box. In a production deployment that imports real data, restrict it to a privileged group (ORDR:cobol:ordr.cob:admin) so a casual operator can't replay the import.
- A REXX equivalent reads the same file via LINEIN and writes via EXEC CICS WRITE. Because both languages share the tmp_dir backend, the sample file works untouched from either side.

F. Menu items gated by QUERY SECURITY

The shipped GUST transaction (runtime/cobol/gust.cob) is a small CRUD UI: A=Add / Q=Query / U=Update / D=Delete / L=List / S=Search. Delete is destructive and only the admin role should be allowed to invoke it. Before 2.7.x the operator had to hard-code IF EIBUSER = 'ADMIN' or IF EIBTRMID = ..., which bypasses the ACL gate that already governs sign-on. QUERY SECURITY lets the program ask the dispatcher's SecurityChecker directly:

```

WORKING-STORAGE SECTION.
COPY DFHAID.
COPY DFHRESP.
COPY DFHVALUE.
...
01 WS-CACR-UP PIC S9(8) VALUE 0.
...
DISPATCH.
...
*> EXEC CICS QUERY SECURITY: refuse D=Delete to callers who
*> are not authorised to UPDATE the GUST resource.
    IF VALID-ACTION = 'Y' AND ACTION = 'D' THEN
        EXEC CICS QUERY SECURITY RESOURCE('GUST')
            UPDATE(WS-CACR-UP) END-EXEC
    IF WS-CACR-UP NOT = DFHVALUE-UPDATABLE THEN
        MOVE 'Delete not permitted for this user.' TO MSG
        MOVE 'N' TO VALID-ACTION
    END-IF
END-IF.

```

The gate runs once per dispatch (the verb is stateless and cheap) and lands before the EVALUATE ACTION switch. Non-admins see the MSG line on the next menu paint; admins fall through to ACT-DELETE unchanged. The same pattern transports to any program with a destructive action — replace 'GUST' with the gated transaction's TRANSID and reuse the four-line block.

See the in-tree `runtime/cobolcopy/DFHVALUE.cpy` for the full constant set, and the security-inquiry sub-chapter of [Chapter 6](#) for the bricks-flavoured access mapping that decides the WS-CACR-UP value.

Appendix A. Adapting to terminal size (mod 2 vs mod 4)

A 3270 connection negotiates one of several screen models — typically mod 2 (24 - 80) or mod 4 (43 - 80). Bricks captures the size from the telnet/3270 handshake into `session.TCB.Rows/Cols`, and exposes it to programs via EXEC CICS ASSIGN:

```
EXEC CICS ASSIGN SCREENHT(SCRH) SCREENWD(SCRW)
                ALTSCRNHT(AH) ALTSCRNWD(AW) END-EXEC
```

SEND MAP always passes the negotiated DevInfo through to `go3270.ScreenOpts.AltScreen`, so the underlying datastream uses Erase/Write Alternate (0x7e) and the terminal clears its full buffer — but a 24-row map painted on a 43-row screen leaves rows 24-42 blank. To use the extra real estate the program has to dispatch to a sized map variant.

Convention

Author one map per model. The mod-2 map keeps its bare name; bigger models add a single-letter suffix:

Model	Suffix	Example map names
mod 2	(none)	HEL01, CUST1, CUST2, CUSTL
mod 3	M	HEL01M, CUST1M, ...
mod 4	L	HEL01L, CUST1L, CUST2L, CUSTLL
mod 5	W	HEL01W, ...

The REXX program builds the suffixed name once and dispatches:

```
EXEC CICS ASSIGN SCREENHT(SCRH) END-EXEC
SUFFIX = ''
IF SCRH >= 43 THEN SUFFIX = 'L'
ELSE IF SCRH >= 32 THEN SUFFIX = 'M'

EXEC CICS SEND MAP('HEL01' || SUFFIX) FROM(SCR.) ERASE END-EXEC
IF EIBRESP = 36 THEN DO /* MAPFAIL - sized variant missing */
    EXEC CICS SEND MAP('HEL01') FROM(SCR.) ERASE END-EXEC
END
```

Three properties make this work without any DSL changes:

1. **Same field names across the family.** `hel01.map` and `hel01l.map` both declare `INF0LINE`, `GREETING`, `FOOTER`. The same `SCR.` stem feeds either one. Bonus tails (e.g. `INF01` / `INF02` / `ACT1` ...) are silently ignored on the smaller map (the renderer only writes values for fields that the map declares).
2. **MAPFAIL fallback.** If the suffixed map isn't on disk, the SEND returns `EIBRESP = 36` and the program retries with the bare name — so an operator who deletes `hel01l.map` doesn't break mod-4 connections; they just see the 24-80 layout.
3. **Paging arithmetic adapts at runtime.** `custl.rexx` reads `SCREENHT` and uses `ROWS_PER_PAGE = 35` on mod 4 vs 15 on mod 2, then picks `CUSTLL` vs `CUSTL` accordingly.

Bricks ships sized variants for every map the demo transactions use:

Mod-2 map (24-80)	Mod-4 sibling (43-80)
runtime/map/helo1.map	runtime/map/helo1l.map — adds system-information + recent-activity panes
runtime/map/cust1.map (menu)	runtime/map/cust1l.map — adds recent-activity history
runtime/map/cust2.map (detail)	runtime/map/cust2l.map — adds an audit-log pane
runtime/map/custl.map (15-row list)	runtime/map/custll.map (35-row list)

To see the difference: connect with `c3270 -model 2 localhost 2300` vs. `c3270 -model 4 localhost 2300`, sign on, and run HEL0 or CUST. The mod-4 view fills the bottom three-quarters of the screen with extra panels.

Appendix B. Pitfalls and idioms

REXX compound-symbol pitfall

STEM.tail with tail an *unset* symbol resolves to STEM.<TAIL> (a literal tail). With tail a *set* symbol it resolves to STEM.<value-of-tail> — so reusing a map field name (OUT.BIRTH = ... when BIRTH is also a local variable) silently writes the wrong tail.

The convention used by runtime/rexx/cust.rexx is to give locals distinct names from map fields (AKT vs ACTION, CKEY vs CUSTNO, BSTR / NDAYS vs BIRTH / DAYS).

COBOL data names across groups

A child name reused across groups is allowed; bare references that match more than one declaration must be qualified with OF / IN (MOVE x OF DET TO y OF SCR). Bare names that are unique across the program continue to work without qualification. Sibling duplicates within a single group are still rejected at parse time because no qualifier can disambiguate them.

COBOL conditions follow IBM, not intuition

Three IBM rules that surprise people, all honoured by bricks (see [Combined conditions and parentheses](#)):

- IF A / 2 > 2 with integer operands truncates the quotient: for A = 5 it compares 2 > 2 and is false. When the fraction matters, give the relation a decimal place outside the divisor: a V field as dividend, or IF A / 2 > 2.0 (true). A decimal divisor does not help — IF A / 2.0 > 2 is still false, because divisors don't count towards dmax.
- IF A NOT = 1 AND 2 means A NOT = 1 AND A NOT = 2 — the NOT carries to the omitted operator.
- STRING ... DELIMITED BY SIZE moves an identifier at its full PIC width, trailing spaces included. Use DELIMITED BY SPACE to stop at the first blank.

3270 has no LF...obviously.

SEND TEXT lays its body out as a flat row-major buffer, every cols bytes wrapping to the next row. Programs that expect newline semantics must pad each logical line to the column width (LEFT(s,80) in REXX, PIC X(80) group children in COBOL) and concatenate.

Kinda obvious, but just making sure it's clear to folks who are new to 3270.

Reset stems before paginated reads

A REXX STEM. accumulator that's reused across pages will leak values from the previous page if not dropped. Use DROP RECS. (with the trailing dot) at the top of each pagination loop.

EIBAID comparison

In REXX, compare C2X(EIBAID) to a hex string: IF C2X(EIBAID) = 'F3' THEN ... (PF3). In COBOL, compare EIBAID directly to a hex literal: IF EIBAID = X'F3'

Appendix C. Quick command reference card

A one-page cheat-sheet of every command family. Use the chapter reference for full syntax, options, and condition codes.

EXEC CICS — terminal I/O (Chapter 4)

Command	Purpose
SEND MAP(name) [FROM(stem.)] [ERASE] [CURSOR(p)]	Paint a BMS map; wait for an AID.
RECEIVE MAP(name) [INTO(stem.)]	Pull the operator's response.
CONVERSE MAP(name) FROM(s) INTO(s) [ERASE]	Fused SEND MAP+RECEIVE MAP.
SEND TEXT FROM(buf) [LENGTH(n)] [ERASE]	Free-form row-major output.
RECEIVE INTO(buf) [LENGTH(v)]	Read the operator's command-line tail.

EXEC CICS — program control (Chapter 5)

Command	Purpose
RETURN [TRANSID(id)] [COMMAREA(d)]	End task; optional chain.
XCTL PROGRAM(name) [COMMAREA(d)]	Transfer control.
LINK PROGRAM(name) [COMMAREA(v)]	Synchronous sub-program call.
ABEND [ABCODE(c)]	Abnormal task termination.
START TRANSID(id) [INTERVAL/TIME] [FROM(d)] [TERMID(t)]	Schedule a deferred fire.
RETRIEVE INTO(v) [LENGTH(v)]	Pull the START payload.

EXEC CICS — system services (Chapter 6)

Command	Purpose
ASSIGN <FIELD>(v) ...	Read EIB / environment fields.
ASKTIME [ABSTIME(v)]	Refresh EIBDATE / EIBTIME (+ optional ABSTIME).
FORMATTIME ABSTIME(s) [DATE / TIME / YYYY... / DAYOFWEEK ...]	Decode an ABSTIME.
QUERY SECURITY RESOURCE(name) [RESCCLASS(c)] [READ/UPDATE/CONTROL/ALTER(v)]	Four-axis access inquiry.
VERIFY PASSWORD USERID(u) PASSWORD(p)	Re-authenticate credentials.
INQUIRE SYSTEM [GMMTEXT(v)]	SIT-level inquiry (v1 honours GMMTEXT only).

EXEC CICS — KSDS files (Chapter 7-8)

Command	Purpose
READ FILE(f) INTO(v) RIDFLD(k) [UPDATE]	Exact-key lookup.
WRITE FILE(f) FROM(d) RIDFLD(k)	Insert record.

Command	Purpose
REWRITE FILE(f) FROM(d)	Replace the READ ... UPDATE record.
DELETE FILE(f) [RIDFLD(k)]	Drop record.
STARTBR FILE(f) [RIDFLD(k)] [GTEQ/EQUAL] [GENERIC KEYLENGTH(n)]	Open browse.
READNEXT FILE(f) INTO(v) [RIDFLD(v)]	Forward one record.
READPREV FILE(f) INTO(v) [RIDFLD(v)]	Backward one record.
RESETBR FILE(f) RIDFLD(k)	Re-position within an open browse.
ENDBR FILE(f)	Close browse.

EXEC CICS – temporary storage / transient data (Chapter 9)

Command	Purpose
`READQ TS QUEUE(q) INTO(v) [ITEM(n)	NEXT]`
WRITEQ TS QUEUE(q) FROM(d) [ITEM(n) REWRITE]	Append or rewrite TS item.
DELETEQ TS QUEUE(q)	Drop a TS queue.
READQ TD QUEUE(q) INTO(v)	Read next line from tmp_dir/q.
WRITEQ TD QUEUE(q) FROM(d)	Append line to tmp_dir/q.
DELETEQ TD QUEUE(q)	Delete a TD file.

EXEC CICS – recovery + condition (Chapter 10)

Command	Purpose
SYNCPOINT	Commit pending unit of work.
SYNCPOINT ROLLBACK	Undo pending unit of work.
HANDLE CONDITION cond(label) ...	Arm condition trap.
IGNORE CONDITION cond ...	Suppress condition trap.
HANDLE AID key(label) ...	Arm AID-key trap.
`HANDLE ABEND LABEL(l)	PROGRAM(p)

EXEC SQL (Chapter 26)

Command	Purpose
EXEC SQL SELECT cols INTO :v, :v FROM ... END-EXEC	Single-row read; +100 = no-data, -811 = >1 row.
EXEC SQL INSERT INTO t (...) VALUES (...) END-EXEC	Insert.
EXEC SQL UPDATE t SET ... WHERE ... END-EXEC	Update; 0 rows is still NORMAL.
EXEC SQL DELETE FROM t WHERE ... END-EXEC	Delete.
EXEC SQL DECLARE c CURSOR FOR <select> END-EXEC	Declare cursor.
EXEC SQL OPEN c END-EXEC	Open cursor.
EXEC SQL FETCH c INTO :v, :v END-EXEC	Next row; +100 at end.
EXEC SQL CLOSE c END-EXEC	Close cursor.
EXEC SQL COMMIT WORK END-EXEC / ROLLBACK WORK	Transaction control.

Command	Purpose
EXEC SQL CONNECT TO 'db' END-EXEC	Per-task database switch.
`EXEC SQL WHENEVER {SQLERROR	NOT FOUND

EXEC CICS WEB — server side (Phase 1)

Command	Purpose
WEB EXTRACT METHOD(v) PATH(v) HOST(v) PORT(v) SCHEME(v) QUERYSTRING(v) ...	Inbound request meta.
WEB EXTRACT CERTIFICATE COMMONNAME(v) ORGANISATION(v) COUNTRY(v) SERIALNUM(v) ISSUER(v)	mTLS peer-cert fields.
WEB EXTRACT TCIPSERVICE(v) PORTNUMBER(v) IPADDRESS(v) CLIENT(v) AUTHENTICATE(v)	Listener inspection.
WEB READ HTTPHEADER(name) VALUE(v) [LENGTH(v)]	Inbound header by name.
WEB STARTBROWSE / READNEXT / ENDBROWSE HTTPHEADER	Walk inbound headers.
WEB READ QUERYPARM(name) VALUE(v)	Query parameter / path capture, percent-decoded.
WEB STARTBROWSE / READNEXT / ENDBROWSE QUERYPARM	Walk parameters.
WEB READ FORMFIELD(name) VALUE(v)	Form-encoded field.
WEB STARTBROWSE / READNEXT / ENDBROWSE FORMFIELD	Walk form fields.
WEB RECEIVE INTO(v) [MAXLENGTH(n)] [LENGTH(v)] [MEDIATYPE(v)] [TYPE(v)]	Read raw request body.
WEB WRITE HTTPHEADER(name) VALUE(v)	Set outbound response header.
WEB SEND FROM(buf) [MEDIATYPE(s)] [STATUSCODE(n)]	Emit response.
WEB PARSE URL URL(s) SCHEMENAME(v) HOST(v) PORT(v) PATH(v) QUERYSTRING(v)	URL splitter.
WEB CONVERTTIME DATESTRING(s) ABSTIME(v)	RFC 1123 / 850 / asctime → ABSTIME.
WEB RETRIEVE DOCTOKEN(v)	Inbound body as DOCUMENT.

EXEC CICS WEB — client side (Phase 2)

Command	Purpose
WEB OPEN HOST(s) [PORT(n)] [SCHEME(s)] SESSTOKEN(v)	Open client session.
WEB OPEN URIMAP(name) SESSTOKEN(v)	Open via URIMAP.
WEB CONVERSE SESSTOKEN(t) METHOD(m) PATH(p) [FROM(b)] INTO(v) [STATUSCODE(v)]	One-shot send + receive.
CONVERSE WEB ...	Alias for WEB CONVERSE.
WEB SEND SESSTOKEN(t) METHOD(m) PATH(p) [FROM(b)] [QUERYSTRING(s)] [MEDIATYPE(s)]	Stage request.
WEB SEND URIMAP(name) METHOD(m) PATH(p) ... INTO(v)	One-shot via URIMAP.
WEB RECEIVE SESSTOKEN(t) INTO(v) [STATUSCODE(v)] [MEDIATYPE(v)] [MAXLENGTH(n)]	Fire staged request.
WEB CLOSE SESSTOKEN(t)	Release session.

Command	Purpose
WEB WRITE HTTPHEADER(name) VALUE(v) SESSTOKEN(t)	Set outbound <i>request</i> header.
WEB READ HTTPHEADER(name) VALUE(v) [LENGTH(v)] SESSTOKEN(t)	Read response header.
WEB STARTBROWSE / READNEXT / ENDBROWSE HTTPHEADER SESSTOKEN(t)	Walk response headers.
WEB EXTRACT SESSTOKEN(t) [SCHEME(v)] [HOST(v)] [PORT(v)] [PATH(v)]	Inspect session endpoint.
WEB EXTRACT URIMAP(name) [SCHEME(v)] [HOST(v)] [PORT(v)] [PATH(v)]	Inspect URIMAP entry.

EXEC CICS DOCUMENT

Command	Purpose
DOCUMENT CREATE DOCTOKEN(v) [SYMBOLLIST(s) [LISTLENGTH(n)]] [DELIMITER(s)]	New empty document.
DOCUMENT INSERT DOCTOKEN(t) FROM(b) [LENGTH(n)] [AT(pos)]	Append / splice bytes.
DOCUMENT INSERT DOCTOKEN(t) SYMBOL(name)	Substitute from symbol list.
DOCUMENT INSERT DOCTOKEN(t) DOCUMENT(other-tok) [AT(pos)]	Splice another document.
DOCUMENT SET DOCTOKEN(t) SYMBOLLIST(s) [DELIMITER(d)]	Re-bind symbol list.
DOCUMENT RETRIEVE DOCTOKEN(t) INTO(b) LENGTH(v)	Read assembled body.
DOCUMENT DELETE DOCTOKEN(t)	Release document.
WEB SEND DOCTOKEN(t) [STATUSCODE(n)] [MEDIATYPE(s)]	Emit document as response.

Standard copybooks

Copybook	Brings in
DFHAID	AID byte constants (ENTER, PF03, PA1, ...) and their DFH... aliases.
DFHRESP	EIBRESP condition constants (RESP-NORMAL / DFHRESP-NORMAL, etc.).
SQLCA	SQLCODE + SQLSTATE constants (SQL-OK, SQLSTATE-DUP-KEY, ...).
DFHWBSC	HTTP status codes (DFHRESP-WB-OK, ...).
DFHWBUH	Common header-name literals (WB-CONTENT-TYPE, ...).
DFHWBMT	MIME-type literals (WB-MT-JSON, ...).
DFHWBMETH	HTTP method literals (WB-GET, WB-POST, ...).
DFHWBSI	Session-token + body buffer templates.
DFHWBUM	URIMAP record fields.
DFHWBHB	Header-browse helpers (DFH-WB-HDR-NAME, ...).
DFHWBCC	TLS / mTLS cert-extract pack.
DFHDCDOC	DOCUMENT token + delimiter constants.

EIB fields (Chapter 11)

Field	Set by	Meaning
EIBAID	SEND MAP / RECEIVE MAP	AID byte.

Field	Set by	Meaning
EIBCP0SN	SEND MAP / RECEIVE MAP	1-based cursor position.
EIBCALEN	dispatcher entry	DFHCOMMAREA length.
EIBTRMID	dispatcher	Terminal ID.
EIBRESP / EIBRESP2	every EXEC CICS	Response code + secondary.
EIBDATE / EIBTIME	ASKTIME	0CYDDDD / HHMMSS.
EIBABCODE	abend trap	4-char abend code.
RC	every EXEC CICS	Mirror of EIBRESP.

SQLCA fields (Chapter 26)

Field	Type	Meaning
SQLCODE	PIC S9(9)	0 = OK, +100 = no-data, -N = error.
SQLSTATE	PIC X(5)	5-character ISO state code.
SQLERRMC	PIC X(70)	Human-readable PG error text.
SQLERRD3	PIC S9(9)	Rows-affected count after INSERT / UPDATE / DELETE.
SQLWARN0	PIC X(1)	'W' if any warning fired (Bricks rarely sets).

End of publication.